

FOR572 | ADVANCED NETWORK FORENSICS:
THREAT HUNTING, ANALYSIS, AND INCIDENT RESPONSE
GIAC Network Forensic Analyst (GNFA)

572.5

Encryption, Protocol Reversing, OPSEC, and Intel



© 2022 Lewes Technology Consulting, LLC and Mat Oldham. All rights reserved to Lewes Technology Consulting, LLC, Mat Oldham, and/or SANS Institute.

PLEASE READ THE TERMS AND CONDITIONS OF THIS COURSEWARE LICENSE AGREEMENT ("CLA") CAREFULLY BEFORE USING ANY OF THE COURSEWARE ASSOCIATED WITH THE SANS COURSE. THIS IS A LEGAL AND ENFORCEABLE CONTRACT BETWEEN YOU (THE "USER") AND SANS INSTITUTE FOR THE COURSEWARE. YOU AGREE THAT THIS AGREEMENT IS ENFORCEABLE LIKE ANY WRITTEN NEGOTIATED AGREEMENT SIGNED BY YOU.

With this CLA, SANS Institute hereby grants User a personal, non-exclusive license to use the Courseware subject to the terms of this agreement. Courseware includes all printed materials, including course books and lab workbooks, as well as any digital or other media, virtual machines, and/or data sets distributed by SANS Institute to User for use in the SANS class associated with the Courseware. User agrees that the CLA is the complete and exclusive statement of agreement between SANS Institute and you and that this CLA supersedes any oral or written proposal, agreement or other communication relating to the subject matter of this CLA.

BY ACCEPTING THIS COURSEWARE, USER AGREES TO BE BOUND BY THE TERMS OF THIS CLA. BY ACCEPTING THIS SOFTWARE, USER AGREES THAT ANY BREACH OF THE TERMS OF THIS CLA MAY CAUSE IRREPARABLE HARM AND SIGNIFICANT INJURY TO SANS INSTITUTE, AND THAT SANS INSTITUTE MAY ENFORCE THESE PROVISIONS BY INJUNCTION (WITHOUT THE NECESSITY OF POSTING BOND) SPECIFIC PERFORMANCE, OR OTHER EQUITABLE RELIEF.

If User does not agree, User may return the Courseware to SANS Institute for a full refund, if applicable.

User may not copy, reproduce, re-publish, distribute, display, modify or create derivative works based upon all or any portion of the Courseware, in any medium whether printed, electronic or otherwise, for any purpose, without the express prior written consent of SANS Institute. Additionally, User may not sell, rent, lease, trade, or otherwise transfer the Courseware in any way, shape, or form without the express written consent of SANS Institute.

If any provision of this CLA is declared unenforceable in any jurisdiction, then such provision shall be deemed to be severable from this CLA and shall not affect the remainder thereof. An amendment or addendum to this CLA may accompany this Courseware.

SANS acknowledges that any and all software and/or tools, graphics, images, tables, charts or graphs presented in this Courseware are the sole property of their respective trademark/registered/copyright owners, including:

AirDrop, AirPort, AirPort Time Capsule, Apple, Apple Remote Desktop, Apple TV, App Nap, Back to My Mac, Boot Camp, Cocoa, FaceTime, FileVault, Finder, FireWire, FireWire logo, iCal, iChat, iLife, iMac, iMessage, iPad, iPad Air, iPad Mini, iPhone, iPhoto, iPod, iPod classic, iPod shuffle, iPod nano, iPod touch, iTunes, iTunes logo, iWork, Keychain, Keynote, Mac, Mac Logo, MacBook, MacBook Air, MacBook Pro, Macintosh, Mac OS, Mac Pro, Numbers, OS X, Pages, Passbook, Retina, Safari, Siri, Spaces, Spotlight, There's an app for that, Time Capsule, Time Machine, Touch ID, Xcode, Xserve, App Store, and iCloud are registered trademarks of Apple Inc.

PMP® and PMBOK® are registered trademarks of PMI.

SOF-ELK® is a registered trademark of Lewes Technology Consulting, LLC. Used with permission.

SIFT® is a registered trademark of Harbingers, LLC. Used with permission.

Governing Law: This Agreement shall be governed by the laws of the State of Maryland, USA.



Encryption, Protocol Reversing, OPSEC, and Intel

©2022 Lewes Technology Consulting, LLC and Mat Oldham | All Rights Reserved | Version # FOR572_H01_02

Authors:

Phil Hagen, Lewes Technology Consulting, LLC

phil@lewesotech.com | @PhilHagen

Mat Oldham

mat.oldham@gmail.com | @roujisecurity

3957833237c0be407c3b1bf10529734f0aee3e8df326d053b1a0a7fccd79365e

TABLE OF CONTENTS

PAGE

Encoding, Encryption, and SSL/TLS	3
Lab 5.1: SSL/TLS Profiling	41
Meddler-In-The-Middle	44
Network Protocol Reverse Engineering	62
Lab 5.2: Undocumented Protocol Features	87
Investigation OPSEC and Threat Intel	90
Lab 5.3: Mini-Comprehensive Investigation	112
Capstone Challenge Preparation	116

This page intentionally left blank.



Encoding, Encryption, and SSL/TLS

This page intentionally left blank.

Encoding

- Encoding – Translates a piece of input data into format more suitable for transport/storage
- Decoding – Translates encoded data to its original format—message is the same as before encoding
- Covers a wide range of needs:
 - Text encoding: ASCII, Unicode, UTF-7, UTF-8, UTF-16
 - Binary encoding: Base64
 - Compression: MPEG2, MPEG4, MP3, H.264

Developers and security professionals often mistakenly use the terms “encoding” and “encryption” interchangeably. Although both operate on a given piece of data, the end results are significantly different. When misused, this error can cause serious security issues.

Encoding is the process of translating a piece of data into a more suitable, usually publicly available format. Decoding is the process of translating the encoded data back to its original format. Encoding is used in many different ways within our day-to-day activities:

- Textual representation of data, such as ASCII, Unicode, UTF-7, UTF-8, UTF-16
- Translating unprintable binary data into a printable format, such as Base64/MIME encoding
- Various compression algorithms: MPEG2, MPEG4, MP3, H.264

Technically, an encoding algorithm is known as a “cipher” that is used to generate an encoded message called “ciphertext”. The decoded equivalent text is referred to as “plaintext”. While technically correct terms, we will use the terms “input message” and “output message” instead to avoid confusion with encryption, which will be addressed in a moment.

Base64 Encoding (I)

- Translates binary data to printable ASCII characters
- Uses 64 output characters:
 - A-Z, a-z, 0-9, (“+” and (“/” or “-”))
 - May use “=” for padding

```
--Apple-Mail=_F387EA08-6CDF-4EFE-B115-E37BE6319094
Content-Type: application/pdf;
    name="LTC Invoice 2018.0079.pdf"
Content-Transfer-Encoding: base64

JVBERI0xLjQKJcfsj6IKNSAwIG9iago8PC9MZW5ndGggNiAwIFIvRmlsdGVyIC9GbGF0ZURlY29k
ZT4+CnN0cmVhbQp4nK1YWXMcRQx+318xVfCwQ7Gdvg/ecAihKxxxtooHigfHhMTgg8QOqfx71KfU
afvnirMQvOB5U0yGc2b15JRmetpern5b/zBvJDNKOrG+mjeCmcCtWv87g54JoLe+xtVz1D1DhZeZ
...
UgovsUQgWzxDNjYzMTM5NENDNDhBQT11NjNGMEU5OTkyOTwRDg1Nj48QzY2MzEZOTRDQzQ4QUEY
NTYzRjBFOTk5Mjk0MEQ4NTY+XQo+PgpzdGFydHhyZWYKMTQ2NTMKJSVFT0YK
--Apple-Mail=_F387EA08-6CDF-4EFE-B115-E37BE6319094
...
```

Base64 encoding is one of the most common encoding formats used on the Internet. Base64 encoding is commonly used to encode binary data that needs to be transferred or stored in a textual (usually ASCII) data format. This ensures the data remains intact during storage and transmission. Some examples include encoding attachments associated with emails within the MIME format, storing complex data within XML structures, and HTTP transactions.

At a high level, base64 encoding takes 3 bytes of input data and encodes it into 4 bytes of output data using a predetermined alphabet. The alphabet consists of a total of 64 “printable” characters:

- 26 uppercase Latin/English letters (A–Z)
- 26 lowercase Latin/English letters (a–z)
- 10 numerical digits (0–9)
- 2 special characters (“+” and “/”); however, the “–” may be used interchangeably with the “/” because some fields may contain base64 data as well as filenames/paths or URI strings

Because this adds one extra byte of output message size for every three bytes of the input message, the overall output message size will be approximately 133% of the input message. In addition, if the number of bits in the input message is not divisible by 24, additional padding is added to the encoded message. One “=” is added to the end of the base64-encoded message for each 8 bits of padding added to the input message to force a 24-bit alignment. Examples of this are below:

- “SANS FOR572 ROCKS!” => “U0FOUyBGT1I1NzIgUk9DS1Mh”
- “SANS FOR572 ROCKS” => “U0FOUyBGT1I1NzIgUk9DS1M=”
- “SANS FOR572 ROCK” => “U0FOUyBGT1I1NzIgUk9DSw==”

Note, however, that the padding is not strictly required to decode the message. The number of missing bytes is trivial to calculate based on the size of the encoded message. Some base64 implementations omit padding entirely.

References:

<https://for572.com/98ug3>

Base64 Encoding (2)

- Alphabet

Index	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	16	17	18	19	20	21	22	23	24	25
Character	A	B	C	D	E	F	G	H	I	J	K	L	M	N	O	P	Q	R	S	T	U	V	W	X	Y	Z
Index	26	27	28	29	30	31	32	33	34	35	36	37	38	39	40	41	42	43	44	45	46	47	48	49	50	51
Character	a	b	c	d	e	f	g	h	i	j	k	l	m	n	o	p	q	r	s	t	u	v	w	x	y	z
Index	52	53	54	55	56	57	58	59	60	61	62	63														
Character	0	1	2	3	4	5	6	7	8	9	+	/														

Input length is not evenly divisible by 3, so padding added to output message for clarity to the decoder. Here, the one byte in last input group results in two “=” padding bytes added to output message

- Example encoding

Input Hex	0x1e	0x08	0xb6	0x84
Binary	0 0 0 1 1 1 1 0 0 0 0 0 1 0 0 0 1 0 1 1 0 1 1 0 1 0 0 0 0 1 0 0 0 0 0 0 0 0			
Output Index	7	32	34	54
Base64	H	g	i	2

As an example, say we want to convert the bytes “0x1e0x08b684” into base64 format with a standard output alphabet. For each character, we first translate the text value into its corresponding ASCII hex representation. The second row in the encoding table reflects this bit stream. Of course, 8 bits are required for each byte of input data. However, the output alphabet consists of only 64 characters. Therefore, each character in the output message requires just 6 bits to represent so we take the first 6 bits (000111) and translate that to a corresponding number, which, in this case, is 7. The value 7 is the index into the base64 output, corresponding to the “H” character. Therefore, the first value of our base64-encoded string is the “H”.

There are still 2 bits (10) leftover that were not used as part of the translation to the “H” character. However, we need a total of 6 bits to represent the base64 output characters, so we take the remaining 2 bits of the first input character and append the first 4 bits of the second input byte, which gives us the next 6-bit value (100000). This translates to the number 32, which is then used as an index into the base64 output alphabet—the character “g”.

This process continues as above for each character, until we reach the final byte of the input message. Because the output alphabet requires 6 bits, we pad the final 2 bits of the input character “0x84” with 0 bits, resulting in a “000000”. This maps to the “A” character in the base64 output alphabet.

Padding, while not explicitly required for the decoding process, is generally included as a means of ensuring that the decoding is accurate. Since the input length is required to be a multiple of three-byte blocks, one equal sign is added for each missing byte in the final input block. In the example above, there is a single byte left in the final input block, meaning two padding bytes will be added to the output message^[1].

Decoding the output message back to its original format would simply use the same process in reverse.

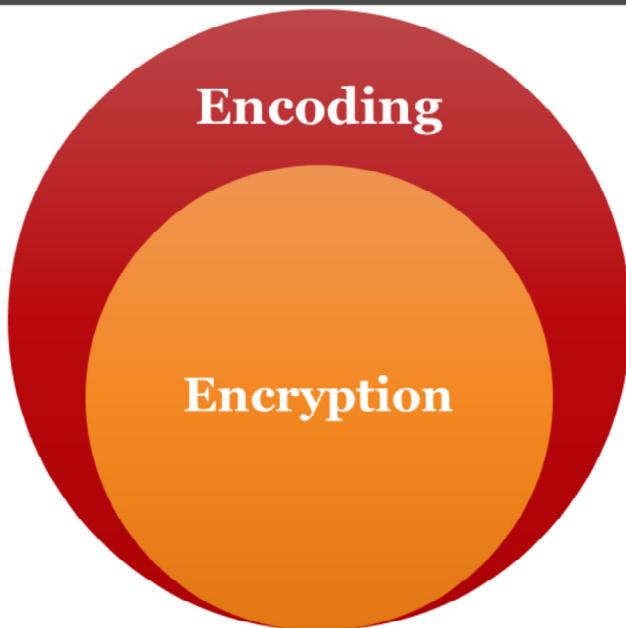
As you may have recognized, successfully encoding and decoding messages using this algorithm fundamentally depends on using the specified base64 alphabet. If different characters were used, or the standard character set was jumbled in some way, only the parties that know that alphabet mapping would be able to decode the message.

Some attackers have used the base64 algorithm and character set with a nonstandard mapping as a sort of rudimentary **encryption** algorithm because the decoding process requires secret information—or a key—consisting of that mapping.

References:

- [1] <https://for572.com/n6l7r>

Encryption (I)



- Subset of encoding algorithms
- Primary purpose is to hide message from unauthorized recipients

As mentioned prior, encoding is typically used to translate data into a format that is more appropriate for a given transport or storage method. A subset of encoding algorithms is used for encryption, where the purpose of the encryption algorithm is to secure the contents of messages between sender and receiver. Although all encryption algorithms fall into the larger set of encoding algorithms, not all encoding algorithms are encryption algorithms.

Encryption is the process of disguising a message in such a way as to hide its contents from unauthorized readers. The technical terms “plaintext” and “ciphertext” still apply in this arena and have the same definitions. However, in this case, the specialized form of encoding process is most accurately called “encryption”, and the reverse—getting plaintext back—is called “decryption”. The encryption process consists of multiple components:

- The message (M) is the data that needs to be hidden from unauthorized users.
- A key (K) is a unique value used by the encryption/decryption algorithms to generate a more secure ciphertext.
- The encryption algorithm (E) is the process used to convert the message into ciphertext.
- The decryption algorithm (D) is the process used to convert the ciphertext back to plaintext.

This essentially allows for the following:

- $D(E(M)) = M$
(Read: The decryption of the encrypted message M is the original message M)

Encryption (2)

- Strength of encrypted message relies on two factors:
- Strength of algorithm
 - Assumes potential public knowledge of the algorithm
 - Depends on low likelihood of mathematical collisions
- The overall strength of the key
 - Depends on the length of the key
 - Depends on the randomness of the key

These concepts have remained unchanged since dawn of cryptography

The strength and resilience of any encrypted message primarily depends on two factors: The strength of the algorithm and the strength of the key used in the algorithm. A weakness in either or both results in a significant weakness risk to the message's secrecy.

The strength of the encryption/decryption algorithm

Custom algorithms that are known only to the sender and receiver of the messages are called restricted algorithms. In these, the strength of an encrypted message is based on how many individuals know the algorithm used to encrypt the message. If someone leaves the group or the algorithm is otherwise compromised, all users of the group must change the algorithm. This method is commonly used in low-security applications. This method can often be defeated easily through reverse engineering of the applications developed. However, a strong encryption algorithm assumes that a loss of the algorithm itself can still leave the integrity of messages intact, as there is a key that remains secret. While restricted algorithms may be considered strong, this is due to the multiple factors used to protect the messages.

Strong algorithms are also fundamentally rooted in the expectation that an encrypted message has sufficient entropy—that is, a properly encrypted message should be indistinguishable from random data. Any repetition in that output message represents an exponential weakness in the overall encryption system.

The strength of the key

In cases where the algorithm is known by users outside the group, a key is used to add an extra layer of security. This is common for mass-produced applications because this allows for the algorithm to be scrutinized for weaknesses, or that the encryption system can withstand the loss of a restricted algorithm. The strength of a key is typically determined by the size and randomness of the key. When multiple keys can produce the same output, a collision occurs.

It should also be noted that these characteristics of a strong encryption system are not new—they have remained foundational truths since the development of any known encryption techniques. While the future may hold new dynamics to such long-proven truths, those scenarios are, at this point, still hypothetical.

A great deal of research (incredible understatement) is available on these topics. While some ranges to the PhD level, quite a lot is very approachable—even those with deep technical insights. Several resources that may be of interest for those wishing to dabble in the fascinating waters of cryptography include Auguste Kerckhoffs's "La Cryptographie Militaire"^[1], from the late 1800s, or the much more recent Simon Singh's "The Code Book"^[2].

In any case, this course will cover an admittedly tiny sliver of this vast topic, focusing on the topics that bear the most impact on network forensics and incident response.

References:

- [1] <https://for572.com/fdkgt>
- [2] <https://for572.com/eifmn>

Symmetric Key Encryption

- Single “private” key is used for both encryption and decryption
- Strength:
Computationally cheap
- Weakness: 100% depends on key security
- Two categories:
 - Stream ciphers
 - Block ciphers

Step 2: Alice provides key to Bob
(Securely and separately from encrypted data)



Step 1: Alice selects key and encrypt data



Step 3: Bob uses key to decrypt data



There are two primary types of key-based algorithms: symmetric and asymmetric.

In a symmetric encryption algorithm, the same key is used by all users to encrypt and decrypt messages. Due to the shared key, typically symmetric key algorithms are less computationally expensive, which is a symmetric key algorithm’s biggest strength.

However, the biggest drawback is that the security of this type of encryption system is wholly dependent on the security of the key. Both the sender and receiver must agree on the key; however, if the key is ever exposed, the encryption system is open to attack. If one individual from the group leaves, every member must change their key. Even if the group does change the key, all previous communications using the previous key value can still be decrypted by any party with knowledge of the old key value.

To demonstrate symmetric key encryption, assume Alice wants to send Bob an encrypted message. Alice types the message she wants to send to Bob and encrypts it with a password (or encryption key). Alice sends the message to Bob via email and calls him to tell him the password. Bob then uses the password to decrypt and read Alice’s message.

However, any party that acquires the key at any time would be able to decrypt all messages ever encrypted with that key value—so key security is absolutely paramount.

There are two specific types of symmetric encryption algorithms:

- Stream ciphers: Algorithms that operate on a single byte of the plaintext message at a time
- Block ciphers: Algorithms that operate on a group of bits or bytes at a time

Cipher Modes

Stream Ciphers:

Encrypt one byte at a time

- Advantages

- Ideal for variable length
- Self-synchronous mode adds randomness

- Disadvantages

- Corrupted data stream breaks cipher
- Typically slower

Block Ciphers:

Encrypt n -bit blocks at a time

- Advantages

- Can validate integrity and confidentiality

- Disadvantages

- Identical blocks produce same results

Stream ciphers are encryption algorithms that operate on one byte or character at a time. Stream ciphers typically start with a key called the seed. The key is then used to generate a keystream of a fixed or variable size. The keystream is then used to either encrypt or decrypt the message one byte at a time. During encryption or decryption, the keystream may be updated. Examples of stream ciphers include the German Enigma, Wireless Communications (WEP), RC4, and several high-sensitivity communication systems.

Overall, there are several advantages to stream ciphers:

- They are ideal for variable length messages.
- They are less susceptible to cryptanalysis attacks because each portion of the message is encrypted with a different portion of the key.

There are several disadvantages as well:

- A transmission error with a stream cipher breaks the rest of the transmission. This requires both data and keystream to be resynchronized to continue the stream cipher.
- Typically, because a stream cipher operates on an individual character, the overall process is slower. However, stream ciphers can typically be implemented in specialized hardware to significantly increase the encryption and decryption process.

There are two common ways to generate the keystream:

- A keystream is generated independently of the message (this is called synchronous ciphertext).
- A keystream is derived from the previous N number of bits (this is called self-synchronous ciphertext).

In synchronous ciphertext, the keystream is generated independently of the message. To decrypt a message, the attacker would not only need to intercept the message, but they also must have knowledge of the keystream used to encrypt the message. A good example would be the German Enigma Machine during World War II. The Enigma required both the transmitter and receiver to reset their machines during a set period of time. The transmissions between two machines required the same exact settings (rotors, wiring, and plugs) that represented the key to be consistent between the two machines. As the sender transmitted messages to the receiver, a keystream was generated based on the rotation of the rotors for each character that was typed over a given time period.

In self-synchronous ciphertext, the keystream is generated based on prior message bytes. This has an added value over synchronous ciphertext where any modification of any part of the stream will affect decryption of the rest of the stream. Essentially, this provides a “tamper-evident” functionality. Given the plaintext data and an initial key, the first byte of data is XOR’d with the initial key. The result of this operation is then used as the key to encrypt the next character of data. For the entire length of the message, each subsequent byte of data is XOR’d with the result of the previous operation. Because the keystream (i.e., the list of bytes used as a key for each XOR operation) is generated based on the prior message, the entire algorithm is self-synchronous.

Block ciphers are encryption algorithms that operate on a fixed-length block of data at a time. Typically, block size is a power of two—most commonly 64 or 128 bits. In most cases, the cipher block used to encrypt the block of data is the same size. When the available input data is shorter than the cipher block, padding—usually 0x00 bytes—is added to the end of the input data to match the necessary cipher block size.

A block cipher’s main advantage over a stream cipher is twofold:

- When encrypting or decrypting a message that is longer than the cipher block length, the message must be partitioned into segments that are the same length as the cipher block. Additional concepts adapted from stream ciphers, such as input vectors and feedback, help increase the overall strength of the cipher.
- Certain cipher block algorithms have built-in modes of operation that can check both the integrity of the data being encrypted or decrypted in addition to its confidentiality. Stream ciphers typically don’t have this ability.

Block ciphers do have some weaknesses, however:

- Identical blocks of data being encrypted typically produce the same results.
- They are more susceptible to cryptanalysis attacks because the key may be the same for all blocks of plaintext, resulting in duplicate inputs creating duplicate outputs.

Block ciphers typically implement one of several different types of modes of operation. A mode of operation essentially modifies how the block cipher algorithm works from one block to another, thus adding randomization to the data. To address shortcomings in a standard block-cipher algorithm, designers have implemented two modes of operation that offer significant improvements: initialization vectors and feedback.

An initialization vector is a random or pseudorandom block of data that is used to “initialize” a cipher block or used as part of a mode of operation. Because a basic block cipher does nothing more than encrypt or decrypt a block of data based on a key, without any modifications, the same block of data would always be encrypted the same way, thus subjecting the cipher to cryptanalysis attacks. An initialization vector adds a factor of randomness to the overall algorithm.

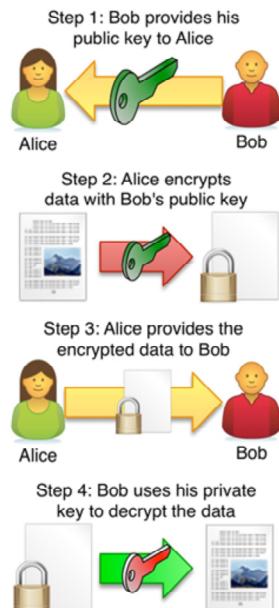
Feedback is the process of the initialization vector and/or “cipher” block changing after each encryption or decryption operation. There are several different modes of feedback that are beyond the scope of this talk. However, it is important to understand that after the encryption operation, either the initialization vector or cipher block is modified. This operation implements stream-cipher-like capabilities into the block cipher algorithm.

The diagram depicts the process of encrypting data with a block cipher that uses an initialization vector and a feedback loop:

- An initialization vector is applied to a cipher block to generate an initial set of random or pseudorandom data.
- A plaintext message is partitioned into one or more blocks of data.
- The block of data is processed through the encryption algorithm generating encrypted text.
- Upon completion, the cipher block is updated based on the feedback or mode of operation determined by the user.
- The next block of data is processed through the algorithm and encrypted with the new cipher block.

Public Key Encryption

- Advantages
 - Private keys stay safe, public key shared
 - Provides authentication using same keys
- Disadvantages
 - More computationally intensive than symmetric key encryption
 - Susceptible to MITM attacks
- Typical use case
 - Use asymmetric to exchange key material
 - Use shared key material for symmetric



Asymmetric encryption systems are more commonly known as “public key encryption” systems.

Public key encryption uses two separate keys, one of which is kept secret and one of which is shared publicly. Although there are two different keys, they are based on a mathematical relationship such that one key can encrypt a message while the other key can decrypt the message. In public key encryption, the public key can be shared with other users, while the private key must remain secret. This is public key encryption’s biggest strength—users do not need a secure means of exchanging key material before using the encryption system.

Another key strength is that these same keys can be used in the opposite direction to provide message authentication. That is, a digest (usually a hash of the message, along with some additional values) that is encrypted with a private key can only be decrypted by the corresponding public key. This provides the recipient with the assurance that the message was sent by a party with the private key material that is mathematically associated to the known public key.

However, there are several weaknesses associated with public key encryption. Due to the unique nature of the keys, they are more computationally costly than their counterparts within symmetric key encryption. Although longer keys reduce the likelihood of a brute-force attack on the system, the longer key increases the computational costs associated with the encryption system. In addition, like any encryption system, public key encryption is also susceptible to meddler-in-the-middle (MITM) attacks, because all parties must still have a high degree of trust in the public key system. (This is why many public keys are signed by multiple parties, which increases the “web of trust” needed for proper cryptographic exchange.)

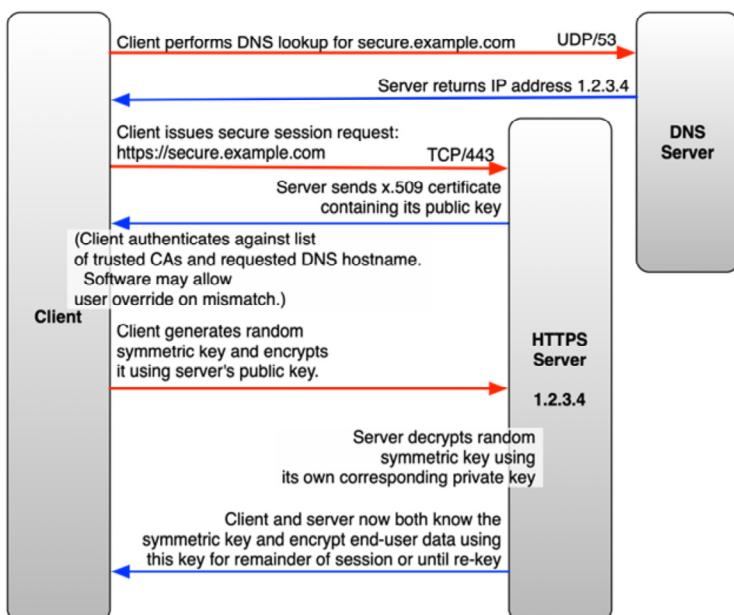
Common examples of public key encryption systems include SSL/TLS, SMIME, and PGP/GPG.

Symmetric and asymmetric encryption are frequently used together to mitigate the mathematical complexity of such a system. Both parties may use the more expensive asymmetric encryption to exchange a small message, without significant computational impact. The message that is exchanged in this phase becomes a symmetric key used for the remainder of the conversation.

The diagram here depicts an example of public key encryption:

- Alice acquires Bob's public key.
- Alice generates a message and encrypts it with Bob's public key.
- Alice sends the encrypted message to Bob.
- Bob decrypts Alice's message using his own private key.

Basic SSL/TLS Process



- Negotiation handshake followed by secure communication over encrypted channel
- In-conversation renegotiation

After identifying the IP address for the intended hostname via DNS or a local cache, the client opens a TCP socket to the intended destination—in this case, what appears to be an HTTPS connection via TCP/443. The SSL/TLS exchange begins with a negotiation called a handshake followed by the client and server passing data between each other through the encrypted channel. The initial handshake must occur in plaintext, yielding some important data points for inspection, profiling, and analysis. Of course, all communications after the connection is set up occur within the encrypted channel. Either party can request an SSL/TLS renegotiation at any time, at which point the initial process repeats. However, the renegotiation occurs over the existing secure channel and is therefore not available for casual inspection.

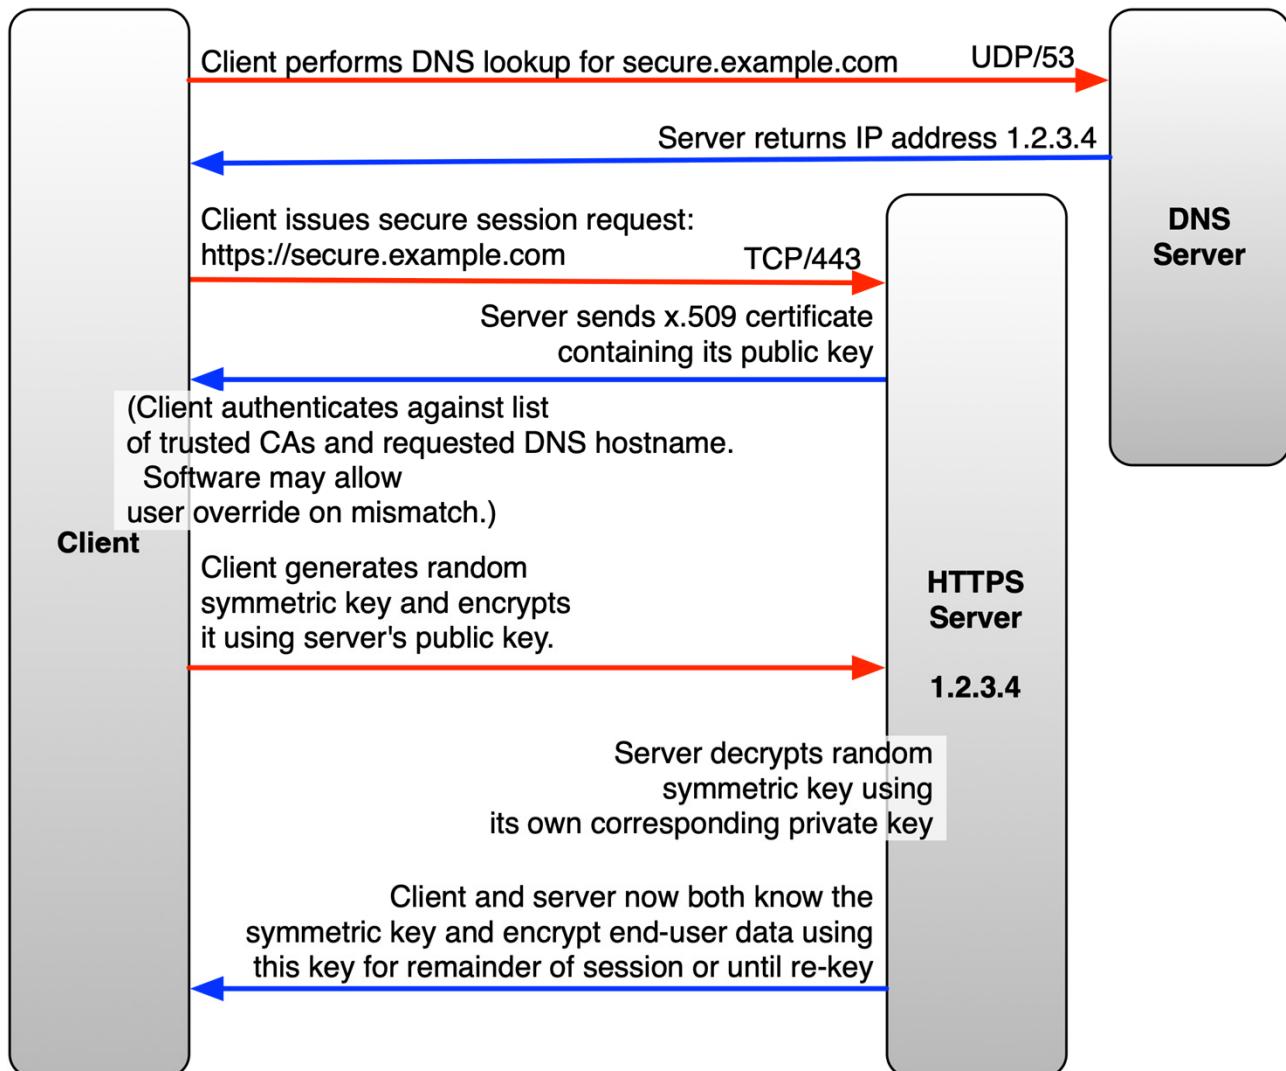
During the typical exchange depicted in the diagram, the client and server exchange a number of useful fields in the plaintext portion of the communication. The initial client request is called the “Client Hello” message. The Client Hello packet includes a list of encryption algorithms (called cipher suites) and compression mechanisms that the client is willing to accept during the handshake. Often, these values can indicate the client operating system or SSL/TLS libraries, because certain platforms offer a known set of cipher specs during this phase.

The server then responds with a “Server Hello” message and selects one of the client’s supported cipher specs and compression mechanisms. By design, the server selects the most highly preferred cipher suite from the client list against its own preferred list of suites, resulting in the most preferred mutually usable cipher suite between the two parties. The server also sends its own TLS certificate. This certificate contains several important data points: The “common name”, or DNS hostname, the certificate claims to certify is an important piece of information that can be used in conjunction with knowledge of the DNS request to ensure a match or isolate the appropriate NetFlow records.

The certificate may also include several parent certificates that have signed the server’s own certificate. The client uses these values to validate the server’s offered certificate against its own core list of “trusted” Certificate Authorities. This evaluation process also involves the client’s confirmation that the common name being certified matches the hostname that it originally requested via DNS. Generally, if the client cannot validate the certificate at

this stage, it displays the familiar warning messages asking whether the user wants to accept the untrusted certificate. (Of course, users may very well just click “OK” to proceed anyway...)

Whether automatically or through the user’s “manual override”, after the certificate is validated, the client continues the negotiation using a public key encryption process. With this process completed, the connection becomes fully encrypted after a few more packets and the connection is considered opaque to the analyst. Once the connection is encrypted, there is little more we can do to examine the contents of the communication. Even subsequent rekeying requests, which either side can request arbitrarily, are accomplished within the encrypted channel.



Perfect Forward Secrecy (PFS)

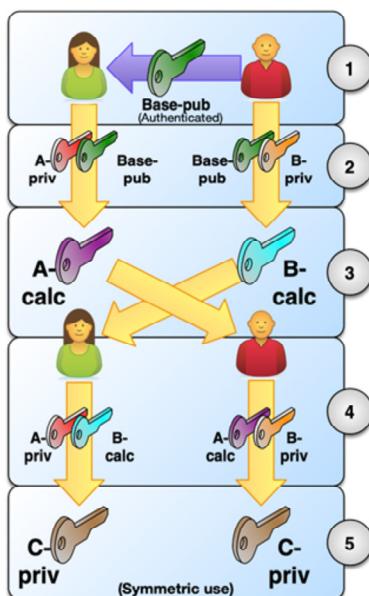
- Any private key compromise may permit later decryption of captured messages
- Public-key system provides PFS by:
 - Generating a random public key per session
 - Keys generated with nondeterministic algorithm
- Compromise of one PFS message cannot lead to the compromise of others
 - Encryption key is truly ephemeral
- Diffie-Hellman key exchange provides PFS

As explained previously, one risk associated with public key cryptography is that any compromise of the private key material will result in potential loss of all messages it previously encrypted. Given the undisputed trend toward use of more resilient encryption on the modern Internet, cryptographers have implemented a mitigation to this weakness called Perfect Forward Secrecy. PFS is a means by which even if the key material from both parties were later acquired, their communications would still remain protected and could not be decrypted.

In a PFS connection, each party generates random public key material to be used for key agreement—without relying on a deterministic algorithm to do so. This means that if the key for one message is ever compromised, it will not cause the compromise of any others. Also, there is no single secret value (such as the system's private key) that would lead to the compromise of past or future messages. This ultimately means that the session keys the server generates are truly ephemeral.

In other terms, parties to a PFS conversation use algorithms that establish a common private key without that key ever being passed over the wire. Perhaps the most commonly cited implementation of PFS is the Diffie-Hellman key exchange.

Diffie-Hellman Key Exchange



1. Share common, public base key
 - (Authenticated with server's pubkey)
2. Apply separate, ephemeral secret keys to the common public key
3. Exchange calculated values
4. Use own ephemeral secret key with other's calculated values
5. Independently derive common, ephemeral secret key

This diagram shows the steps involved in a Diffie-Hellman Key Exchange. As you can see, the “C-priv” keys have been mutually calculated, without any private key material being passed across the untrusted public medium.

The steps of this key exchange, which provides perfect forward secrecy for the derived private key, is as detailed below:

1. Alice and Bob share on a common, public base key. (Typically, this is generated by the server and sent to the requesting client.)
2. Alice and Bob apply separate, ephemerally generated secret keys to the common public key, which results in two new calculated public values.
3. Alice and Bob publicly exchange the calculated values.
4. Alice and Bob use the same ephemeral secret keys against the exchanged calculated values to derive a common secret value without ever transmitting secret key material over the untrusted medium in either encrypted or unencrypted form.
5. Alice and Bob can use the independently calculated shared secret as needed—for example, as a symmetric encryption key.

Obviously, a PFS-protected communication will be a big hurdle to a network forensicator trying to investigate the connection, because even a later acquisition of the private key will not allow decryption. The only means of debugging these connections involves configuring one endpoint or the other to log the derived private key material to a file for later decryption processing.

For those who prefer a more mathematical approach to this process, the below steps correspond to the diagram with the equations used in each.

1. Agree on known g, p
2. Pick hidden a, b
3. $A = g^a \text{ mod } p$
 $B = g^b \text{ mod } p$
4. Exchange A, B values
5. $C = A^b \text{ mod } p$
 $C = B^a \text{ mod } p$

To give a little more detail on this, the initial "key" that Alice and Bob agree on is actually two numbers. One is a large prime " p " and the other is a value " g ", which is a "primitive root modulo p " (also sometimes called a "generator"). The value " g " has the property that for any value V between 0 and p , there is some exponent x such that ($V = g^x \text{ mod } p$). Generally, g and p are well-known values that are used over and over again for a given PFS implementation.

Having agreed on g and p , Alice and Bob now each choose their own secret a and b values. Alice calculates a new value ($A = g^a \text{ mod } p$), and Bob calculates ($B = g^b \text{ mod } p$). They then exchange their A and B values with each other publicly.

Bob can now calculate ($C = A^b \text{ mod } p$) and Alice calculates ($C = B^a \text{ mod } p$). Because of the mathematical relationship between g and p , both Bob and Alice end up with the same C value. Only somebody who knows a or b could calculate this value.

However, the only "secret" components in this communication are the chosen a and b values. Figuring out a and b requires solving the "discrete logarithm problem" for which there is no known efficient algorithm ("computationally intractable" as the computer scientists like to say). However, an organization with heavy computing resources could precalculate much of the work to solve the discrete logarithm problem for a given pair of g and p . Once A and B are known, this precalculation effort would allow solving for a and b in perhaps a minute with modern computing power—fast enough to eavesdrop on the communication in near real time.

Many thanks to Hal Pomeranz for taking the time to assemble this technical explanation!

References:

<https://for572.com/chapw>
<https://for572.com/jg21r>

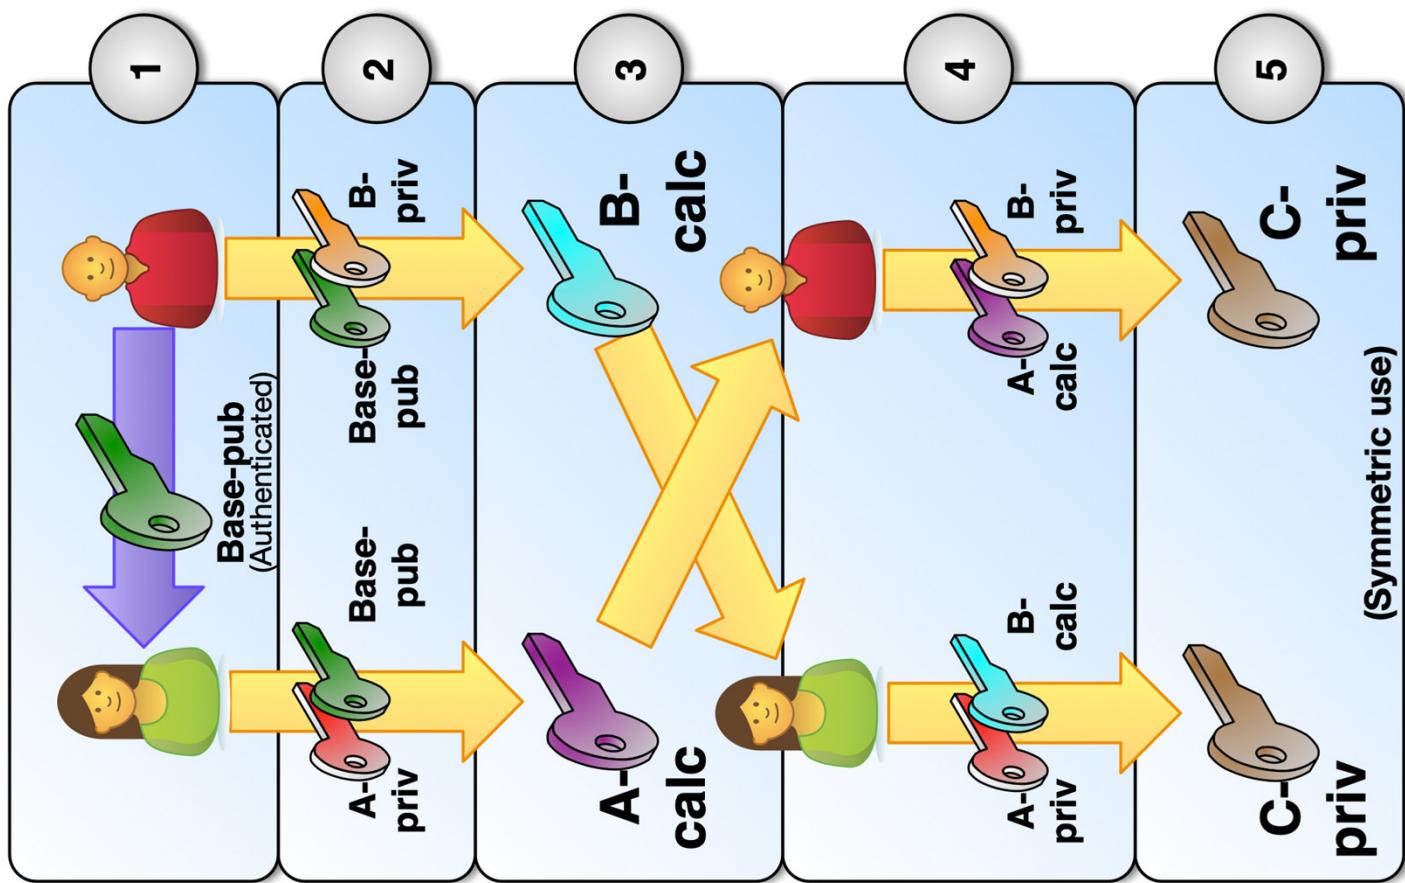
Step 1: Alice and Bob share a common, public base key

Step 2: Alice and Bob apply separate, ephemerally-generated secret keys to the common public key, which results in two new calculated public values

Step 3: Alice and Bob publicly exchange the calculated values

Step 4: Alice and Bob use the same ephemeral secret keys against the exchanged calculated values to derive a common secret value without ever transmitting secret key material over the untrusted medium in either encrypted or unencrypted form

Step 5: Alice and Bob can use the independently-calculated shared secret as needed - for example as a symmetric encryption key



HTTPS, in a Nutshell

- Just a TLS “wrapper” over plain old HTTP
- Some unique characteristics to consider



We previously spent a good deal of time on HTTP—the good news is that for the most part, all of what we just covered also applies to encrypted HTTP. HTTPS is basically the same HTTP you know and love—just wrapped into an encrypted connection. However, by virtue of the encryption, it can't be easily examined without special tools and techniques. Despite that potentially work-stopping hurdle, there is still a good deal to learn from just external observation of an HTTPS connection.

In this module, we'll primarily focus on the differences in how HTTP behaves between typical unencrypted mode and encrypted HTTPS. We'll also briefly touch on some of the more useful artifacts that can still be captured from HTTPS communications—even if we can't observe the content itself.

TLS Handshake Details (1a)



```

TLSv3.2 Record Layer - Handshake Protocol: Client Hello
Content Type: Handshake (22)
Version: TLS 1.0 (0x0301)
Length: 512
- Handshake Protocol: Client Hello
- Handshake Type: Client Hello (1)
Length: 512
Version: TLS 1.2 (0x0303)
Session ID Length: 32
Session ID: 0x848c0d3591777f8fad4e01f30f0504bf0fa12568f68355af090a770ff2
Cipher Suites (27 suites)
Cipher Suite: Reserved (GREASE) (0x1a1a)
Cipher Suite: TLS_AES_128_GCM_SHA256 (0x1301)
Cipher Suite: TLS_AES_256_GCM_SHA384 (0x1302)
Cipher Suite: TLS_CHACHA20_POLY1305_SHA256 (0x1303)
Cipher Suite: TLS_ECDHE_ECDSA_WITH_AES_256_GCM_SHA384 (0xc02c)
Cipher Suite: TLS_ECDHE_ECDSA_WITH_AES_128_GCM_SHA256 (0xc02b)
Cipher Suite: TLS_ECDHE_ECDSA_WITH_CHACHA20_POLY1305_SHA256 (0xccaa9)
Cipher Suite: TLS_ECDHE_RSA_WITH_AES_256_GCM_SHA384 (0xc030)
Cipher Suite: TLS_ECDHE_RSA_WITH_AES_128_GCM_SHA256 (0xc02f)
Cipher Suite: TLS_ECDHE_RSA_WITH_CHACHA20_POLY1305_SHA256 (0xccaa8)
Cipher Suite: TLS_ECDHE_ECDSA_WITH_AES_256_CBC_SHA384 (0xc024)
Cipher Suite: TLS_ECDHE_ECDSA_WITH_AES_128_CBC_SHA256 (0xc023)
Cipher Suite: TLS_ECDHE_ECDSA_WITH_AES_256_CBC_SHA (0xc025)
Cipher Suite: TLS_ECDHE_ECDSA_WITH_AES_128_CBC_SHA (0xc024)
Cipher Suite: TLS_ECDHE_RSA_WITH_CHACHA20_POLY1305_SHA256 (0xc025)
Cipher Suite: TLS_ECDHE_RSA_WITH_AES_256_CBC_SHA384 (0xc026)
Cipher Suite: TLS_ECDHE_RSA_WITH_AES_128_CBC_SHA256 (0xc025)
Cipher Suite: TLS_ECDHE_RSA_WITH_AES_256_CBC_SHA (0xc024)
Cipher Suite: TLS_ECDHE_ECDSA_WITH_AES_256_CBC_SHA384 (0xc027)
Cipher Suite: TLS_ECDHE_ECDSA_WITH_AES_128_CBC_SHA256 (0xc026)
Cipher Suite: TLS_ECDHE_ECDSA_WITH_AES_256_CBC_SHA (0xc025)
Cipher Suite: TLS_ECDHE_RSA_WITH_AES_256_CBC_SHA384 (0xc028)
Cipher Suite: TLS_ECDHE_RSA_WITH_AES_128_CBC_SHA256 (0xc027)
Cipher Suite: TLS_ECDHE_RSA_WITH_AES_256_CBC_SHA (0xc026)
Cipher Suite: TLS_ECDHE_RSA_WITH_AES_128_CBC_SHA (0xc025)
Cipher Suite: TLS_RSA_WITH_AES_256_GCM_SHA384 (0x009d)
Cipher Suite: TLS_RSA_WITH_AES_128_GCM_SHA256 (0x009c)
Cipher Suite: TLS_RSA_WITH_AES_256_CBC_SHA256 (0x003d)
Cipher Suite: TLS_RSA_WITH_AES_128_CBC_SHA256 (0x003c)
Cipher Suite: TLS_RSA_WITH_AES_256_CBC_SHA (0x0035)
Cipher Suite: TLS_RSA_WITH_AES_128_CBC_SHA (0x002f)
Cipher Suite: TLS_ECDHE_ECDSA_WITH_3DES_EDE_CBC_SHA (0xc008)
Cipher Suite: TLS_ECDHE_RSA_WITH_3DES_EDE_CBC_SHA (0xc012)
Cipher Suite: TLS_RSA_WITH_3DES_EDE_CBC_SHA (0x000a)
[Community ID: 1.SS0ly90WQ2P1L40VVK8N7B0ZL]
[Compression Method: 0x01]
[Priority: 0x01]
[Extensions: 14 extensions]
Extension: reserved (GREASE) [len=0]
Extension: session_id [len=0]
Extension: extended_master_secret [len=0]
Extension: renegotiation_info [len=1]
Extension: renegotiation_info [len=1]
Extension: ec_point_formats [len=2]
Extension: application_layer_protocol_negotiation [len=14]
Extension: status_request [len=5]
Extension: supported_versions [len=24]
Extension: signed_certificate_timestamp [len=0]
Extension: key_share [len=43]
Extension: supported_groups [len=2]
Extension: supported_versions [len=12]
Extension: reserved (GREASE) [len=1]
[Community ID: 1.SS0ly90WQ2P1L40VVK8N7B0ZL]
```

```

TLSv3.2 Record Layer - Handshake Protocol: Client Hello
Content Type: Handshake (22)
Version: TLS 1.0 (0x0301)
Length: 512
- Handshake Protocol: Client Hello
- Handshake Type: Client Hello (1)
Length: 512
Version: TLS 1.2 (0x0303)
Session ID Length: 32
Session ID: 0x848c0d3591777f8fad4e01f30f0504bf0fa12568f68355af090a770ff2
Cipher Suites (27 suites)
Cipher Suite: Reserved (GREASE) (0x1a1a)
Cipher Suite: TLS_AES_128_GCM_SHA256 (0x1301)
Cipher Suite: TLS_AES_256_GCM_SHA384 (0x1302)
Cipher Suite: TLS_CHACHA20_POLY1305_SHA256 (0x1303)
Cipher Suite: TLS_ECDHE_ECDSA_WITH_AES_256_GCM_SHA384 (0xc02c)
Cipher Suite: TLS_ECDHE_ECDSA_WITH_AES_128_GCM_SHA256 (0xc02b)
Cipher Suite: TLS_ECDHE_ECDSA_WITH_CHACHA20_POLY1305_SHA256 (0xccaa9)
Cipher Suite: TLS_ECDHE_RSA_WITH_AES_256_GCM_SHA384 (0xc030)
Cipher Suite: TLS_ECDHE_RSA_WITH_AES_128_GCM_SHA256 (0xc02f)
Cipher Suite: TLS_ECDHE_RSA_WITH_CHACHA20_POLY1305_SHA256 (0xccaa8)
Cipher Suite: TLS_ECDHE_ECDSA_WITH_AES_256_CBC_SHA384 (0xc024)
Cipher Suite: TLS_ECDHE_ECDSA_WITH_AES_128_CBC_SHA256 (0xc023)
Cipher Suite: TLS_ECDHE_ECDSA_WITH_AES_256_CBC_SHA (0xc025)
Cipher Suite: TLS_ECDHE_RSA_WITH_AES_256_CBC_SHA384 (0xc026)
Cipher Suite: TLS_ECDHE_RSA_WITH_AES_128_CBC_SHA256 (0xc025)
Cipher Suite: TLS_ECDHE_RSA_WITH_AES_256_CBC_SHA (0xc024)
Cipher Suite: TLS_ECDHE_ECDSA_WITH_AES_256_CBC_SHA384 (0xc027)
Cipher Suite: TLS_ECDHE_ECDSA_WITH_AES_128_CBC_SHA256 (0xc026)
Cipher Suite: TLS_ECDHE_ECDSA_WITH_AES_256_CBC_SHA (0xc025)
Cipher Suite: TLS_ECDHE_RSA_WITH_AES_256_CBC_SHA384 (0xc028)
Cipher Suite: TLS_ECDHE_RSA_WITH_AES_128_CBC_SHA256 (0xc027)
Cipher Suite: TLS_ECDHE_RSA_WITH_AES_256_CBC_SHA (0xc026)
Cipher Suite: TLS_ECDHE_RSA_WITH_AES_128_CBC_SHA (0xc025)
Cipher Suite: TLS_RSA_WITH_AES_256_GCM_SHA384 (0x009d)
Cipher Suite: TLS_RSA_WITH_AES_128_GCM_SHA256 (0x009c)
Cipher Suite: TLS_RSA_WITH_AES_256_CBC_SHA256 (0x003d)
Cipher Suite: TLS_RSA_WITH_AES_128_CBC_SHA256 (0x003c)
Cipher Suite: TLS_RSA_WITH_AES_256_CBC_SHA (0x0035)
Cipher Suite: TLS_RSA_WITH_AES_128_CBC_SHA (0x002f)
Cipher Suite: TLS_ECDHE_ECDSA_WITH_3DES_EDE_CBC_SHA (0xc008)
Cipher Suite: TLS_ECDHE_RSA_WITH_3DES_EDE_CBC_SHA (0xc012)
Cipher Suite: TLS_RSA_WITH_3DES_EDE_CBC_SHA (0x000a)
[Community ID: 1.SS0ly90WQ2P1L40VVK8N7B0ZL]
[Compression Method: 0x01]
[Priority: 0x01]
[Extensions: 14 extensions]
Extension: reserved (GREASE) [len=0]
Extension: session_id [len=0]
Extension: extended_master_secret [len=0]
Extension: renegotiation_info [len=1]
Extension: renegotiation_info [len=1]
Extension: ec_point_formats [len=2]
Extension: application_layer_protocol_negotiation [len=14]
Extension: status_request [len=5]
Extension: supported_versions [len=24]
Extension: signed_certificate_timestamp [len=0]
Extension: key_share [len=43]
Extension: supported_groups [len=2]
Extension: supported_versions [len=12]
Extension: reserved (GREASE) [len=1]
[Community ID: 1.SS0ly90WQ2P1L40VVK8N7B0ZL]
```

Now we'll look at portions of the first few packets exchanged as they appear in Wireshark's packet details pane.

This screen shows the Client Hello message. This contains the parameters the client is willing to accept for the remainder of the session. This includes cipher suites, compression algorithms, and (optionally) any number of TLS extensions. The cipher suites and compression algorithms are listed in order of the client's priority for acceptance.

Most important to the forensicator is the list of cipher suites, as this can be a key indicator of how recent the client TLS library is. In this case, the request was made with Safari Version 14.0.3 (16610.4.3.1.7). This represented the most recent release at the time the screenshot was taken, and the list of ciphers represents all current and non-deprecated protocols at the top. Lower in the list are several older ciphers, included to ensure backward compatibility with servers using older TLS libraries.

The meaning of any individual cipher suite is not generally of significant consequence for investigative purposes, but the SSL/TLS RFCs^{[1][2]} specify the meaning of each if necessary.

References:

- [1] <https://for572.com/jz0dt>
- [2] <https://for572.com/3jusg>

TLS Handshake Details (1b)



- Requested SSL/TLS version + cipher suites + SSL/TLS extensions = Client profile
- “JA3 Hash” combines these into easy-to-use value
 - Standalone; Integrated to Arkime, Zeek, NetworkMiner

```
$ cd /cases/for572/sample_pcaps/
$ ja3.py -j tls_dfir.com_session.pcap | jq -cr '.[]'
{"destination_ip":"70.32.97.206","destination_port":443,"ja3":"771,4865-4867-4866-49195-49199-52393-52392-49196-49200-49162-49161-49171-49172-51-57-47-53-10,0-23-65281-10-11-35-16-5-51-43-13-45-28,29-23-24-25-256-257,0","ja3_digest":"334da95730484a993c6063e36bc90a47","source_ip":"192.168.211.134","source_port":48584,"timestamp":1544121582.707847}
...
$ grep 334da95730484a993c6063e36bc90a47 /usr/local/for572/lib/ja3fingerprint_ua.json | jq -cr '.useragent'
Mozilla/5.0 (Android 10; Mobile; rv:68.0) Gecko/68.0 Firefox/68.0
...
$ grep b50f81ae37fb467713e167137cf14540 /usr/local/for572/lib/ja3fingerprint_ua.json | jq -cr '.useragent'
Malware: TBot / Skynet Tor Botnet
```

The distinct profile a client-side SSL/TLS library provides has always been a useful means of fingerprinting common encrypted traffic. The security team at Salesforce commoditized this profiling opportunity and released an open-source utility named “JA3”^[1] (pronounced JA-three). The JA3 process calculates a simple hash value of the requested SSL/TLS version, list and order of requested Cipher Suites, list and order of supported elliptic curve values, and any available SSL/TLS extensions.

After computing the JA3 hash, an analyst can build profiles of normal activity in their environments or query databases of known hash values to attempt client software identification—specifically the SSL/TLS library. This is especially useful, as malware has evolved to use common platforms for C2 purposes. For example, one malware variant was documented to be using comments on popular Instagram accounts^{[2][3]}, which would make it extremely difficult to separate from normal, benign traffic to the same site/service through traditional SSL/TLS analytic mitigations such as NetFlow and DNS profiling.

The fingerprint itself is useful for forensic purposes, as demonstrated in the slide above. However, the value of this approach increases significantly when integrated to network security monitoring platforms such as Arkime and Zeek. In this model, the JA3 hashes are calculated and stored at the time of observation, allowing much longer-term use.

References:

- [1] <https://for572.com/5uoeb>
- [2] <https://for572.com/epr2x>
- [3] <https://for572.com/v7bl8>

TLS Handshake Details (2a)



```
- TLSv1.2 Record Layer: Handshake Protocol: Server Hello
  Content Type: Handshake (22)
  Version: TLS 1.2 (0x0303)
  Length: 97
- Handshake Protocol: Server Hello
  Handshake Type: Server Hello (2)
  Length: 92
  Version: TLS 1.2 (0x0303)
  Random: uefa56aebca3e4e9ed7a812c7f0abdda3d220585bfcf4f5ace797879ec34de3f
  Session ID Length: 32
  Session ID: 6b000d8bf12e7e2b16d258c100e750bf14abc4285f20680255270795d30167d0
  Cipher Suite: TLS_ECDHE_RSA_WITH_AES_128_GCM_SHA256 (0xc02f)
  Compression Method: null (0)
  Extensions Length: 21
  Extension: server_name (len=0)
  Extension: renegotiation_info (len=1)
  Extension: ec_point_formats (len=4)
  Extension: status_request (len=0)
[Community ID: 1:51D1y80HNhqrLtg4WYVX5NT0ZCE=]
```

Here we see the server's response to the Client Hello shown previously. This is similarly named—it is the Server Hello message. As we have highlighted, the server chose one of the client's supported ciphers, selecting the most preferable on its own list that the client offered in the previous message. This was the second priority in the Client Hello message, which indicates that the server cannot (or will not) negotiate the client's most preferred cipher suite. While not a definitive means of profiling the server, examining how a server responds to clients' negotiation requests can be used to suggest the server's SSL/TLS libraries similarly to the client side.

TLS Handshake Details (2b)

- Chosen SSL/TLS version + cipher suite + extensions = Server profile
- “JA3S Hash” combines and reports just like JA3
 - Fewer known signatures, but growing in popularity

```
$ cd /cases/for572/sample_pcaps/
$ ja3s.py -j tls dfir.com session.pcap | jq -c '.[]'
{"destination_ip": "192.168.211.134", "destination_port": 48514, "ja3": "771,49199,0-65281-11-35-473", "ja3_digest": "a9f50de6377ae7c6dbaabcb53f089ac30", "source_ip": "70.32.97.206", "source_port": 443, "timestamp": 1544121579.456
473}
...
```

The same team that refined the JA3 hash also created a similar means of fingerprinting the server side of an SSL/TLS negotiation. This JA3S (JA-three-ess) hash uses the server’s chosen parameters for fingerprint generation. As with the client side, a growing number of databases for known server configurations are being created, but coverage is notably less extensive than is available for client libraries.

Even without a public fingerprint, the ability to profile typical server behaviors within an environment can be a valuable means of baselining, and therefore, anomaly detection.

```

▼ TLSv1.3 Record Layer: Handshake Protocol: Client Hello
  Content Type: Handshake (22)
  Version: TLS 1.0 (0x0301)
  Length: 512
  ▼ Handshake Protocol: Client Hello
    Handshake Type: Client Hello (1)
    Length: 300
    Version: TLS 1.2 (0x0303)
    Random: 0014000000000000000000000000000000000000000000000000000000000000
    Session ID Length: 32
    Session ID: 19d4f973e066b213573d47125eda1b6011dff08f851580ce...
    Cipher Suites Length: 52
    ▼ Cipher Suites (26 suites)
      Cipher Suite: TLS_AES_128_GCM_SHA256 (0x1301)
      Cipher Suite: TLS_AES_256_GCM_SHA384 (0x1302)
      Cipher Suite: TLS_CHACHA20_POLY1305_SHA256 (0x1303)
      Cipher Suite: TLS_ECDHE_ECDSA_WITH_AES_256_GCM_SHA384 (0xc02c)
      Cipher Suite: TLS_ECDHE_ECDSA_WITH_AES_128_GCM_SHA256 (0xc02b)
      Cipher Suite: TLS_ECDHE_ECDSA_WITH_AES_256_CBC_SHA384 (0xc024)
      Cipher Suite: TLS_ECDHE_ECDSA_WITH_AES_128_CBC_SHA256 (0xc023)
      Cipher Suite: TLS_ECDHE_ECDSA_WITH_AES_256_CBC_SHA (0xc00a)
      Cipher Suite: TLS_ECDHE_ECDSA_WITH_AES_128_CBC_SHA (0xc009)
      Cipher Suite: TLS_ECDHE_ECDSA_WITH_CHACHA20_POLY1305_SHA256 (0xccaa9)
      Cipher Suite: TLS_ECDHE_RSA_WITH_AES_256_GCM_SHA384 (0xc030)
      Cipher Suite: TLS_ECDHE_RSA_WITH_AES_128_GCM_SHA256 (0xc02f)
      Cipher Suite: TLS_ECDHE_RSA_WITH_AES_256_CBC_SHA384 (0xc028)
      Cipher Suite: TLS_ECDHE_RSA_WITH_AES_128_CBC_SHA256 (0xc027)
      Cipher Suite: TLS_ECDHE_RSA_WITH_AES_256_CBC_SHA (0xc014)
      Cipher Suite: TLS_ECDHE_RSA_WITH_AES_128_CBC_SHA (0xc013)
      Cipher Suite: TLS_ECDHE_RSA_WITH_CHACHA20_POLY1305_SHA256 (0xccaa8)
      Cipher Suite: TLS_RSA_WITH_AES_256_GCM_SHA384 (0x009d)
      Cipher Suite: TLS_RSA_WITH_AES_128_GCM_SHA256 (0x009c)
      Cipher Suite: TLS_RSA_WITH_AES_256_CBC_SHA256 (0x003d)
      Cipher Suite: TLS_RSA_WITH_AES_128_CBC_SHA256 (0x003c)
      Cipher Suite: TLS_RSA_WITH_AES_256_CBC_SHA (0x0035)
      Cipher Suite: TLS_RSA_WITH_AES_128_CBC_SHA (0x002f)
      Cipher Suite: TLS_ECDHE_ECDSA_WITH_3DES_EDE_CBC_SHA (0xc008)
      Cipher Suite: TLS_ECDHE_RSA_WITH_3DES_EDE_CBC_SHA (0xc012)
      Cipher Suite: TLS_RSA_WITH_3DES_EDE_CBC_SHA (0x00aa)
    Compression Methods Length: 1
    Compression Method: (1)
    ▼ Compression Method (1)
      Compression Method: null (0)
      Random: 0014000000000000000000000000000000000000000000000000000000000000
      ▼ Extension (194)
        Extension: renegotiation_info (len=1)
        Extension: server_name (len=32)
        Extension: extended_master_secret (len=0)
        Extension: signature_algorithms (len=24)
        Extension: status_request (len=5)
        Extension: signed_certificate_timestamp (len=0)
        Extension: application_layer_protocol_negotiation (len=14)
        Extension: ec_point_formats (len=2)
        Extension: key_share (len=38)
        Extension: psk_key_exchange_modes (len=2)
        Extension: supported_versions (len=9)
        Extension: supported_groups (len=10)
        Extension: padding (len=194)

```

```

▼ TLSv1.3 Record Layer: Handshake Protocol: Server Hello
  Content Type: Handshake (22)
  Version: TLS 1.2 (0x0303)
  Length: 122
  ▼ Handshake Protocol: Server Hello
    Handshake Type: Server Hello (2)
    Length: 110
    Version: TLS 1.2 (0x0303)
    Random: 0014000000000000000000000000000000000000000000000000000000000000
    Session ID Length: 32
    Session ID: 5b2253e858253a...
    Cipher Suite: TLS_AES_256_GCM_SHA384 (0x1302)
    Compression Method: null (0)
    Extensions Length: 46
    ▶ Extension: supported_versions (len=2)
    ▶ Extension: key_share (len=36)
  ▶ TLSv1.3 Record Layer: Change Cipher Spec
  ▶ TLSv1.3 Record Layer: Application Data Protocol: http-over-tls

```



TLS Handshake Details (3a)

```

- TLSv1.2 Record Layer: Handshake Protocol: Certificate
  Content Type: Handshake (22)
  version: TLS 1.2 (0x0303)
  Length: 2475
- Handshake Protocol: Certificate
  Handshake Type: Certificate (11)
  Length: 2471
  Certificates Length: 2468
- Certificates (2468 bytes)
  Certificate (2460 bytes)
    Certificate: 0002053130020419a0030201020212035421d15e0db9d4e425c18261e097f6509a300d00... (id-at-commonName=sansgear.com)
    - signedCertificate
      - version: v3 (2)
        serialNumber: 0x035421d15e0db9d4e425c18261e097f6509a
      - signature: sha256WithRSAEncryption
      - issuer: rdnSequence (0)
    - validity
      - notBefore: utcTime (0)
        utcTime: 2021-01-23 08:54:28 (UTC)
      - notAfter: utcTime (0)
        utcTime: 2021-04-23 08:54:28 (UTC)
    - subject: rdnSequence (0)
      - rdnSequence: 1 item (id-at-commonName=sansgear.com)
        RDNSequence item: 1 item (id-at-commonName=sansgear.com)
          - RelativeDistinguishedName item (id-at-commonName=sansgear.com)
            Id: 2.5.4.3 (id-at-commonName)
            - DirectoryString: printableString (1)
              printableString: sansgear.com
      - subjectPublicKeyInfo
      - extensions: 9 items
        Extension (id-ce-keyUsage)
        Extension (id-ce-extKeyUsage)
        Extension (id-ce-basicConstraints)
        Extension (id-ce-subjectKeyIdentifier)
        Extension (id-ce-authorityKeyIdentifier)
        Extension (id-ce-subAltName)
        Extension Id: 2.5.29.17 (id-ce-subjectAltName)
        GeneralNames: 2 items
          - GeneralName: dNSName (2)
            dNSName: sansgear.com
          - GeneralName: dNSName (2)
            dNSName: www.sansgear.com
      - Extension (id-ce-certificatePolicies)
      - Extension (SignedCertificateTimestampList)
    algorithmIdentifier (sha256WithRSAEncryption)
    Padding: 0
  encrypted: 1aac56f21b609c507de3bed7cad41dd75ef74487404ff2d6071b386bc43ca2b93b3da826...

```

Finally, we see the certificates the server sends on this and the following slide. Above is the server's own certificate – confirmed by observing the “`id-at-commonName`” value of “`sansgear.com`” in the packet details hierarchy. This is sometimes identified as the “subject name”, “SN”, or simply “subject” of the certificate.

There are several useful artifacts in the certificate. Aside from the subject name, the serial number provides a unique way of identifying the certificate. Each certificate issuer (or “Certificate Authority”) can issue exactly one certificate with a given serial number. Therefore, this can be a method for unmasking fast flux DNS functionality, or an attacker’s infrastructure that reuses key material across multiple servers or server instances.

Lastly, note the “validity” value. These establish the time frame in which the certificate is determined valid. In this case, the 90-day window suggests the certificate may be issued through the Let’s Encrypt! service, which mandates a 90-day value for all of their certificates. However, the “`issuer`” field above the validity window is a more definitive means of identifying the Certificate Authority.

TLS Handshake Details (3b)



```
▼ TLSv1.2 Record Layer: Handshake Protocol: Certificate
  Content Type: Handshake (22)
  Version: TLS 1.2 (0x0303)
  Length: 2475
  ▼ Handshake Protocol: Certificate
    Handshake Type: Certificate (11)
    Length: 2471
    Certificates Length: 2468
    Certificates (2468 bytes)
      Certificate Length: 1333
      ▶ Certificate: 3082053130820419a0030201020212035421d15e0db9d4e425c18261e097f6509a300d06... (id-at-commonName=sansgear.com)
      Certificate Length: 1129
      ▶ Certificate: 308204653082034da0030201020210400175048314a4c8218c84a90c16cddf300d06092a... (id-at-commonName=R3, id-at-organizationName=Let's Encrypt)
        ▼ signedCertificate
          version: v3 (2)
          serialNumber: 0x400175048314a4c8218c84a90c16cddf
        ▶ signature (sha256WithRSAEncryption)
        ▶ issuer: rdnSequence (0)
        ▶ validity
          subject: rdnSequence (0)
            ▼ rdnSequence: 3 items (id-at-commonName=R3, id-at-organizationName=Let's Encrypt, id-at-countryName=US)
              ▶ RDNSequence item: 1 item (id-at-countryName=US)
              ▶ RDNSequence item: 1 item (id-at-organizationName=Let's Encrypt)
              ▶ RDNSequence item: 1 item (id-at-commonName=R3)
            ▶ subjectPublicKeyInfo
            ▶ extensions: 8 items
        ▶ algorithmIdentifier (sha256WithRSAEncryption)
        Padding: 0
        encrypted: d94ce0c9f584883731dbbb13e2b3fc8b6b62126c58b7497e3c02b7a81f2861ebcee02e73...
```

In addition to the certificate for the “sansgear.com” hostname, this message includes the full public certificate for its issuer, or Certificate Authority. As noted previously, the Lets Encrypt! service issued this certificate and its subject name is highlighted.

It is quite common to observe multiple upstream certificates sent in the same message. This is known as a “certificate chain”. The use of certificate chains is intended to maximize the chances that a hostname’s certificate is trusted by the client’s TLS configuration. Each client has a list of hardcoded trusted “root” certificates, and as long as there is a pathway from those root certificates to any signer in the chain, the client will validate the offered certificate for the hostname.

```

- TLSv1.2 Record Layer: Handshake Protocol: Certificate
  Content Type: Handshake (22)
  Version: TLS 1.2 (0x0303)
  Length: 2475
  - Handshake Protocol: Certificate
    Handshake Type: Certificate (11)
    Length: 2471
    Certificates Length: 2468
    - Certificates (2468 bytes)
      Certificate Length: 1333
      - Certificate: 3082053130820419a0030201020212035421d15e0db9d4e425c18261e097f6509a300d06... (id-at-commonName=sansgear.com)
        signedCertificate
          version: v3 (2)
          serialNumber: 0x035421d15e0db9d4e425c18261e097f6509a
          signature (sha256WithRSAEncryption)
          issuer: rdnSequence (0)
        validity
          notBefore: utcTime (0)
            utcTime: 2021-01-21 08:54:28 (UTC)
          notAfter: utcTime (0)
            utcTime: 2021-04-21 08:54:28 (UTC)
        subject: rdnSequence (0)
          rdnSequence: 1 item (id-at-commonName=sansgear.com)
            RDNSequence item: 1 item (id-at-commonName=sansgear.com)
              relativeDistinguishedName item (id-at-commonName=sansgear.com)
                Id: 2.5.4.3 (id-at-commonName)
                DirectoryString: printableString (1)
                  printableString: sansgear.com
        subjectPublicKeyInfo
        extensions: 9 items
          Extension (id-ce-keyUsage)
          Extension (id-ce-extKeyUsage)
          Extension (id-ce-basicConstraints)
          Extension (id-ce-subjectKeyIdentifier)
          Extension (id-ce-authorityKeyIdentifier)
          Extension (id-ce-authorityInfoAccess)
          Extension (id-ce-subjectAltName)
            Extension Id: 2.5.29.17 (id-ce-subjectAltName)
              GeneralNames: 2 items
                GeneralName: dnsName (2)
                  dnsName: sansgear.com
                GeneralName: dnsName (2)
                  dnsName: www.sansgear.com
          Extension (id-ce-certificatePolicies)
          Extension (SignedCertificateTimestampList)
        algorithmIdentifier (sha256WithRSAEncryption)
        Padding: 0
        encrypted: 1aac56f21b609c507de3bed7cad41dd75ef74487404ff2d6071b386bc43ca2b93b3da826...

```

Handshake Details (3a)

```

- TLSv1.2 Record Layer: Handshake Protocol: Certificate
  Content Type: Handshake (22)
  Version: TLS 1.2 (0x0303)
  Length: 2475
  - Handshake Protocol: Certificate
    Handshake Type: Certificate (11)
    Length: 2471
    Certificates Length: 2468
    - Certificates (2468 bytes)
      Certificate Length: 1333
      - Certificate: 3082053130820419a0030201020212035421d15e0db9d4e425c18261e097f6509a300d06... (id-at-commonName=sansgear.com)
      Certificate Length: 1129
      - Certificate: 308204653082034da0030201020210400175048314a4c8218c84a90c16cddf300d06092a... (id-at-commonName=R3,id-at-organizationName=Let's Encrypt)
        signedCertificate
          version: v3 (2)
          serialNumber: 0x0400175048314a4c8218c84a90c16cddf
          signature (sha256WithRSAEncryption)
          issuer: rdnSequence (0)
        validity
          subject: rdnSequence (0)
            rdnSequence: 3 items (id-at-commonName=R3,id-at-organizationName=Let's Encrypt,id-at-countryName=US)
              RDNSequence item: 1 item (id-at-countryName=US)
              RDNSequence item: 1 item (id-at-organizationName=Let's Encrypt)
              RDNSequence item: 1 item (id-at-commonName=R3)
        subjectPublicKeyInfo
        extensions: 8 items
        algorithmIdentifier (sha256WithRSAEncryption)
        Padding: 0
        encrypted: d94ce0c9f584883731dbbb13e2b3fc8b6b62126c58b7497e3c02b7a81f2861ebcee02e73...

```

Handshake Details (3b)

HTTPS “Gotcha”

- “Keep-Alive” connections allow multiple request/response pairs over a single socket
 - Provided by HTTP (not HTTPS-specific)
 - Complicates byte counts over session lifetime
- Consider a 15,516,267-byte HTTPS connection observed in NetFlow evidence
 - Single file, approximately 15.5 MB?
 - 100 web page loads, approximately 155 KB each?
 - Some other combination?
- HTTP/2 channels amplify this challenge

When analyzing an encrypted communications channel, one of the useful pieces of metadata is the amount of data transferred. This methodology might be useful in estimating the amount of data an attacker has stolen from a victim’s network. However, remember that the underlying protocol, in this case, HTTP, provides the “Keep-Alive” capability, which allows multiple request/response exchanges in a single socket. This feature was first devised to improve efficiency with regard to normal TCP handshaking—saving a few dozen bytes per connection. You can imagine that there is a significant improvement with TLS, which requires hundreds or thousands of bytes to establish the connection!

Similarly, HTTP/2’s introduction of up to 100 channels per request message almost entirely obscures the volume of individual request/response pairs from that of the overall session.

Although these new technologies certainly improve efficiency, they complicate volume-based analysis because there is no way to identify how many exchanges were conducted over a single TLS session, nor the amount of data for each exchange.

Differences from HTTP

- Multiple subject names per certificate: Subject Alternative Names (SAN)
 - Common for domain parking and CDNs
- Name-based virtual hosting requires SNI
 - Server Name Indication: Hostname in Client Hello
- Transparent proxying is much more complex
 - Requires explicit proxy configuration
 - Squid 3.5+ provides “Peek and Splice” features
 - Interception decision is made after Client Hello and Server Hello messages are exchanged

The nature of the encapsulating TLS protocol presents some other challenges that cause the enclosed HTTP traffic to diverge from its typical behavior when unencrypted. Understanding these differences can prove extremely useful in finding investigative value in the aspects of HTTPS traffic that we can still observe without decrypting the contents.

The Subject Alternative Name, or SAN^[1], field is another TLS variation that has become ubiquitous after Google Chrome’s insistence that the field be present in all certificates. (Note that this is no relation to Storage Area Networking.) The TLS SAN field is a list of one or more names for which the certificate can validate. This allows a site owner with multiple hostnames to use a single certificate and key to encrypt the transmissions. Domain parking providers and Content Delivery Networks (CDNs) often include numerous hostnames for which they are providing service. Occasionally, observing this field in conjunction with suspicious or malicious activity may also include insight to an attacker’s other hostnames.

Possibly the most pervasive difference is that “name-based virtual hosting” requires the presence of a field in the Client Hello message called “Server Name Indication” or “SNI”^[2]. This extension of the TLS standard, along with many others, is defined by RFC 6066. Over time, this field has become a nearly-universal standard for most browsers and their respective TLS libraries.

Regardless of how the TLS-related hostname is identified, it is often a useful lead in searching out similar behavior on other clients, mining DNS logs, and more.

There is a related difference between behavior of HTTP and HTTPS traffic when it comes to transparent proxying of web traffic. Transparent proxying takes place when a client’s web requests are diverted to a proxy server by means of a NATting firewall or similar routing device. This enforces proxying even for devices or applications that are not configured with or aware of the environment’s proxy settings.

When the client’s outgoing request is unencrypted (i.e., regular HTTP), the proxy can freely inspect the request to broker the exchange between the client and the intended server. The proxy has access to the client’s request headers and can therefore identify the hostname being requested from the “Host:” header. However, since a TLS-protected request is encrypted before the headers are sent, the traditional methods of request inspection are not possible.

Most proxy software can now dynamically generate TLS certificates for connections. Version 3.2 of the Squid proxy introduced the Certificate Daemon, which accomplishes the dynamic certificate generation internally. There are external solutions for Squid as well, and commercial proxy solutions such as BlueCoat also provide this functionality.

It should also be noted that to do this, the client systems must be configured to trust the certificates that the proxy dynamically generates. In any environment where the clients' system configuration can be controlled, this is trivial to accomplish.

There are currently two primary methods of proxying HTTPS traffic:

- Although not functional in a purely transparent model, a client that is given an explicit proxy setting for HTTPS requests uses the “CONNECT” method to request that the proxy retrieve a particular resource via HTTPS^[3]. The client then connects to the proxy to request the HTTPS resource. At this point, the proxy initiates the HTTPS connection with the intended destination and brokers the remainder of the connection. In this model, the proxy is aware of the hostname and port of the client's requests, but not the full URL, query string, headers, or content! The proxy acts as a tunneling NAT device in this configuration. However, when operating in this mode, the proxy can optionally use dynamic certificate generation to perform full content inspection.
- The proxy can be configured to leverage a series of features to inspect the first few TLS negotiation messages and then make a decision of whether to intercept, pass along, or block the rest of the exchange. In Squid version 3.5, this feature was introduced and became known as the “Peek and Splice” operations^[4]. While a full discussion on the use and pitfalls of these features is beyond the scope of this course, the end result is that transparently proxying TLS-protected traffic is possible. The proxy uses the SNI field, as well as other fields in the Client Hello message and various fields from the Server Hello message to provide protocol-compliant but forged and intercepted TLS connections to the client. This permits logging, caching, allow/block decisions, and virtually all of the same benefits a proxy provides for plaintext HTTP traffic to HTTPS traffic as well.

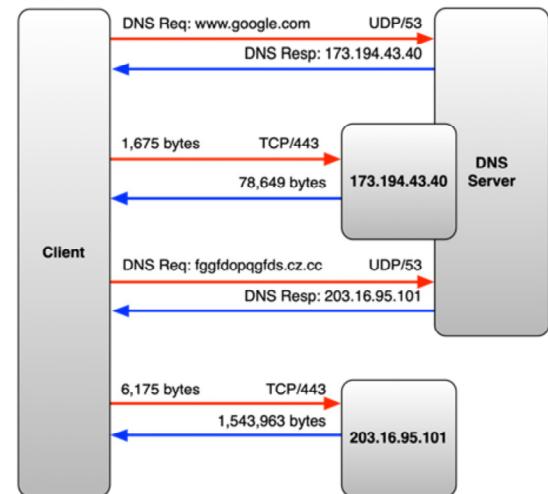
Does this sound a little bit shady? It should—this is commonly called a meddler-in-the-middle, or MITM, operation. As such, it is important to note that intercepting any TLS traffic may be illegal in your jurisdiction, company, or network environment. Consult with legal counsel before experimenting with a solution like this!

References:

- [1] <https://for572.com/-o0od>
- [2] <https://for572.com/vidf0>
- [3] <https://for572.com/8e2xh>
- [4] <https://for572.com/ftyen>

Analytic Mitigation: Supplement NetFlow with DNS Correlation

- Hostname needed to validate certificate, also used during TLS negotiation
- DNS traffic/logs, Client Hello SNI, certificate subject/SAN fields can all provide insight
 - Don't limit your view of hostname evidence to only DNS traffic!
 - Zeek and other NSMs often log these values

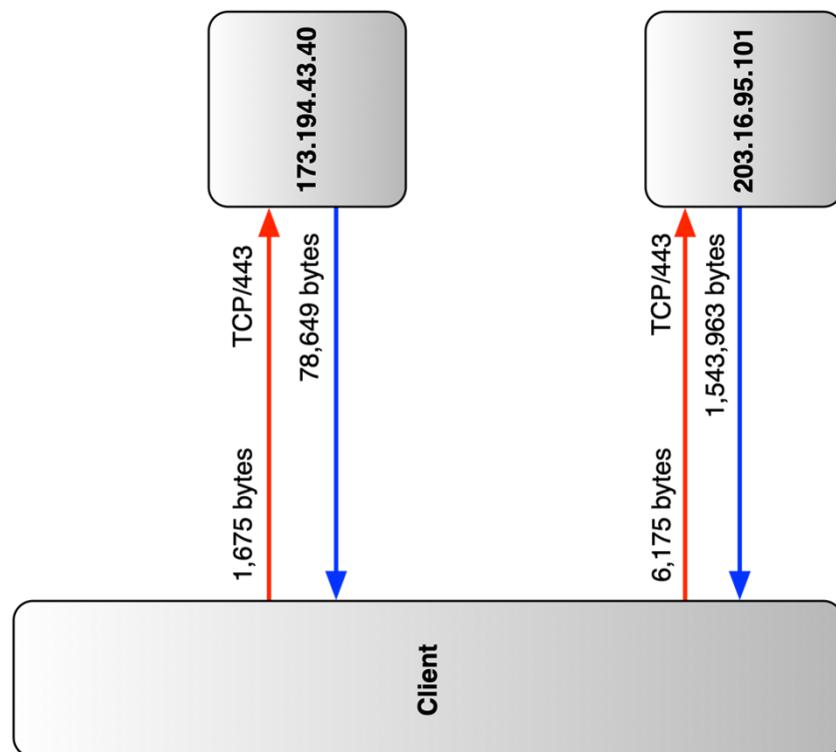
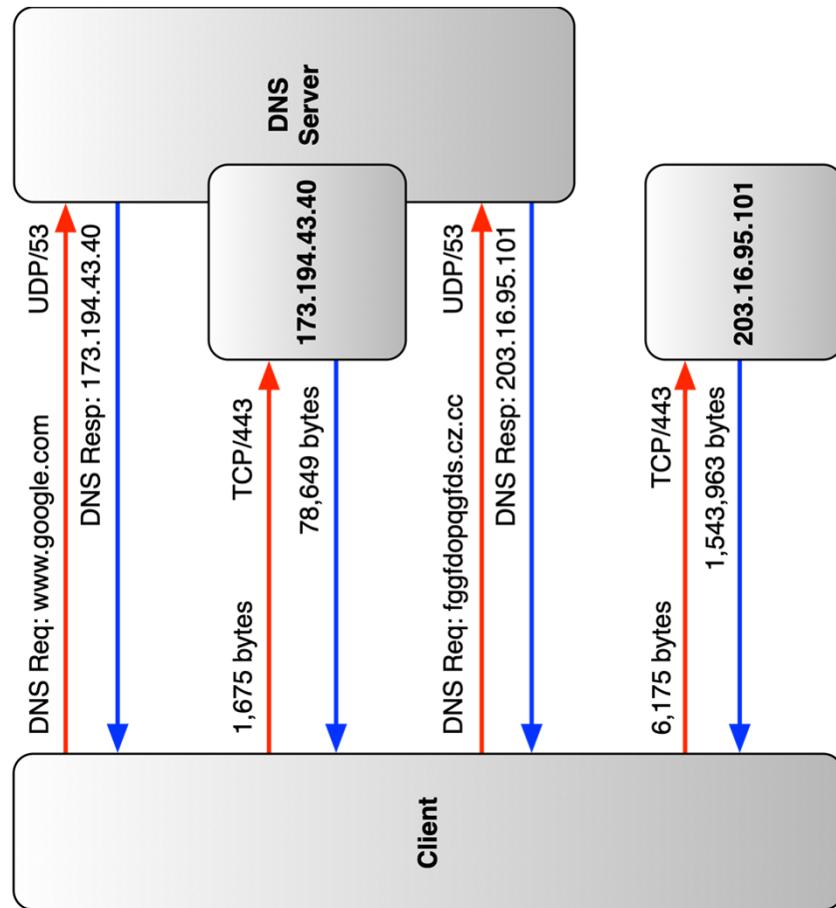


With the exception of the first few packets in the TLS negotiation, the remainder of the TLS traffic will be functionally opaque. For this reason, NetFlow offers an exceptionally valuable way to characterize the connection without incurring the storage overhead of a full-packet capture system. However, even without the full visibility to the TLS negotiation process or more complex TLS interception scenarios, there are several ways that a network forensicator can mitigate the impact of TLS on his or her ability to conduct a thorough investigation.

Consider the sequence of events that occurs when a user hits the "Return" key after typing "`https://example.com`" in their browser. The system first performs a DNS lookup to get the proper IP address for the "`example.com`" hostname, then initiates a connection to that IP on port 443 (by default). If we have any insight into the client's DNS activity immediately before it initiates the HTTPS connection, we could learn the host they are attempting to connect with, better characterizing evidence from noncontent sources such as NetFlow.

This is no guarantee, of course, because the server could provide a certificate with a subject name that does not match what was expected from the DNS request, but in most cases, this will be a very useful way to combine different sources of evidence. If the user or client software ignores the mismatch and failed certificate validation, correlating the DNS activity to a particular certificate would be much more difficult.

The use of hostnames in general is a valuable correlative data point for inherently imperfect evidence such as NetFlow. Their presence in the pre-encrypted TLS negotiation phase provides opportunities to log the requested and delivered hostnames. Zeek and other network security monitoring platforms typically log these values. In any case, remember that hostnames are quite prevalent in more than just DNS traffic itself.



Analytic Mitigation: Certificate Metadata and Extraction



- Capture certificates, parse and extract fields
 - Zeek's x509.log contains certificate metadata
 - NetworkMiner, Zeek automatically extract certificates

```
$ zeek for572 -r tls_dfir.com_session.pcap
$ cat x509.log | jq '{id, "certificate.serial", "san.dns"}'

{
  "id": "FkaAld3lb8AnjqfgA",
  "certificate.serial": "03B608F6F16375CBA1FFD65C8C532EA27C99",
  "san.dns": [
    "dfir.com",
    "www.dfir.com"
  ]
}

$ grep FkaAld3lb8AnjqfgA files.log | jq '{ fuid, extracted }'
{
  "fuid": "FkaAld3lb8AnjqfgA",
  "extracted": "extract-1544121579.457714-SSL-FkaAld3lb8AnjqfgA"
}
$ openssl x509 -text -noout -inform der -in extract_files/extract-1544121579.457714-SSL-FkaAld3lb8AnjqfgA
...
Serial Number:
03:b6:08:f6:f1:63:75:cb:a1:ff:d6:5c:8c:53:2e:a2:7c:99
Issuer: C = US, O = Let's Encrypt, CN = Let's Encrypt Authority X3
Subject: CN = dfir.com
...
```

Another method that may be useful in addressing the hurdle of analyzing TLS traffic would be to extract important certificate details to a text file for further examination. Because the certificate is provided in the clear during the TLS handshake sequence, we can extract fields such as the common/subject names, serial numbers, and more. This would establish a certificate log, which could be a useful tool for later correlation.

Of course, because the subject of each certificate is a hostname, DNS query logs would be invaluable to perform this correlation, but even with the timestamp, source/destination IPs/ports, and a few key certificate fields, we could efficiently characterize the nature of otherwise opaque HTTPS traffic.

Zeek's features in this area are quite impressive. The x509.log file contains a significant amount of metadata fields for each observed certificate. Additionally, Zeek's file extraction features include certificates exchanged during the TLS negotiation process.

Yet another solution would be to use NetworkMiner^[1], which performs certificate extraction consistent with its other file extraction operations. Of course, since these are written out to disk as the input data is processed, an analyst could examine them with filesystem-based methods.

References:

[1] <https://for572.com/dlakp>

The screenshot shows the NetworkMiner 2.6 interface. At the top, there is a menu bar with File, Tools, and Help. Below the menu is a toolbar with various icons. The main window has several tabs: Hosts (60), Files (6), Images, Messages, Credentials, Sessions (16), DNS (159), Parameters (255), Keywords, and Anomalies. A search bar is located at the top of the main pane. The main pane displays a table of certificate details. The columns are: Frame nr., Filename, Extension, Size, Source host, S. port, and Destination. There are 180 rows in the table. Below the table, there is a navigation bar with buttons for back, forward, and search, followed by tabs for NetworkMiner, AssembledFiles, 70.32.97.206, TCP-443, and a dropdown menu. At the bottom, there is a section for Recent, Starred, and Home files, each with a preview icon and a tooltip.

Frame nr.	Filename	Extension	Size	Source host	S. port	Destination
91	dfir.com.cer	cer	1 549 B	70.32.97.206 [dfir.com] [for572.com]	TCP 443	192.168.21
91	Let's Encrypt Authority X3.cer	cer	1 174 B	70.32.97.206 [dfir.com] [for572.com]	TCP 443	192.168.21
178	dfir.com[1].cer	cer	1 549 B	70.32.97.206 [dfir.com] [for572.com]	TCP 443	192.168.21
178	Let's Encrypt Authority X3[1].cer	cer	1 174 B	70.32.97.206 [dfir.com] [for572.com]	TCP 443	192.168.21
180	dfir.com[2].cer	cer	1 549 B	70.32.97.206 [dfir.com] [for572.com]	TCP 443	192.168.21
180	Let's Encrypt Authority X3[2].cer	cer	1 174 B	70.32.97.206 [dfir.com] [for572.com]	TCP 443	192.168.21

Analytic Mitigation: Wireshark Decryption (I)



The Wireshark interface displays two frames of a TLS session. Frame 25620 (highlighted in red) shows encrypted application data (Content Type: Application Data (23), Version: TLS 1.0 (0x0301), Length: 576, Encrypted Application Data: e7cf772516041745 [Application Data Protocol: http-over-tls]). Frame 25629 (highlighted in red) shows the decrypted HTTP request and response. The request is a GET /wp-content/plugins/woocommerce/assets/js/frontend/woocommerce.min.js?ver=2.0.13 HTTP/1.1\r\n and the response is a 200 OK [text/javascript]. Both frames are part of a TCP stream (Source: 85.10.210.250, Destination: 172.16.62.100, Protocol: TCP).

Fortunately for us, Wireshark also provides inline TLS decryption functionality that can be used in several specific scenarios.

These screenshots show a "before" and "after" of Wireshark's SSL/TLS decryption. At first, the selected frame is a part of a TCP/443 connection with effectively opaque data as a part of a TLS stream (previously established). However, the second screenshot was taken after supplying the required key material detailed in the next slides. This shows that Wireshark is now able to parse the same frame into its various TLS and HTTP components, despite still being a part of the same encrypted TCP/443 connection.

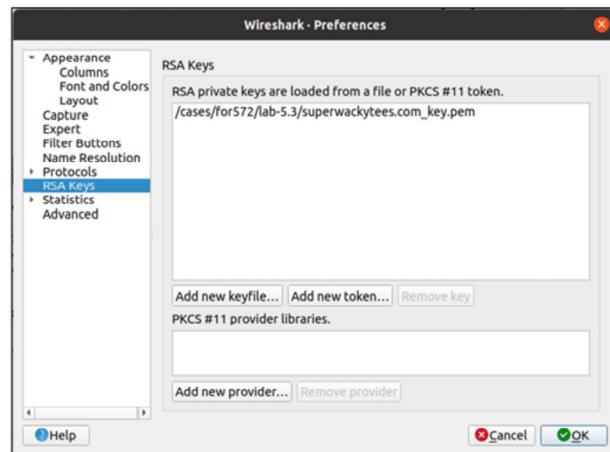
When using this functionality in both Wireshark and tshark, the analyst can use display filters against the decrypted traffic as well as perform field and object extraction using the respective protocol dissector's features.

No.	Time	Source	Destination	Protocol	Length	Info
25616	2013-08-29 17:39:19.856648	172.16.62.100	85.10.210.250	TCP	1514	443 → 49333 [ACK] Seq=7446 Ack=845
25617	2013-08-29 17:39:19.856666	172.16.62.100	85.10.210.250	TLSv1	888	Application Data
25618	2013-08-29 17:39:19.856683	172.16.62.100	85.10.210.250	TCP	60	443 → 49333 [FIN, ACK] Seq=9740 Ack=845
25619	2013-08-29 17:39:19.857236	85.10.210.250	172.16.62.100	TCP	54	49333 → 443 [ACK] Seq=845 Ack=9741
25620	2013-08-29 17:39:19.857484	85.10.210.250	172.16.62.100	TLSv1	672	Application Data, Application Data
25621	2013-08-29 17:39:19.857599	172.16.62.100	85.10.210.250	TCP	60	443 → 49334 [ACK] Seq=146 Ack=845
25622	2013-08-29 17:39:19.859429	172.16.62.100	85.10.210.250	TLSv1	1514	Application Data
25623	2013-08-29 17:39:19.859475	172.16.62.100	85.10.210.250	TCP	1514	443 → 49334 [ACK] Seq=1606 Ack=845
Frame 25620: 672 bytes on wire (5376 bits), 672 bytes captured (5376 bits)						
Ethernet II, Src: PcsCompu_3f:07:10 (08:00:27:3f:07:10), Dst: PcsCompu_01:f4:8a (08:00:27:01:f4:8a)						
Internet Protocol Version 4, Src: 85.10.210.250, Dst: 172.16.62.100						
Transmission Control Protocol, Src Port: 40224, Dst Port: 442, Seq: 227, Ack: 146, Len: 610						
Transport Layer Security						
TLSV1 Record Layer: Application Data Protocol: http-over-tls						
Content Type: Application Data (23)						
Version: TLS 1.0 (0x0301)						
Length: 32						
Encrypted Application Data: e7cf77251604174515d73d57c9f0110949f4da9a1e010ef6fb9fb32aec5fee3						
[Application Data Protocol]						
TLSV1 Record Layer: Application Data Protocol: http-over-tls						
Content Type: Application Data (23)						
Version: TLS 1.0 (0x0301)						
Length: 576						
Encrypted Application Data: 25620						
[Application Data Protocol]						
2013-08-29 17:39:19.857484						
17:39:19.857599						
172.16.62.100						
85.10.210.250						
HTTP						
672 GET /wp-content/plugins/woocommerce						
60 443 → 49334 [ACK] Seq=146 Ack=845						
[Application Data Protocol]						
25622						
2013-08-29 17:39:19.859429						
17:39:19.859475						
172.16.62.100						
85.10.210.250						
TCP						
1514 443 → 49334 [ACK] Seq=1606 Ack=845						
Transport Layer Security						
[2 Resassembled TLS segments (548 bytes): #25620(1), #25620(547)]						
Hypertext Transfer Protocol						
GET /wp-content/plugins/woocommerce/assets/js/frontend/woocommerce.min.js?ver=2.0.13 HTTP/1.1\r\n						
Accept: application/javascript, */*,q=0.8\r\n						
Referer: https://www.superwackytees.com/\r\n						
Accept-Language: en-US\r\n						
User-Agent: Mozilla/5.0 (compatible; MSIE 10.0; Windows NT 6.1; WOW64; Trident/6.0)\r\n						
Accept-Encoding: gzip, deflate\r\n						
Host: www.superwackytees.com\r\n						
DNT: 1\r\n						
Connection: Keep-Alive\r\n						
Cookie: wc_session_cookie_40dfb3c31fd6f81fec9b328d3df0497a=bPYK5km0jdzmL2f9VxLgs0rjD8nhYALF%7C%7C1377970759%7C%7C						
[Full request URI: https://www.superwackytees.com/wp-content/plugins/woocommerce/assets/js/frontend/woocommerce.mi						
[HTTP request 1/1]						
[Response in frame: 25624]						

Analytic Mitigation: Wireshark Decryption (2)



- Wireshark can decrypt some SSL/TLS traffic
- Non-PFS (e.g., RSA negotiation) requires:
 - Initial and re-key negotiations, private key
- tshark uses settings from Wireshark preferences file



```
$ tshark -n -r https.eth1_combined.pcap -Y 'tcp.port == 443 and http'  
26209 14287.776475 172.16.62.100 → 85.10.210.250 HTTP 1257 HTTP/1.0 200 OK (JPEG JFIF image)  
...
```

For Wireshark to decrypt TLS sessions where perfect forward secrecy is not used (for example, when using RSA negotiation process), the following conditions must be met:

- You have a full-packet capture for the entire session, including the initial key exchange and all subsequent re-key message exchanges
- You have the private key material for the connection (and the passphrase if needed)

To enable this feature, use the Wireshark GUI to specify one or more private key files in the “Edit | Preferences | RSA Keys” dialog. After this configuration is applied, Wireshark transparently decrypts and parses the contents of matching TLS sessions, allowing the common and convenient display filters, protocol inspection features, etc. These settings also persist across Wireshark sessions, as they are stored in a preference file on disk.

Naturally, tshark provides the same decryption functionality as well. This occurs transparently after using the Wireshark configuration dialog as indicated above.

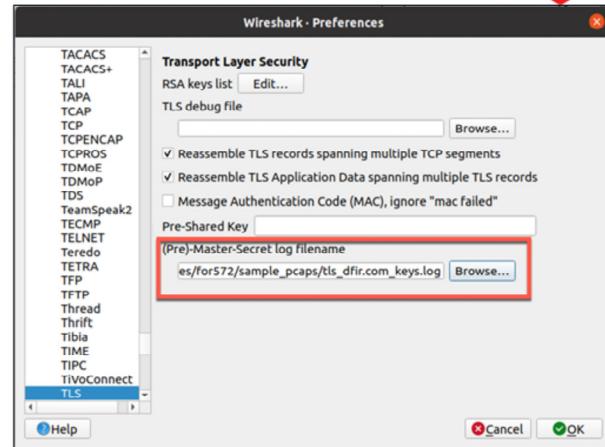
References:

<https://for572.com/p8jdo>

Analytic Mitigation: Wireshark Decryption (3)



- PFS (e.g., Diffie-Hellman) requires:
 - Initial and re-key negotiations, ephemeral secret key dump
- tshark uses settings from Wireshark preferences file



```
$ tshark -n -r tls_dfir.com_session.pcap -Y 'tcp.port == 443 and http'  
928 33.312070 192.168.211.134 → 70.32.97.206 HTTP 812 GET /product-category/posters/ HTTP/1.1  
...
```

However, if a TLS session does use perfect forward secrecy, we cannot use the server's private key to decrypt the contents—after all, this is the reason PFS was created. Instead, the following conditions must be met:

- You have a full-packet capture for the entire session, including the initial key exchange and all subsequent re-key message exchanges.
- You created a log file containing all ephemeral keys generated during the PFS negotiation.

Creating the log file is detailed in Sally Vandeven's SANS GIAC Gold paper^[1]. Two samples of both pcap file and corresponding key log files have been provided in your SIFT VM, in the “/cases/for572/sample_pcaps/” directory.

Enabling this functionality simply requires you to point Wireshark's SSL protocol parser to the Pre-Master-Secret log file, which it will then use to perform the decryption.

For command-line use, tshark also references this setting, allowing easy access to the decrypted data. Alternatively, you can specify this on the command-line with the “-o tls.keylog_file” option.

References:

[1] <https://for572.com/05dru>



Lab 5.1



SSL/TLS Profiling

This page intentionally left blank.

Lab 5.1 Objectives: SSL/TLS Profiling



- Use pre-encryption phase of SSL/TLS traffic to establish profiles of SSL/TLS clients
- Extract certificate subject names
- Extract specific SSL/TLS certificates for analysis

This page intentionally left blank.

Lab 5.1 Takeaways: SSL/TLS Profiling



- Even encrypted traffic can provide useful insight to the participating endpoints
- Building baseline profiles of normal activity and identifying specific instances of anomalous traffic are still attainable with encrypted traffic
- Certificates and their fields are useful IOCs

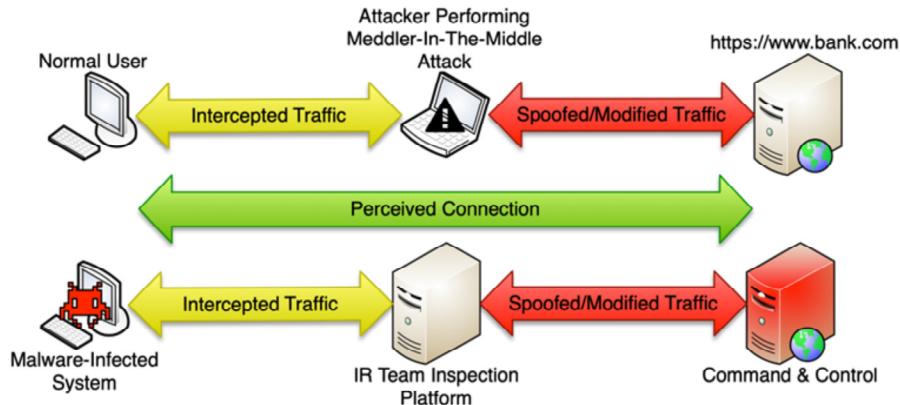
This page intentionally left blank.

Meddler-In-The-Middle

This page intentionally left blank.

MITM: Dual Use Technique

- When a system sits between two others
 - Allows for decrypting, filtering, manipulation
 - Used against plaintext and encrypted communications



While Meddling-In-The-Middle, or MITM, is generally thought of as an attacker methodology, it is actually a dual-use technique that is used in incident response operations as well. Although an attacker would need to manipulate the environment in some way to accomplish this, IR-based MITM is often as easy as a NAT rule or activation of a web proxy. When encryption is used, this complicated both attacker and IR MITM, but can be overcome with social engineering or system configuration, respectively. In that case, either party is performing TLS interception.

TLS inspection is a method by which an intermediary device sits between two systems (for example, an attacker and their target system or a user and an external TLS site). The interception/inspection device acts as both a client and server for any traffic that it encounters. This allows the device to decrypt the traffic and perform inspection of the content for malicious activity. To establish the perception of a successful connection, the device re-encrypts the traffic using its own private key and forwards the traffic to its final destination. This process could be applied to all manners of encrypted protocols but is most conducive to those using asymmetric or public key encryption.

Of course, the main challenge of TLS interception is that the intermediary presents a certificate that does not validate in the client's software. While there are enterprise solutions that can accomplish this, they rely on control over the client system's trust database. A far more nefarious possibility would be if the attacker leverage an existing certificate authority, then issued a forged but "legitimate" certificates that didn't suffer validation failure. This is why the compromises of the Comodo and DigiNotar certificate authorities were so alarming for SSL/TLS users.^[1]

As the Internet becomes more heavily encrypted, the use of TLS inspection will become more prevalent in many corporate environments. However, the use of such devices comes with technical and legal issues. From a technical standpoint, encryption and decryption are computationally expensive operations. Attempting to perform MITM at scale is a significant technical challenge. From a legal standpoint, legitimate websites (banks, email, etc.) are encrypted for a reason: They typically contain personal information that a normal user would not want revealed! When performing TLS inspection, an organization will likely have purview of their employees' personal data.

References:

[1] <https://for572.com/lo65k>

ARP Spoofing

- Floods LAN with spurious ARP replies so attacker's MAC address is associated with target's IP
 - More likely used to spoof the gateway IP address
- Victim(s) traffic framed to attacker's system
- Mitigations
 - Layer 2: DHCP Snooping + dynamic ARP inspection
 - Static ARP entries
 - OS-specific configurations

Of course, the Address Resolution Protocol (ARP) is used to translate network layer addresses (Layer 3 IP addresses) to link layer addresses (Layer 2 MAC addresses). When an IP datagram is sent from one host to another on a local area network, the sending host must determine the MAC address for the recipient system before the Layer 2 framing header can be applied. When the MAC address for that destination IP address is unknown to the sending system, it must discover that MAC address using ARP. The sending host will broadcast an “ARP request” packet and the destination machine will respond with an “ARP reply” that maps the IP address to the MAC address. The sending system will retain this pairing result in its own local ARP cache for a period of time, so the request/reply processes don't overwhelm the local network segment.

The weakness with ARP is that it is a stateless protocol with no means to validate the ARP replies. For efficiency, network devices cache ARP replies regardless of whether they actually requested them. These weaknesses lead to a MITM attack known as ARP spoofing.

ARP spoofing is the process by which an attacker with access to the Local Area Network segment floods the LAN with forged ARP replies associating a target's IP address with the attacker's MAC address. The systems on the local network segment will see and cache this malicious association for efficiency, then frame any packets destined for that target to the attacker's system instead. Typically, an attacker will use this method against a default gateway, but ARP spoofing can also be done on a per-host basis. When the traffic is redirected to the attacker's host, the attacker has the option of inspecting, manipulating, or denying traffic to/from the destination system. If the attacker carries this attack method out properly, the victim(s) will not even know their traffic is being intercepted, as the attacker can simply act as a router and pass the traffic along its intended path—after potentially inspecting, capturing, or manipulating the contents.

A number of defenses may be available to mitigate the risk of ARP spoofing. From the infrastructure's point of view, SOHO and enterprise switches may provide a mitigation feature that observes DHCP traffic and monitors ARP activity to identify erroneously or maliciously injected ARP activity^[1]. Endpoint-based mitigations include static ARP entries and configuration settings for various operating systems and network devices that cache only the

replies for which the system sends a request. While effective, these are often labor-intensive. Often, these are best applied in a targeted manner; for example, when a specific segment is observed to be under attack, or perhaps on a longer-term basis for high-risk assets.

References:

- [1] <https://for572.com/hnk5x>

Port Stealing

- Attacker manipulates a switch into receiving traffic that is destined to a victim system
 - Attacker sends fake Ethernet frames with the victim's MAC address
 - With enough data, the switch binds the attacker's machine to the victim's MAC address
- Mitigations
 - “Enterprise” grade switch with “port security”

“Port Stealing” is a MITM attack in which the attacker’s system successfully manipulates a network switch into diverting traffic originally destined for a victim’s system to their own. This is accomplished through a common vulnerability within Layer 2 Ethernet switches. These switches can learn what hardware is connected to each of its ports by listening to traffic sent to the switch port. As soon as a system sends any Ethernet frame to the switch, the switch caches this information into its internal state table (called a CAM table). Later, when the switch is making a determination on where to send a newly arrived packet, it inspects the Ethernet header to identify the destination MAC address. If the packet is destined for a computer that is cached in the CAM table, it is sent only to the port to which the destination system is connected.

However, if an attacker is connected to another port on the switch, they can send fake Ethernet frames with the victim’s MAC address forged into the source address field, essentially impersonating the victim machine. As an attacker sends data, the switch becomes confused and repeatedly alters the MAC address binding to either of the two ports by referencing the information within the packets. If the attacker sends data faster, the switch will send the attacker the packets intended for the victim host. When the traffic is redirected to the attacker’s host, the attacker has the option of inspecting, manipulating, or denying traffic to/from the destination system.

This method generally incurs a bit of collateral damage, however. It results in the victim seeing intermittent connectivity or being booted off the network entirely.

This attack is most likely to occur on smaller networks as most enterprise grade switches offer “port security”, which essentially locks a given MAC address to a single port at a time after it is first observed. Port security can also be enforced on some switch platforms by an administrator manually associating the MAC address and a specific switch port—but again, this is a very labor-intensive option.

UDP First Response Wins (I)

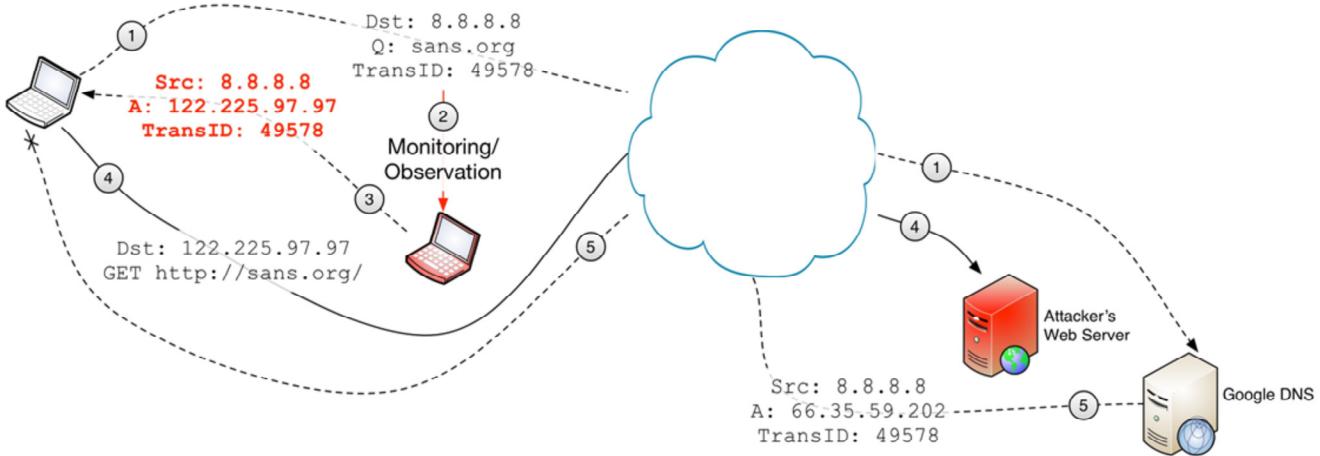
- UDP traffic lacks session tracking information
 - No session setup/teardown
 - No sequencing information
- Several protocols rely on this fact
 - UDP multicast and response
- Mitigations
 - Application layer validation

For all the efficiencies UDP provides by being “unreliable”, it also suffers from various security-related issues due to the lack of validation of the protocol itself. Traffic sent via TCP enjoys reliability by establishing an individual session through the three-way handshake and sequencing information. These ensure proper ordering and overall integrity of delivered payloads.

UDP, on the other hand, does not have these checks in place. Therefore, when a UDP-based application listens for a response to an outstanding request, it generally accepts the first response it receives—provided the IP and port information match what’s expected. A spoofed packet crafted to meet the expected parameters will be happily passed to the calling application. It is up to the application to validate the response and determine the truthfulness of the source of the information.

Several UDP-based protocols benefit from this paradigm, including DNS, DHCP, and more.

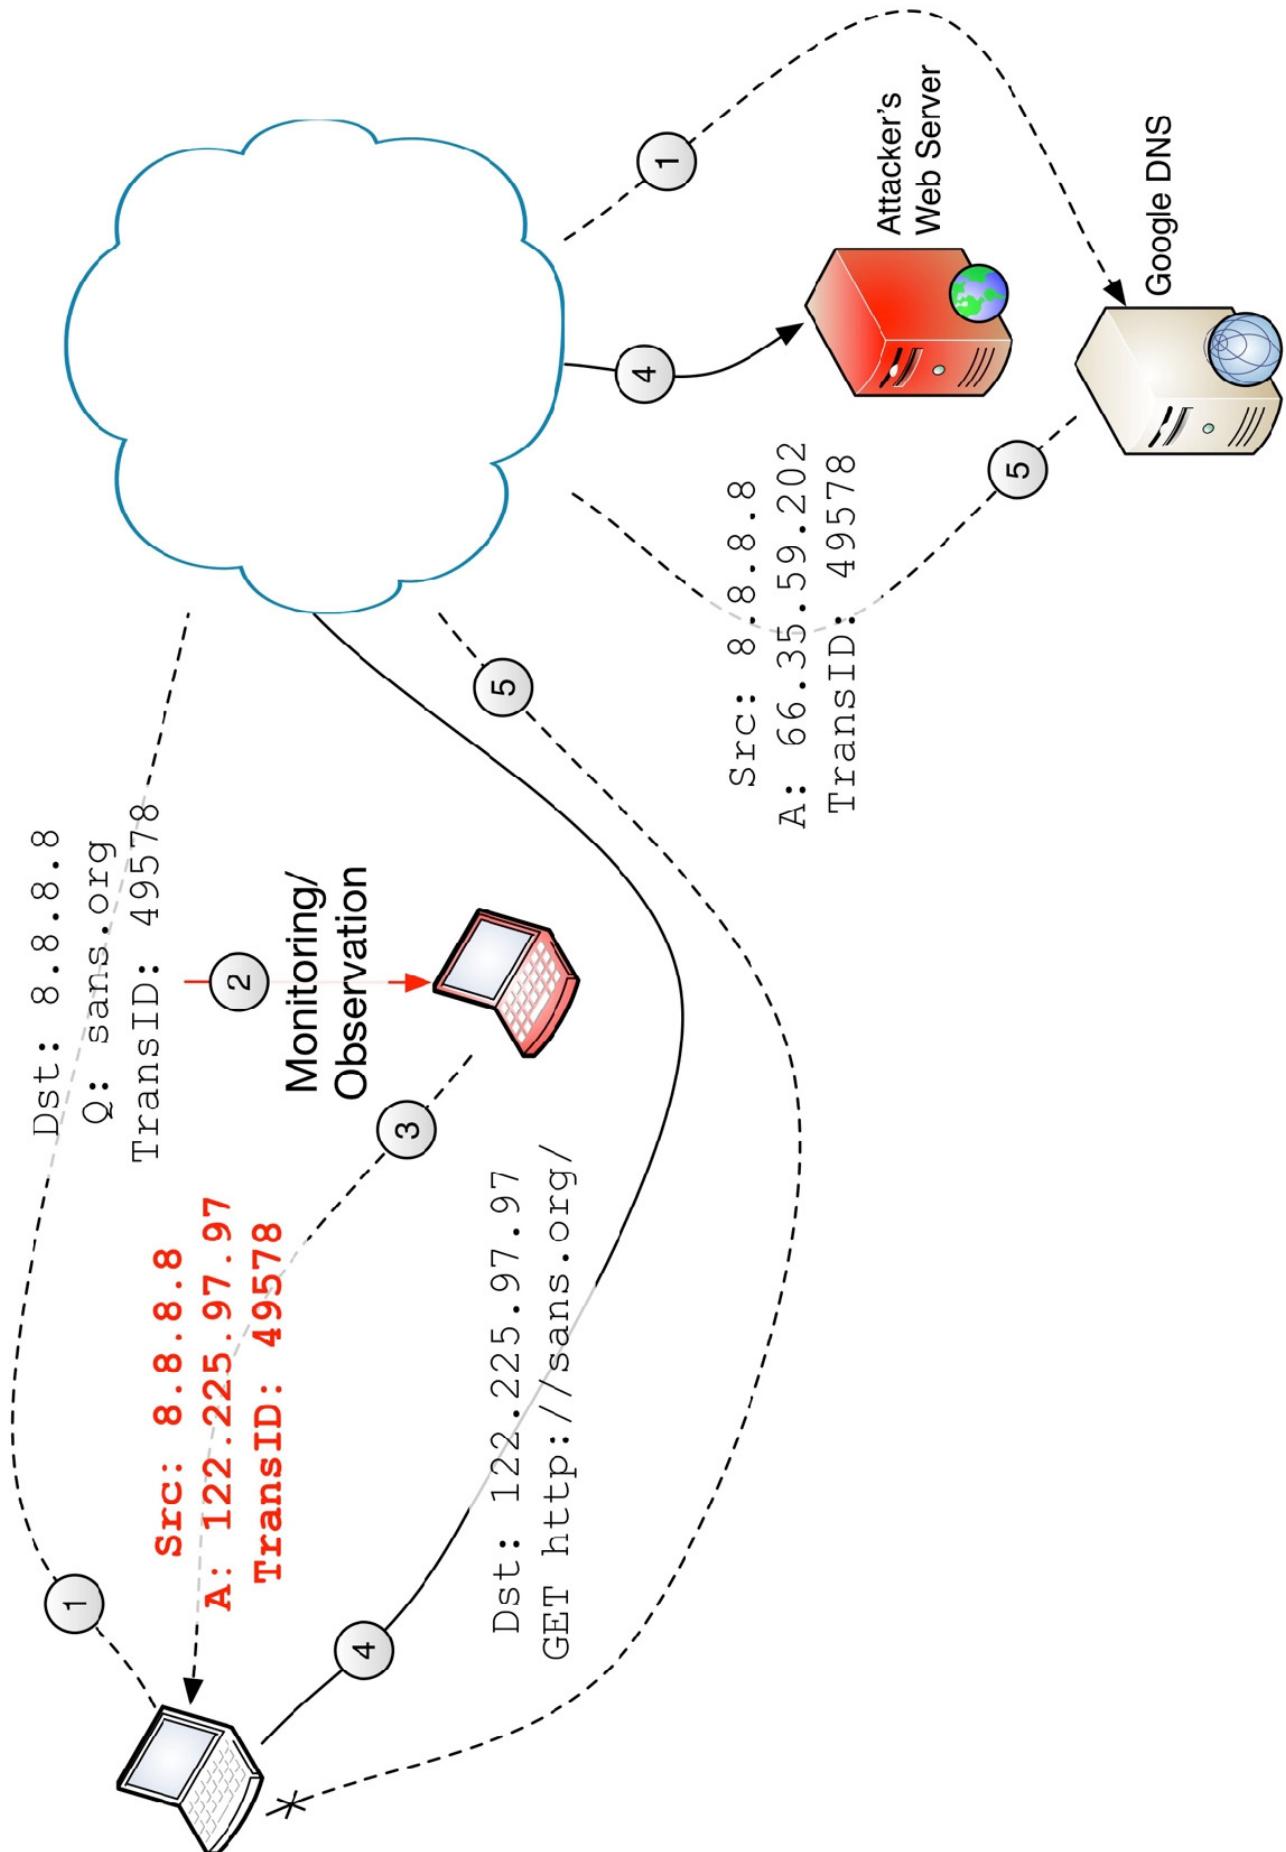
UDP First Response Wins (2)



An example of the “UDP first response wins” problem:

1. A client application sends a DNS request to a server using a third-party DNS provider outside the client’s network.
2. An attacker, monitoring the client’s local network, examines the DNS request.
3. The attacker crafts a response that corresponds to the request and crafts a spoofed reply that matches. The client receives the DNS response and interprets the results as if it had come from a DNS server.
4. The client establishes communication based on the contents of the spoofed DNS reply.
5. At some point, the client receives the response from the true DNS server it queried. However, because the outstanding DNS request has already been satisfied with the spoofed response, the client ignores the content.

As the attacker is monitoring the user’s network locally and is closer to the client application, the specially crafted DNS response would be received by the client before the DNS server providing responses outside the client’s network, thus, the client application would receive whatever response arrives first (in this case, most likely the attacker’s).



MITM Themes

- Network devices cache data for performance
 - Caches are **great** targets for manipulation
- Systems broadcast requests for dynamic services (ARP, DHCP)
 - Responses not typically validated
- Many network settings change dynamically
 - Enough data may manipulate those settings

Through the previous three examples, there are three common themes:

- Flooding a network device with information will ultimately modify the device's cache with information provided by the attacker, thus allowing the attacker to redirect traffic to their system. Caches are used by network devices to improve overall performance, making them an easy target for attack.
- Responding to broadcast requests faster than the legitimate respondent device allows the attacker to provide the first response and ultimately manipulate the target machine. The majority of broadcast requests are used for initial setup and dynamic configuration of hosts on a network. Because most responses to broadcast requests (ARP, DHCP, etc.) are not validated, the victim machine trusts the first reply.
- Sending enough data to a network device may manipulate dynamic settings associated with the device. Many network devices have several dynamic settings, as the Internet as a whole is a fairly dynamic and fluid place. Instead of statically implementing every network device setting, many organizations will allow devices to inherit settings based on the data that is seen across the network. Although this can reduce the overall likelihood of a static setting conflicting with a change in an upstream device, this also allows attackers to manipulate those settings, depending on the amount of data sent by the attacker.

There are several additional MITM attacks not covered here but that also follow the same three themes:

- STP Mangling: Allows an attacker to manipulate the spanning tree protocol configuration of a network device on a switched network to obtain data.
- DHCP Spoofing: Because DHCP requests are made as broadcasts, if an attacker responds faster than the real DHCP server, the attacker can manipulate the IP address, Gateway, and DNS assigned to the victim.
- ICMP Redirect: An attacker forges an ICMP redirect packet to redirect traffic to the attacker's machine.
- IRDP Spoofing: An attacker can forge an advertisement packet pretending to be a router on the LAN. The attacker can send a preference level and lifetime to ensure the attacker's system will be the preferred router.

Open-Source MITM Software: Bettercap and dsniff

- **Bettercap**

- ARP poisoning, dynamic host discovery
- URL, HTTP POST, credential dumping, transparent proxy



- **dsniff**

- Collection of tools focused on various protocols
- Can use ARP, MAC, or DNS spoofing for MITM
- Includes tools to MITM SSH, SSL/TLS, pull credentials, URLs, and more



Bettercap^[1] is a new incarnation of the venerable (but long since abandoned) Ettercap MITM tool. It primarily uses ARP poisoning to facilitate dynamic host discovery and traffic interception on a per-host or per-network basis.

After subverting ARP to intercept the traffic, Bettercap can employ any of a handful of modules to acquire information from the user's target(s). These allow the attacker to intercept, manipulate, or block the victims' traffic. These plugins include the following features out of the box, and the tool can be extended easily with Ruby code to accommodate new protocols and payloads.

- Credential collection (IMAP, POP, SMTP, SMB v1/v2, IRC, HTTP Basic/Digest)
- URL acquisition
- HTTPS hostname
- Transparent proxy

dsniff^[2] is an open-source collection of tools originally designed for penetration testing and auditing. It includes a number of applications designed to monitor various protocols for interesting data such as emails, passwords, and files. It also includes several applications (`arpspoof`, `dnsspoof`, `macof`) to facilitate the interception of network traffic normally unavailable to an attacker due to Layer 2 switching, allowing for the user to perform MITM attacks against targeted hosts. Also, `dsniff` includes two additional applications that allow for exploiting weak key bindings associated with some SSL/TLS and SSH connections.

References:

- [1] <https://for572.com/njzrs>
- [2] <https://for572.com/c5mfd>

Open-Source MITM Software: Yersinia

- Appropriately named after “Yersinia pestis” bacteria
- Linux/Unix only
- Focused on a large number of Layer 2 protocols
- Demonstrates even older tools can be ruthlessly effective against legacy protocols



Yersinia is another open-source tool that implements various MITM attacks against Layer 2 protocols that result in allowing the targeted system's traffic to be redirected through the attacker's.

Yersinia is capable of attacks against protocols that are fundamental to most networks—including many that are seldom monitored for misuse:

- Spanning Tree Protocol
- Cisco Discovery Protocol
- DHCP
- Hot Standby Router Protocol (HSRP)
- Dynamic Trunking Protocol
- 802.1Q
- 802.2X
- VLAN Trunking Protocol

References:

<https://for572.com/94siw>

MITM for Investigative Benefit

- Intercept traffic for inspection or blocking
 - Forward network proxies
 - Network Security Monitoring (NSM)
 - Intrusion Detection/Data Loss Prevention (IDS/DLP)
- Inspecting encrypted communications requires understanding their specific weaknesses
 - Weak keys
 - Weak key exchanges
 - SSL/TLS inspection and subverting trust

Meddler-in-the-middle “attacks” are not just used by attackers to inject, manipulate, or filter network traffic. Several legitimate technologies perform this to aid investigative efforts as well as for traditional network defense.

- Forward(ing) proxy: A network proxy that forwards requests from a client to the target server. Typically, forward proxies are capable of inspecting outbound requests and their subsequent responses to determine whether requests or responses should be blocked or filtered. Some common methods used for content filtering include URL or DNS blocking, IP blocking, filtering URLs or MIME content, or content keyword filtering. More sophisticated proxies may implement content analysis techniques to look for patterns of activity such as beaconing.
- Network Security Monitoring (NSM) platforms: Content inspection systems that log metadata from the observed protocols, with a focus on proactive data collection rather than reactively searching for signatures.
- Intrusion Detection System (IDS) and Data Loss Prevention (DLP): Broad content inspection attempting to identify patterns of activity that match known signatures of “interesting” behavior. These systems perform deep content inspection, alerting on known signatures. IDS and DLP solutions are similar in concept. Where they differ is in the directionality of traffic typically being inspected. An IDS generally observes inbound traffic, seeing known-malicious patterns. However, a DLP system inspects outbound traffic, using signatures for sensitive information to identify when that data leaves the environment. Both then raise an alert when the data of interest is observed, providing a tip to the proper team(s) to take responsive action.

When dealing with encrypted traffic, most traditional devices that perform network traffic inspection are foiled, because encryption is inherently intended to prevent inspection. However, depending on the environmental architecture and nature of the encrypted traffic, there are several options available that may allow security teams to be aware of what the attacker is doing. These include weak keys, weak key exchanges, or attacking at the termination point of encrypted traffic and then re-encrypting it to the target system. For example, the interception device might replace a client’s list of requested cipher suites with a single, weak cipher in what is called a downgrade attack^[1].

References:

[1] <https://for572.com/e7bzf>

When a weak encryption key is used, it makes a cipher more vulnerable to an overall attack. As mentioned during the prior module, the overall strength of a cipher is dependent upon multiple factors, including the key that is used to generate a more secure ciphertext. One design goal of encryption algorithm inventors is to have a “flat” key space where all keys are equally strong. However, for some algorithms (such as DES) a known (and mathematically insignificant) number of weak keys may be acceptable if the cryptosystem is built to prevent their use. In the case of DES, there are exactly four keys where the key for encryption and decryption is the exact same. Too many weak keys is a serious flaw in a cipher’s design because there is too high a chance that a key generated randomly will be a weak one, thus compromising the algorithm’s overall security.

Attackers implementing encryption functions can inadvertently use weak keys as part of their algorithm through the following methods:

- Using dictionary-based words as keys. This subsequently makes the entire ciphertext susceptible to dictionary-based attacks.
- Using keys that fill a portion of the key buffer (e.g., using a 20-bit key when the algorithm uses a 32-bit buffer). Typically, encryption algorithms will pad the rest of the buffer with “0x00” bytes, making the key easier to determine, which would then reveal the underlying ciphertext.

In certain circumstances, regardless of how strong a key is, if the key exchange between two cryptographic systems is weak, then the key is susceptible to being recovered and used for decryption purposes. It is not uncommon for attackers to use embedded hardcoded passwords within malware or transmit an encryption key through an alternate means.

Commercial Inspection Solutions: Consider Market Needs

- Most commercial demand for MITM is for HTTPS
 - May include other protocols over TLS
 - May include SSH/SFTP
- Web proxies
 - Zscaler, Symantec, Forcepoint, Squid*
- Firewalls/NGIPS
 - Cisco Firepower, Palo Alto
- Decrypting taps
 - Gigamon
 - NETRESEC Polar Proxy*



Several commercial solutions offer varying ranges of encrypted traffic inspection capabilities, and a few open-source solutions have entered the space as well. Although most large-scale solutions typically provide the ability to perform TLS inspection, some also allow for the same functionality against other secure protocols such as SSH/SCP and SFTP. Commercial solutions also provide additional benefits over open-source solutions, including a focus on performance and speed, in addition to offering a number of options combined into a single product. Commercial solutions offering TLS inspection capabilities typically fall into a few top-level infrastructure categories.

Web proxies handle HTTPS (and often any TLS-wrapped protocols) such as Zscaler's cloud offering, Symantec WebFilter platform (formerly Blue Coat WebFilter, now owned by Broadcom), and Forcepoint's Web Security (formerly Websense) solution. As an open-source alternative, Squid can provide this capability as well, although focused on HTTPS.

Firewalls, IDS, and so-called "next-gen" IPS devices have grown to accommodate TLS interception as well. Cisco's acquisition of the SourceFire company has grown into their current Firepower offering, and Palo Alto Networks has provided decryption capability in their firewalls for many years.

Lastly, consider the more DFIR- and Security Operations-focused decrypting tap. These devices perform the same function as a traditional tap; however, by breaking the TLS connection such as the devices above, they provide a decrypted data stream for plaintext analysis. This means existing tools that cannot handle encrypted data as well as plaintext can still be used on the decrypted revenue from these devices. Gigamon provides this feature in their suite of solutions, and a recent open-source solution from NETRESEC named PolarProxy^[1] is a capable free alternative for small-to-medium scale requirements.

Because these inspection solutions represent legitimate uses of the MITM model, they also assume a level of trust between the client systems and the TLS inspection platform. One way most of these systems provide a transparent user experience is through dynamic certificate generation. The inspection platform identifies the hostname the client is requesting, then uses its own CA (which is already trusted by client systems) to generate a certificate that allows the TLS negotiation to complete.

A newer approach to this model involves applying the infrastructure-based decryption and functionality in the form of a tap^[2]. This allows incident responders to use legacy platforms such as Arkime, Zeek, or similar solutions with the decrypted data flow, since they benefit most from examining plaintext traffic. Of course, such approaches require careful legal consideration and will likely require exclusions for certain types of traffic that generally carry information deemed too personal for such inspection and retention.

Many organizations have taken the approach to limit this level of more intrusive monitoring and capture to just the highest-security enclaves within their environments. This limits risk due to the decreased scope of collection.

References:

- [1] <https://for572.com/6kegg>
- [2] <https://for572.com/rdsf8>

Limitations of Infrastructure-Based TLS Inspection

- Certificate and public key pinning
 - Hardcoded certificate details in the client
 - Deviations from expected certificate/CA cause error
 - Mostly implemented in client binary applications
- Browsers now moving toward supporting certificate transparency list (CTL) validation
 - Locally-managed CA certificates generally exempt

One noteworthy technique that will all but prevent TLS inspection often performed in an enterprise environment is known as “certificate pinning” or “public key pinning”. These can take several forms, but all rely on some knowledge of the acceptable certificates or certificate authorities being hardcoded into the client software. When used, any mismatched data during the TLS negotiation process can be detected and avoided.

While most often implemented in mobile or native applications, regular web browsers may also provide this functionality to some extent. In either case, the client software has a preconfigured list of hashes for allowed certificates based on their domain or common/subject name. Since an intercepting party (with benign or evil intent) would issue a certificate that does not match that expected hash value, the client can easily identify that it is not communicating with the intended host. This allows the client to display an error and/or refuse to continue the communication, depending on the developer’s implementation.

In an enterprise environment, this may require category- or domain-based exceptions to TLS interception. For example, the Zscaler company provides a list of services that are known to be incompatible with their TLS interception product^[1]. Of note is that many of the affected services on that list (such as Dropbox, Evernote, and others) are only impacted with the mobile or native applications, while browser-based usage can be transparently intercepted as usual, provided the adequate trust relationships have been enabled by an administrator. Other environments may need to except known banking sites based on domain or host categorization.

Note that this type of pinning is different from the well-intentioned but now extinct “HTTP Public Key Pinning” (HPKP) standard. With HPKP, a browser would remember the first certificate it received for an HTTPS host, then raise an error if that certificate changed without appropriate re-negotiation steps. This caused countless problems, including the core reliance on a first-received certificate being valid, as well as many more. While an interesting approach, all major browser developers have now removed HPKP support^[2].

During 2018, major browser developers announced their mandate for TLS servers to all support Certificate Transparency^{[3][4][5]}, which aims to prevent subverted certificates such as a TLS inspection device would issue. However, to permit this in an enterprise environment, such restrictions are generally limited to the system's core trusted CA list, not locally-managed certificates. Therefore, TLS inspection of browser-based traffic is expected to be functional for the foreseeable future.

References:

- [1] <https://for572.com/zxq8g>
- [2] <https://for572.com/cbouf>
- [3] <https://for572.com/c7ed3>
- [4] <https://for572.com/8wsze>
- [5] <https://for572.com/908bp>

More Encryption + No Inspection Does Not Mean Lost Hope

- As encryption becomes more prolific and TLS inspection remains complicated, hope is not lost
 - Examination of metadata fields, profiling with JA3 hashes
 - NetFlow-based analysis applies equally to encrypted and unencrypted communications
 - Logs from endpoints of encrypted communications
 - Malware-focused and protocol-focused reverse engineering can yield critical insights
- Job doesn't necessarily get "harder"—just "different"

Although the trend toward more extensive use of encryption has been long underway and arguably irreversible, the use of TLS inspection technologies remains legally and technically complicated. However, although these issues will take time and investment to fix, there are still significant benefits DFIR professionals can provide even when faced with functionally opaque communications.

Examining the fields that are observable during the encryption negotiation process can provide a wealth of useful artifacts. The numerous fields from TLS certifications provide ample profiling opportunities. Techniques such as generating JA3 hashes can provide valuable and scalable insight to the client behaviors from a given network or data set.

NetFlow analysis is a content-agnostic investigative medium that can provide equally valuable insights into any communications, regardless of whether they are encrypted or not. Additionally, advances in the state of reverse engineering and protocol analysis often yield deep insight into the nature of encrypted communications that can often give a forensicator the critical foothold needed to pick apart an encryption implementation based on its uncovered weaknesses.

Logs generated at the client or server endpoints of an encrypted session can provide valuable insight into the nature of the communication. These are in plaintext form, since the endpoints are the terminal nodes of the encrypted interaction.

Although these both represent a change in our collective investigative approach, they do not necessarily make our work "more difficult". To the contrary, early investment in developing these skills will pay tremendous dividends in the future.

References:

- <https://for572.com/dre39>
- <https://for572.com/odbqz>

Network Protocol Reverse Engineering

This page intentionally left blank.

Network Protocol Reverse Engineering

- Process of extracting structure, attributes, and data from network protocol implementations
- Two primary purposes
 - Generate network-based indicators to identify additional malicious activity
 - Develop protocol decoders to identify and explain attacker activity

Network protocol reverse engineering can best be defined as the process of extracting structure, attributes, and data from a network protocol implementation without access to its specifications. Typically, the content and purpose of a protocol is determined through "conventional" reverse engineering of the malicious binary using static or dynamic analysis. This process can be time-consuming and for some protocols may be unnecessary. In certain circumstances, the goal is to obtain as much information as possible as quickly as possible to support one of two goals: the generation of additional network-based indicators to identify additional malicious activity and the development of protocol decoders to identify what actions the attacker carried out.

Protocol Attributes

- Protocol structure
- Protocol flow
- Encapsulation
- Protocol functionality
- Encoding/encryption routines

When performing protocol reverse engineering, analysts typically focus on attributes that can be used to generate signatures or uniquely identify the protocol in question. These attributes commonly include the following:

- **Protocol structure:** The layout of control signaling, metadata, and payload data for each command issued as part of the protocol
- **Protocol flow:** The timing, order, size, and directionality of each complete command and corresponding response
- **Encapsulation:** The protocol encapsulating the user's data (i.e., if the carrier protocol is ICMP, UDP, TCP, SSL/TLS, HTTP, etc.)
- **Protocol functionality:** The set of functions and commands that may be issued to a client by the attacker
- **Encoding/encryption:** The means by which each protocol datagram is transformed prior to encapsulation, commonly used to evade generic IDS signatures by attackers

Approaching Protocol Reverse Engineering

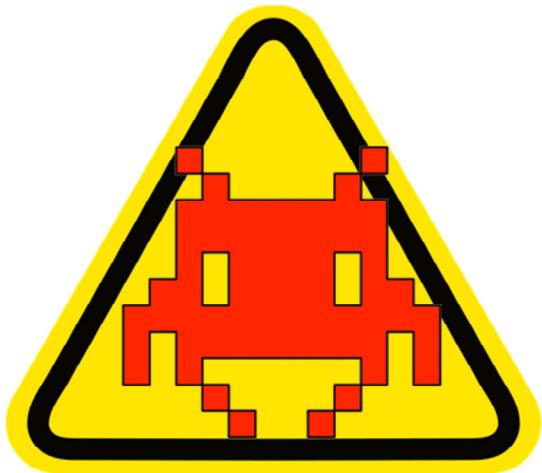
- Requires one of three data sources
 - Client binary or source code
 - Server binary or source code
 - Captured network traffic
- Identify what you are attempting to achieve
 - Balance goals against time and effort to achieve them
 - You don't need to know everything to determine how malware works or how it's most relevant to your task

Performing protocol reverse engineering typically involves using at least one of three sources of data: the client binary or source installed on the compromised machine, the server binary or source code used by the attacker, and the captured network traffic. Ideally, all three items are available but in reality, typically only the client-side binary and/or network traffic are available.

Outside of requiring data sources to perform the protocol reverse engineering, an analyst must also determine what their ultimate goal is. Many protocols today have levels of complexity that might make this goal impossible without a significant amount of time and effort. However, realistic goals in understanding how a protocol works can help to develop additional network-based indicators to identify future malicious traffic on a network. What this also means is that an analyst may need to become comfortable with several principles that typically go against digital forensics.

Through protocol reverse engineering, you may not need to know everything about a particular protocol to identify both how to detect the protocol for additional malicious activity or identify the capabilities of the malware itself (such as file uploads and downloads, etc.). If, as an analyst, you identify a unique string within the protocol that always exists among a large enough population of traffic, that information is good enough to write an indicator that can identify subsequent malicious activity.

Disclaimer



All protocols discussed in this section are real-life samples used by Advanced Persistent Threat Actors

All the protocols contained within this module are real-life samples used by Advanced Persistent Threat Actors and documented by various security companies. Although the protocols are the same, the names and faces have been modified to ensure no organizational or actual APT indicator information is released. You can read additional information about each sample based on the report linked within the notes section.

Similarly, we will follow a path of a traditional APT-based compromise, something that has been repeated thousands of times at organizations all across the world. Assume an organization has been targeted by a spear phishing email. One of the recipients assumed the email was legitimate and opened the attached PDF, infecting their system with a first-stage downloader communicating to an external website.

Compromised Web Pages/Sites

- Commonly used by attackers to
 - Host malicious executables
 - Embed commands read by malicious software
 - Participate in Fast-Flux DNS architectures
- Hides in plain sight on legitimate sites
 - Complicates use of endpoint threat intelligence
- Analysis requires:
 - Knowledge of the HTTP Protocol
 - Knowledge of HTML standards
 - JavaScript/CSS coding skills

Attackers commonly use legitimate websites and domains to host malicious information and executables. This allows the attackers to make use of resources like processing power, bandwidth, and the hosting availability of compromised web servers in addition to being able to mask the true source of their location. However, because in many cases the websites are legitimate websites, an analyst may need to troll through a significant amount of information to identify what is actually happening. For example, does the site contain malicious JavaScript that is redirecting to another site to download a file, does the web page contain embedded information that is read by a malicious backdoor, etc.? Having a thorough understanding of both the HTTP protocol and HTML assists in the identification of what is actually happening.

Example 1: Request



- One day in late June, 2011...

```
GET /index.htm?1308772046 HTTP/1.1
User-Agent:
    IPHONE8.5(host:malwarehunter, ip:192.168.1.135)
Accept: */*
Host: bad.domain.com
Connection: Keep-Alive
```

```
$ date -u -d @1308772046
Wed 22 Jun 19:47:26 UTC 2011
```

This first example covers a request that is issued by a compromised system after the user opens a malicious email. You spoke with the victim and identified the time the user opened the email and started your search for additional activity. Immediately after opening the email, you identify some suspicious HTTP-based traffic. Although basic, you start to identify additional oddities. For example, in this case, the GET request states the User-Agent is IPHONE8.5; however, the user claims to be using their laptop, not an iPhone. Similarly, embedded within the request is the name of the compromised host and the IP address. Additional research indicates that there are no known User-Agents that start with IPHONE8.5.

This request has at least one attribute (the unique user-agent) that would be able to be turned into a signature to identify additional malicious activity on the network. Similarly, additional information is embedded within the User-Agent that can assist the analyst in identifying additional compromised systems. It is not uncommon for attackers to embed identifying information into common request headers. This allows attackers to look through log files and differentiate between which systems may be connecting back. Although sometimes the information is not encoded, many times it is.

Example 1: Response



```
HTTP/1.1 200 OK
Server: Microsoft-IIS/5.0
X-Powered-By: ASP.NET
Connection: close
Content-Type: text/html
Accept-Ranges: bytes
Content-Length: 2550

<yahoo sb="Sv|@yC2wyMxwy1#XC|wm(26322)"></yahoo>

<html>...
```

When reviewing traffic, you examine the response associated with the issued request. When looking at the response, you easily identify the response header indicating the server responding is a Microsoft-IIS/5.0 server and the application that is providing the response is Microsoft's Active Server Page programming language, which indicates the page may be capable of generating dynamic content. Lastly, the total content length for the returned information is 2550 bytes.

However, when reviewing the actual data returned, you identify an odd tag `<yahoo sb=></yahoo>` that looks like normal HTML tags but comes before the `<html>` tag associated with the returned document. With an understanding of HTML structure, you find this odd because typically tags are contained between `<html>` and `</html>` for the browser to render the information effectively. Similarly, the usage of the word “yahoo” raises suspicions, as in most cases that is not a typical name for an HTML tag, nor built within the HTML standard. It’s further suspicious through the random characters that appear between the tags.

Once again, without knowing exactly what the malicious software is doing, an analyst can easily pick out additional unique information that may be searchable—in this case, specifically the `<yahoo sb=>` and `</yahoo>` tags. This information can lead to further signatures or allow an analyst to look for additional unique information contained within the captured traffic.

Example I: Expanded Scope



- Additional samples:

```
<yahoo sb="Sv|@yC2wyMxwy1#XC|wm(26322)"></yahoo>
<yahoo sb="HCwSk6Y#kj3#kUg4fw#G(4378)"></yahoo>
<yahoo sb="qNH#Z|YM6GiKhe|Y3USeWD0ln(12336)"></yahoo>
...

```

Through your analysis, you decide to check for additional “yahoo sb” tags within the captured network traffic. From the results returned, you find at least two other instances, also calling out to the same malicious IP address. As with most network analysis (and analysis in general), the observations you make help identify patterns. Although you may still not be sure what the actual malware does, you have gotten to the point where you believe the malware dropped by the spear phishing email appears to interact somehow with information contained within the <yahoo> tags. Simultaneously, you can start to build some assumptions about the encoding of the data contained between the yahoo tags. Specifically, all three variants have an integer between parentheses. Two of the three variants have two pipe characters (|) between the tags; the third does not. Based on those assumptions, it appears that although it is not exact, the malware probably parses something along the lines of the following statement:

```
<yahoo sb=<encrypted data>(<integer>)"></yahoo>
```

Because the server response contained the phrase ASP.NET, which is a programming language used for generating dynamic web pages, you may also assume that the encoded statement is probably dynamically generated using ASP.NET. Because the data appears to be dynamically generated and is encoded, there is a likelihood that a key is needed to both encode and decode the data. The key may be contained within the malware itself or because there appears to be an integer field that changes dynamically but is contained between parentheses, the key itself could be a key that is used for encoding.

Example 1: Meta-Temporal Analysis



- “Meta” activity
 - After receiving the “<yahoo>...”-tagged page, each system made an additional download from the same IP address
 - Contents of the second download appear to be a Windows executable based on file magic (“MZ” bytes)

Finally, through your analysis, you decide to look at traffic coming and going from the malicious IP address. After each request, you see another request from the compromised system to download a file from the malicious IP address. This downloaded file appears to contain a Microsoft Windows portable executable header.

Example 1: Summary

- By analyzing just the protocol and not knowing anything about the malicious software, we can determine that the malicious software reads encoded commands from a legitimate web page between <yahoo></yahoo> tags and has the capability to download and execute a malicious file on a compromised system

Based on the analysis conducted so far, we've arrived at the following observations:

- The get request contains information about the compromised host, including its hostname and IP address within the User-Agent string.
- The User-Agent string is a nonstandard value that makes for a good indicator when looking for additional traffic.
- The response contains <yahoo> tags, which are nonstandard HTML tags.
- The information contained between the <yahoo> tags appears encoded and may be encoded using some algorithm with a dynamically generated key, potentially the integer value identified between the parentheses.
- The command, when decoded, most likely downloads a file from a location on the compromised server.

By analyzing just the protocol and without analyzing the malicious software, we can theorize that the malicious software reads encoded commands from a legitimate web page between <yahoo> tags and has the capability to download and execute a malicious file on a compromised system.

From the Mandiant APT 1 Report:

WEBC2-YAHOO enters a loop where every ten minutes it attempts to download a web page that may contain an encoded URL of a file to download and execute. The encoded URL will be found in the pages returned inside an attribute named “sb” or “ex” within a tag named “yahoo”. This tag may be placed anywhere in the page and be usable by the malware but in practice, it is normally placed at the very beginning of the page by the attacker.

The malware will send beacon requests every two or three minutes. When the malware receives a response from the server, it then searches the returned content for “<yahoo sb=”<%text%>(<%int%>) ”></yahoo>”. The “<%int%>” variable is an integer used to index into a crypto array when decrypting the “<%text%>” portion.

References:

<https://for572.com/5aqly>

Backdoors Hidden in Common Protocols

- Malware may use common protocols for command and control purposes
 - Often uses protocols with arbitrary and optional fields
 - Know normal to spot outliers for further inspection
- Allows hiding in plain sight
- Analysis requires
 - In-depth knowledge of the protocol
 - Knowledge of usage of the protocol
 - Pattern recognition

Although malicious software can contain custom protocols, oftentimes malicious software uses existing libraries and common protocols for command and control purposes. As seen within the prior example where the attacker added host information to the User-Agent string, in many cases such information may be encoded within the GET or POST statements, cookies, or other fields of the HTTP backdoor. This allows the attacker to embed identifying information within a common protocol to “hide in plain sight” and potentially make it past a company’s web-based proxy, something that may be harder for a typical binary protocol to achieve. Similarly, an attacker may wrap a custom protocol within HTTP headers to ensure their communications can flow over common protocols and avoid detection.

Example 2: Request



```
GET /1.html HTTP/1.1
Cookie: Y29tbWFuZD1HZXRDb21tYW5kO2NsawVudGtleT0zOTU0O2hvc
3RuYW1lPW1hbHdhcmVodW50ZXI7
User-Agent: Mozilla/4.0 (compatible; MSIE 6.0;
Windows NT 5.1; SV1; .NET CLR 2.0.50727; .NET CLR
3.0.04506.648; .NET CLR 3.5.21022; .NET CLR
3.0.4506.2152; .NET CLR 3.5.30729)
Host: 123.123.123.123:8080
Connection: Keep-Alive
```

As part of your analysis, you look for additional traffic that originated from the compromised system—specifically traffic that occurred immediately after the malicious executable was downloaded (and likely executed). You filter out all the known malicious traffic and come across the following GET request, which is the first GET request after the executable. When looking at the GET request, you initially go through all the parts of the request itself. The URI information and format appear to be correct, the user-agent string appears to be valid, and the host information appears to contain an IP address and port. However, one oddity that does just not look right within the header is the Cookie field.

The HTTP protocol does not define the layout of fields contained within HTTP headers. However, as an analyst armed with basic knowledge of how browsers work in general, you identify the Cookie field as typically contained at the bottom of a request if issued by a web browser. For example:

```
Safari:
GET /favicon.ico HTTP/1.1
Host: www.google.com
User-Agent: Mozilla/5.0 (Macintosh; Intel Mac OS X 10_8_4)
AppleWebKit/536.30.1 (KHTML, like Gecko) Version/6.0.5 Safari/536.30.1
Accept: */*
Referer: http://www.google.com/
Accept-Language: en-us
Accept-Encoding: gzip, deflate
Cookie: SID=DQAAAO8AACI3BDW391zpe8Zra-J7nDJg3t0<snip>
```

Chrome:

Host: www.google.com
Connection: keep-alive
Accept: text/html,application/xhtml+xml,application/xml;q=0.9,*/*;q=0.8
User-Agent: Mozilla/5.0 (Macintosh; Intel Mac OS X 10_8_4) AppleWebKit/537.36 (KHTML, like Gecko) Chrome/28.0.1500.95 Safari/537.36
X-Chrome-Variations: CNO1yQEIl7bJAQiptskBCMS2yQEIl0TKAQj7hMoBCLiFygE=
Accept-Encoding: gzip,deflate, sdch
Accept-Language: en-US,en;q=0.8
Cookie: PREF=ID=5138dfe5

Firefox:

Host: www.google.com
User-Agent: Mozilla/5.0 (Macintosh; Intel Mac OS X 10.8; rv:22.0)
Gecko/20100101 Firefox/22.0
Accept: image/png,image/*;q=0.8,*/*;q=0.5
Accept-Language: en-US,en;q=0.5
Accept-Encoding: gzip, deflate
DNT: 1
Referer: http://www.google.com/
Cookie: NID=67=sQ8UktBtJk4QGQgS3HJro

Example 2: Request (Decoded)



The screenshot shows the CyberChef interface with the following configuration:

- Operations:** Regular expression
- Input:** GET /1.html HTTP/1.1
Cookie: Y29tbmF0ZD1HZXRDb21tYW5kO2NsawWvudGtleT0zOTU002hvc3RuYw1PWhbhHcmVodw50ZXI
User-Agent: Mozilla/4.0 (compatible; MSIE 6.0; Windows NT 5.1; SV1; .NET CLR 2.0.50727; .NET CLR 3.0.4506.640; .NET CLR 3.5.21022; .NET CLR 3.0.4506.2152; .NET CLR 3.5.30729)
Host: 123.123.123.123:8080
Connection: Keep-Alive
- Regular expression:** ^Cookie: (.*)\$
- Options:** Case insensitive, ^ and \$ match at newlines
- From Base64:** Alphabet: A-Za-z0-9+=, Remove non-alphabet chars
- Split:** Split delimiter: ;, Join delimiter: \n
- Output:** command=GetCommand
clientkey=3954
hostname=malwarehunter

Armed with this knowledge, you identify that this information is unlikely to be requested by a web browser and potentially by some other application, maybe the newly downloaded backdoor. When looking at the Cookie information, you remember that HTTP protocols commonly encode information in an encoding scheme that is guaranteed to return ASCII text such as base64, UTF-7. Trying some of the more popular encoding routines, you stumble across the following when decoding via base64.

```
command=GetCommand;clientkey=3954;hostname=malwarehunter;
```

This can also be decoded with the following CyberChef recipe:

```
[{"op": "Regular expression",  
 "args": ["User defined", "^Cookie: (.*)$",  
         true, true, false, false, false,  
         false, "List capture groups"] },  
 {"op": "From Base64",  
 "args": ["A-Za-z0-9+=", true] },  
 {"op": "Split",  
 "args": [";", "\n"] }]
```

With absolutely no knowledge of the malware but experience based on knowledge of protocols and understanding of HTTP, you are successfully able to identify and start to recover commands issued by the backdoor.

References:

<https://for572.com/jler7>

Last build: 6 hours ago - Enigma, Typex and the Bombe operations added for GCHQ's Centenary

Operations		Recipe	Input	Output
Search...		Regular expression	<input checked="" type="radio"/> II Built in regexes User defined Regex <code>^Cookie: (.*)\$</code>	GET /1.html HTTP/1.1 Cookie: Y29tbWFlZD1tYw5k02NsawUvdGtlet0zOTU02hvc3RuyW1lPw1hbhdhcwvodm50ZXI7 User-Agent: Mozilla/4.0 (compatible; MSIE 6.0; Windows NT 5.1; SV1; .NET CLR 2.0.50727; .NET CLR 3.0.0456.648; .NET CLR 3.5.21022; .NET CLR 3.0.4306.2152; .NET CLR 3.5.30729) Host: 123.123.123.123:8080 Connection: Keep-Alive
Favourites			<input checked="" type="checkbox"/> Case insensitive <input checked="" type="checkbox"/> ^ and \$ match <input type="checkbox"/> at newlines <input type="checkbox"/> Unicode support <input type="checkbox"/> Astral support <input type="checkbox"/> Dot matches all <input type="checkbox"/> Display total	
Data format				
Encryption / Encoding				
Public Key				
Arithmetic / Logic				
Networking				
Language				
Utils				
Date / Time				
Extractors				
Compression				
Hashing				
Code tidy			<input checked="" type="checkbox"/> Remove non-alphabet chars	
Forensics				
Multimedia				
Other				
Flow control				
STEP				Auto Bake

Mixed Protocol Analysis

- Mixes common protocols with binary protocols
- Example may include binary data wrapped within HTTP headers
- Analysis requires:
 - In-depth knowledge of the protocol
 - Knowledge of usage of the protocol
 - Knowledge of binary data structures
 - Pattern-recognition skills

Analyzing mixed protocols (such as binary-based information submitted over HTTP) provides attackers the ability to transmit custom-encoded data within headers that appear to be legitimate and may bypass traditional detection mechanisms such as web proxies. Analyzing mixed protocols requires knowledge of both the common protocol (in our case, this will be an HTTP header) and knowledge of identifying patterns within binary output.

Example 3: HTTP POST Request Headers



```
POST /search25548?h1=FIFEFDAHAPGDENCNFNFFFNAGAH
HTTP/1.1
User-Agent: Mozilla/5.0 (compatible;Windows NT 5.1)
Host: 123.123.123.123
Content-Length: 144
Connection: Keep-Alive
Cache-Control: no-cache
```

When looking at additional traffic, you identify yet another potential odd request header. When looking at the header itself, you see another User-Agent string that doesn't appear to be a normal User-Agent. When looking at your favorite resources for identifying common User-Agent strings, you notice that most strings have a space after a semicolon for information contained between parentheses. While not necessarily indicative of malicious activity, you decide to extract the POST data to a file so you can look more closely at its contents.

00000000	50 4f 53 54 20 2f 73 65 61 72 63 68 32 35 35 34	POST /search2554
00000010	38 3f 68 31 3d 46 49 46 45 46 44 41 48 41 50 47	8?h1=FIFEFDAHAPG
00000020	44 45 4e 43 4d 46 4e 46 46 46 4e 41 47 41 48 20	DENCMFNFNAGAH
00000030	48 54 54 50 2f 31 2e 31 0a 55 73 65 72 2d 41 67	HTTP/1.1.User-Ag
00000040	65 6e 74 3a 20 4d 6f 7a 69 6c 6c 61 2f 35 2e 30	ent: Mozilla/5.0
00000050	20 28 63 6f 6d 70 61 74 69 62 6c 65 3b 57 69 6e	(compatible;Win
00000060	64 6f 77 73 20 4e 54 20 35 2e 31 29 0a 48 6f 73	dows NT 5.1).Hos
00000070	74 3a 20 31 32 33 2e 31 32 33 2e 31 32 33 2e 31	t: 123.123.123.1
00000080	32 33 0a 43 6f 6e 74 65 6e 74 2d 4c 65 6e 67 74	23.Content-Lengt
00000090	68 3a 20 31 34 34 0a 43 6f 6e 6e 65 63 74 69 6f	h: 144.Connectio
000000a0	6e 3a 20 4b 65 65 70 2d 41 6c 69 76 65 0a 43 61	n: Keep-Alive.Ca
000000b0	63 68 65 2d 43 6f 6e 74 72 6f 6c 3a 20 6e 6f 2d	che-Control: no-
000000c0	63 61 63 68 65 0d 0a 0d 0a f3 d7 dd cc d1 cd d1	cache.....
000000d0	d8 ca 9e e9 d7 d0 da d1 c9 cd 9e e6 ee 9e e5 e8
000000e0	db cc cd d7 d1 d0 9e 8b 90 8f 90 8c 88 8e 8e e3
000000f0	b3 b4 96 fd 97 9e fd d1 ce c7 cc d7 d9 d6 ca 9e
00000100	8f 87 86 8b 93 8c 8e 8e 8f 9e f3 d7 dd cc d1 cd
00000110	d1 d8 ca 9e fd d1 cc ce 90 b3 b4 b3 b4 fd 84 e2
00000120	fa d1 dd cb d3 db d0 ca cd 9e df d0 da 9e ed db
00000130	ca ca d7 d0 d9 cd e2 cb cd db cc d0 df d3 db e2
00000140	fa db cd d5 ca d1 ce 80 be be be be be be
00000150	be be be be be be be



Example 3: HTTP POST Request Body

```
$ cat http_post_payload.bin | hexdump -C
00000000  f3 d7 dd cc d1 cd d1 d8 ca 9e e9 d7 d0 da d1 c9  | .....|
00000010  cd 9e e6 ee 9e e5 e8 db cc cd d7 d1 d0 9e 8b 90  | .....|
00000020  8f 90 8c 88 8e e3 b3 b4 96 fd 97 9e fd d1 ce  | .....|
00000030  c7 cc d7 d9 d6 ca 9e 8f 87 86 8b 93 8c 8e 8f  | .....|
00000040  9e f3 d7 dd cc d1 cd d1 d8 ca 9e fd d1 cc ce 90  | .....|
00000050  b3 b4 b3 b4 fd 84 e2 fa d1 dd cb d3 db d0 ca cd  | .....|
00000060  9e df d0 da 9e ed db ca ca d7 d0 d9 cd e2 cb cd  | .....|
00000070  db cc d0 df d3 db e2 fa db cd d5 ca d1 ce 80 be  | .....|
00000080  be  | .....|
```

Because the header information is suspicious, there's very little to go on and indicate the activity is malicious. However, the POST request itself appears to be binary information. One unusual observation is that not even a single byte of this content is printable ASCII.

When analyzing binary protocols, it is common to look for repeated patterns. This may suggest a key used for encryption or decryption because it will appear whenever the key is XORed with the null byte (0x00), which is often used for padding a message to a block boundary (usually 64 or 128 bits). Although in some cases a key may be easily recovered, testing against basic encoding or encryption techniques using permutations of the associated pattern can oftentimes lead to breakthroughs.

In this example, the entire last row contains the hex value “0xbe” repeated over and over. As the value appears to be used repeatedly at the end of the message, and the message itself aligns with a 128-bit boundary, the “0xbe” may indicate a one-byte key used for encoding the associated data.

References:

<https://for572.com/qqu12>



Example 3: HTTP POST Information Decoded

The screenshot shows the CyberChef interface with the following configuration:

- Operations:** Hex
- Recipe:**
 - From Hex (Delimiter: Auto)
 - Strip HTTP headers
 - XOR (Key: be, Scheme: Standard, Null preserving)
- Input:** A long hex string (length: 690 lines: 1) starting with 584f535420273656172636832353534383768313d4649464546444148415984744454e43d464e46464e41474148204854.
- Output:** Microsoft Windows XP [Version 5.1.2600]
(C) Copyright 1985-2001 Microsoft Corp.
C:\Documents and Settings\username\Desktop>.....

After XOR-decrypting the HTTP payload, its contents reflect plaintext ASCII. Examining the decrypted file contents reflects what appears to be a Microsoft Windows shell prompt in what was the POST data.

The CyberChef recipe for this decode is below:

```
[{"op": "From Hex", "args": [{"option": "Auto"}]}, {"op": "Strip HTTP headers", "args": []}, {"op": "XOR", "args": [{"option": "Hex", "string": "be"}, {"Scheme": "Standard", "Null preserving": false}]}]
```

This is an example of a basic (yet absurdly effective) web shell, in which the HTTP request contains the prompt. The server, which is under the attacker's control, responds with the command to be run.

Binary Protocol Analysis

- Intended purely for machine reading
- Typically contains message headers or metadata encoded within the protocol
- Analysis requires:
 - Basic binary data structure knowledge
 - Pattern recognition skills

A binary protocol is never expected to be human-read. Instead, software handles the parsing, providing denser messages with a wider variety of encoding options. This translates into speed of transmission and handling. Detecting a binary protocol requires some form of infrastructure visibility that can identify the protocol. For example, configuring a web proxy to block all non-HTTP-based traffic over port 443 will most likely result in blocking a binary protocol from communicating with its command and control server.

When analyzing binary protocols, typically additional metadata about the payload must be passed between the client and server. This information commonly includes the size of the following message, such that the server can translate the appropriate amount of data.

Example 4: GH0ST RAT (I)

- GH0ST RAT is a remote administration tool
 - Used in compromises of embassies, political opponents
 - Source code released by unknown parties
- Implements binary command and control protocol

```
struct gh0st_header {  
    char    gh0st_word[5];  
    DWORD   packetLength;  
    DWORD   uncompressedLength;  
    char    compressed_payload[];  
};
```

The GH0ST Remote Administration Tool (RAT)^[1] is a backdoor whose source code has been widely distributed online. It was used by a nation-state-like actor known as “Gh0stNet” to compromise numerous embassies and various offices of the Dalai Lama.

The compiled source code provides attackers with many ways to control a victim system, including the ability to create/manipulate/delete/launch/transfer files, perform screen capture, perform audio capture, enable a webcam, list/kill processes, open a command shell, and wipe event logs. The GH0ST RAT source code uses a custom binary protocol that is compressed using the ZLIB compression library and prepended with the `gh0st_header` structure shown here. It includes the following three fields:

- A static string appears in the first five bytes
- The size of the packet including the headers and compressed payload byte array
- The size of the decompressed payload byte array

References:

[1] <https://for572.com/lgm61>



Example 4: GH0ST RAT (2)

- Encoded Message

`0x8c000000 = 140`

`0xe0000000 = 224`

```
$ cat ghostrat.tcpflow.bin | hexdump -C
00000000  41 64 6f 62 65 8c 00 00  00 e0 00 00 00 78 9c 4b  |Adobe.....x.K|
00000010  63 66 60 98 c3 c0 c0 c0  0a c4 8c 40 ac c1 c5 c0  |cf`.....@....|
00000020  c0 04 a4 83 53 8b ca 32  93 53 15 02 12 93 b3 15  |....S..2.S.....|
00000030  8c 19 18 1c 58 5a 19 78  19 20 40 86 c1 02 ac 66  |....XZ.x. @....f|
00000040  01 50 40 01 88 05 1f 33  c0 01 23 d8 14 06 86 48  |.P@....3..#....H|
00000050  20 7e 00 55 c3 c3 07 64  28 30 30 5c 00 9a c1 c8  |~.U...d(00\....|
00000060  80 07 30 32 82 1d 01 52  c3 0c e5 ef e5 64 60 a8  |..02...R....d`..|
00000070  07 b2 7c 1d 7d c2 1d 83  5c 3d 42 fd 42 5c 83 f0  |...|...}\=B.B\...|
00000080  99 81 00 85 fc 40 02 00  66 ac 12 27  |.....@..f...' |
```

140 bytes

The initial beacon packet contains basic system information about the compromised host. This includes hostname, system stats (OS version and CPU speed), and whether any capture devices (i.e., webcams) are available. Because GH0ST RAT is an open-source tool, it is not difficult to modify the source code to implement additional features. In this case, the attacker modified the source code to remove the “Gh0st” string that is typically used for default GH0ST RAT communications at the beginning of the packet header and replaced it with the string “Adobe”. Following the string, Adobe is the compressed size of the message (`0x8C000000` in the second box) and the uncompressed size (`0xE0000000` in the third box). Through analyzing the first few bytes of the message, you could determine that `0x8C000000` translates to the decimal 140. You identify the payload size here is also 140 bytes. Translating `0xE0000000` from hex to a digit indicates the uncompressed size should be 224 bytes.

```
$ cat ghostrat.tcpflow.bin | hexdump -C
00000000  41 64 6f 62 65 8c 00 00  00 e0 00 00 00 78 9c 4b  |Adobe.....x.K|
00000010  63 66 60 98 c3 c0 c0 c0  0a c4 8c 40 ac c1 c5 c0  |cf`.....@....|
00000020  c0 04 a4 83 53 8b ca 32  93 53 15 02 12 93 b3 15  |....S..2.S.....|
00000030  8c 19 18 1c 58 5a 19 78  19 20 40 86 c1 02 ac 66  |....XZ.x. @....f|
00000040  01 50 40 01 88 05 1f 33  c0 01 23 d8 14 06 86 48  |.P@....3..#....H|
00000050  20 7e 00 55 c3 c3 07 64  28 30 30 5c 00 9a c1 c8  |~.U...d(00\....|
00000060  80 07 30 32 82 1d 01 52  c3 0c e5 ef e5 64 60 a8  |..02...R....d`..|
00000070  07 b2 7c 1d 7d c2 1d 83  5c 3d 42 fd 42 5c 83 f0  |...|...}\=B.B\...|
00000080  99 81 00 85 fc 40 02 00  66 ac 12 27  |.....@..f...' |
```



Example 4: GH0ST RAT (3)

The screenshot shows the CyberChef interface with a recipe for decompressing a GH0ST RAT payload. The input is a series of hex bytes. The recipe consists of four operations: 'From Hexdump', 'Drop bytes' (start 0, length 5), 'Drop bytes' (start 0, length 4), and 'Zlib Inflate' (start index 0, initial output buffer 0, buffer expand Adaptive). The output shows the decompressed ASCII strings, which include recognizable patterns like 'Adobe.....x.K', 'Service Pack 3', and 'MALWAREHUNTER'.

Decompressing this payload is straightforward using CyberChef and the knowledge from the source code. The following recipe provides the necessary functionality:

```
[{"op": "From Hexdump", "args": []}, {"op": "Drop bytes", "args": [0, 5, false]}, {"op": "Drop bytes", "args": [0, 4, false]}, {"op": "Drop bytes", "args": [0, 4, false]}, {"op": "Zlib Inflate", "args": [0, 0, "Adaptive", false, false]}, {"op": "To Hexdump", "args": [16, false, false], "disabled": true}]
```

After performing decompression, you identify the decompressed message is 224 bytes and can pick out several additional strings of interest within the decompressed payload that look familiar, such as “Service Pack 3” and “MALWAREHUNTER”.

For a comprehensive look at GH0ST RAT and its command structure, see the MITRE chopshop project’s^[1] decoding script for the protocol^[2].

References:

- [1] <https://for572.com/uhrxo>
- [2] <https://for572.com/3o1rn>



Example 4: GH0ST RAT (4)

- Partial protocol decode

Output		OpCode 0x66 = IMPLANT_LOGIN
000000000	66 03 00 00 9c 00 00 00 05 00 00 00 01 00 00 00 f.....	time: 3ms length: 1891 lines: 14
00000010	28 0a 00 00 02 00 00 00 53 65 72 76 69 63 65 20 (.....Service	Major/Minor Win Version 0x05/0x01 = 5.1
00000020	50 61 63 6b 20 33 00 00 40 04 85 00 0d 00 00 00 Pack 3..@....	Build Number 0x280a = 2600
00000030	00 00 00 00 1c 00 38 00 02 00 00 00 a0 0d 00 00 8.....	Service Pack Service Pack 3
00000040	20 0d 00 00 11 e3 00 00 00 00 00 00 00 00 00 00 ã.....	CPU Speed 0xbd09 = 2493Mhz
00000050	01 00 01 00 00 00 00 00 59 00 00 00 e0 00 00 00 Y..à..	IP Address 127.0.0.1
00000060	a0 0d 00 00 0c 0e 00 00 20 00 00 d0 04 85 00 b....	Hostname MALWAREHUNTER
00000070	01 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 	
00000080	00 00 00 00 00 00 00 00 00 00 00 00 00 01 01 00 	
00000090	00 01 00 00 01 00 00 00 03 00 00 00 00 01 01 00 	
000000a0	bd 09 00 00 7f 00 00 01 4d 41 4c 57 41 52 45 48 MALWAREH	
000000b0	55 4e 54 45 52 00 00 00 00 00 00 00 00 00 00 00 UNTER.....	
000000c0	00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 	
000000d0	00 00 00 00 00 00 00 00 00 00 00 71 0f 00 00 00 q....	

If protocol decoding is required, it may be a manual process. A SANS Gold Paper by David M. Martin^[1] details the decoding process for the Gh0st RAT protocol using the Dshell framework, a part of which is overlaid on the decompressed extraction from the previous slide. You can see a number of values that are almost intuitively understood, but a great deal of additional protocol decoding would be needed to fully express the contents of this particular packet data. The scale of this task becomes even more daunting when you consider the large number of OpCode values the protocol provides – each with its own structure.

References:

- [1] <https://for572.com/408dq>

Lab 5.2



Undocumented Protocol Features

This page intentionally left blank.

Lab 5.2 Objectives: Undocumented Protocol Features



- Review real-world malware traffic
- Reverse engineer network protocols to identify:
 - Artifacts for detection of additional malicious activity and signature generation
 - Encoding/encryption algorithms used to mask the attacker's activity
 - Host information in undocumented protocol fields

This page intentionally left blank.

Lab 5.2 Takeaways: Undocumented Protocol Features



- Identifying IOCs from undocumented protocols is a key goal of protocol reverse engineering
- Find known fields to start, then use ancillary data from DNS, etc., to fill knowledge gaps
- Identify unique and consistent byte streams to generate useful IOCs
- Fully documenting an unknown protocol can take a lot of time and effort—pace and scoping are critical

This page intentionally left blank.

Investigation OPSEC and Threat Intel

This page intentionally left blank.

The Network Is Alive and OPSEC Is Critical

- Network environment brings unique challenges
- Research can tip off attackers
 - DNS lookups, blocking C2 traffic, pwned email/proxy/etc., servers
 - Attackers track the hunters (AKA you)
- Live collection and analysis needed
 - Offline analysis not always feasible
 - Understand implications of actions
 - Build “live” incident response plans



The concepts behind disk and even memory forensics are relatively sound and have been tried and tested extensively. We generally acquire a "snapshot" of data at a point in time, then perform repeated analysis in an offline fashion against that data. However, performing forensic analysis in the network realm brings with it a number of concerns that must be considered before engaging in the analytic process.

For one, our research itself—or even the collection process—can cause events to occur that can let an attacker know we’re looking for them. These could include something as seemingly benign as a DNS lookup, or implementing a perimeter firewall block on a host known to be involved with malware command and control channels. If an attacker has successfully breached a proxy server, they could easily observe all web research being conducted on their malware, identities, or tactics/techniques/procedures. Of course, loading URLs associated with a malware infestation could quite likely be hazardous, bringing downloads of yet more unknown binary evil into the environment. For these reasons, we must be highly attuned to the operational security used by anyone involved with an investigation.

Another somewhat different dynamic is that given the sources of evidence in the network forensic world, much of our acquisition is done in a "live fashion". This requires human interaction with an operational device or system that often cannot be removed from service for very long—if at all. While forensically sound procedures for disk-based acquisition of high-availability systems have been used for some time, there is no easy analog for switches, routers, firewalls, or other devices the network investigator might be interested in. Snapshots and restore points are great for this on disk-based systems, but there is not likely to be such a feature built into switch firmware, for example.

For these reasons, we must consider how conducting an investigation in such a live environment can be accomplished without compromising the integrity of its results.

First, let’s talk about operational security, or OPSEC. Just as militaries and governments have a vested interest in protecting their information from the eyes and ears of adversaries, so must the forensic investigator. While OPSEC is a key factor for disk-based forensic investigations as well, it is fundamentally built into network-based investigations from the point of collection onward. Poor or nonexistent OPSEC safeguards could effectively stop an investigation dead in its tracks.

Ideally, network capture devices and sensors would access the network only through a one-way, read-only network tap, and subsequent analysis would be performed offline, in a segregated enclave, where no manner of network traffic could be released to the Internet at large. This would almost completely mitigate any risk of an adversary knowing what the investigation team was up to. However, reality must intercede. This is not always possible, nor is it always practical. Investigators can gain great benefits from live or real-time access to Internet resources—in some cases, such access is paramount to success.

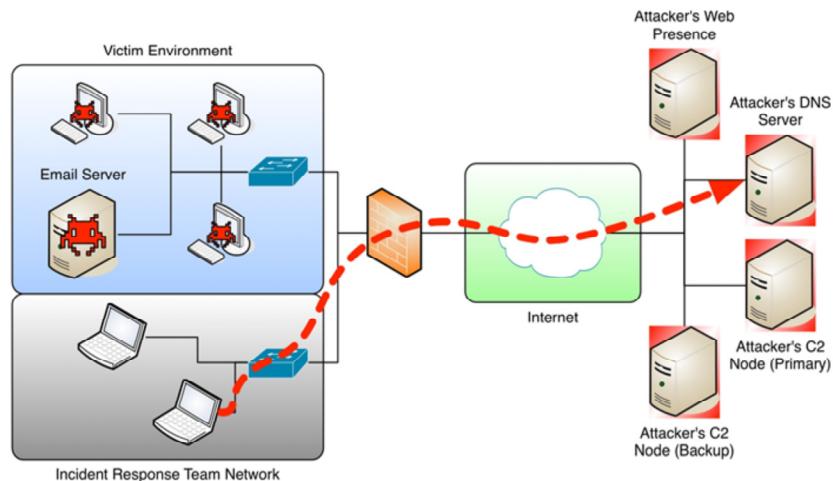
For example, maintaining duplicate databases and other records in an offline enclave sounds straightforward, but simply cannot achieve Internet scale with any degree of currency. It will never be realistic to have a near-live, offline duplicate of the entire DNS or WHOIS hierarchy, for example.

Another aspect of how important online access can be is seen in scoping an intrusion. By definition, the iterative, dynamic nature of the scoping process—determining which assets are compromised and the details surrounding an adversary’s compromise and continued access of the environment—require continual access to that environment as well as the Internet. How else would it be possible to determine whether a particular type of HTTP traffic was related to an advertising company’s cookie tracking or an attacker’s spear phishing campaign?

We will examine some of the more common concerns with regard to investigation OPSEC in turn.

Live Research Complications: DNS Activity

- May reach attacker's own DNS server
- Uniqueness of DNS query and source network may be clear giveaway
 - Great use case for self-collected Passive DNS



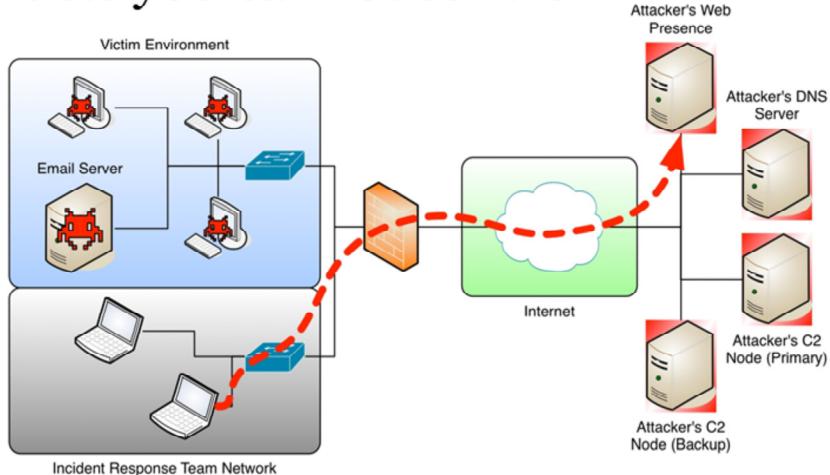
Let's start with DNS lookups. As we discussed previously, DNS traffic is some of the most common, fundamental traffic on our networks today. However, attackers often control their own DNS servers, giving them full and complete visibility to who is querying their domains, when, from where, and how often. Simply making a query to determine the current IP address for the “c2.badguy.evil” hostname could immediately let an attacker know that an in-house or contracted incident response team is on their trail.

Because malware often uses time-based hostnames through the use of deterministic name generation algorithms, an investigator performing open-source research on a domain name that was only valid last Tuesday could let the adversary know which generation of malware has been put under the microscope. The attacker could then adapt to this new landscape, or at least know how far they are ahead of the investigation.

This is also a risk to consider when building and configuring capture and monitoring devices. A system that performs real-time DNS lookups on all observed IP addresses will not only cause a great deal of network traffic, it would also give the adversary a good understanding of the sophistication of the victim's defensive posture.

Live Research Complications: OSINT Research

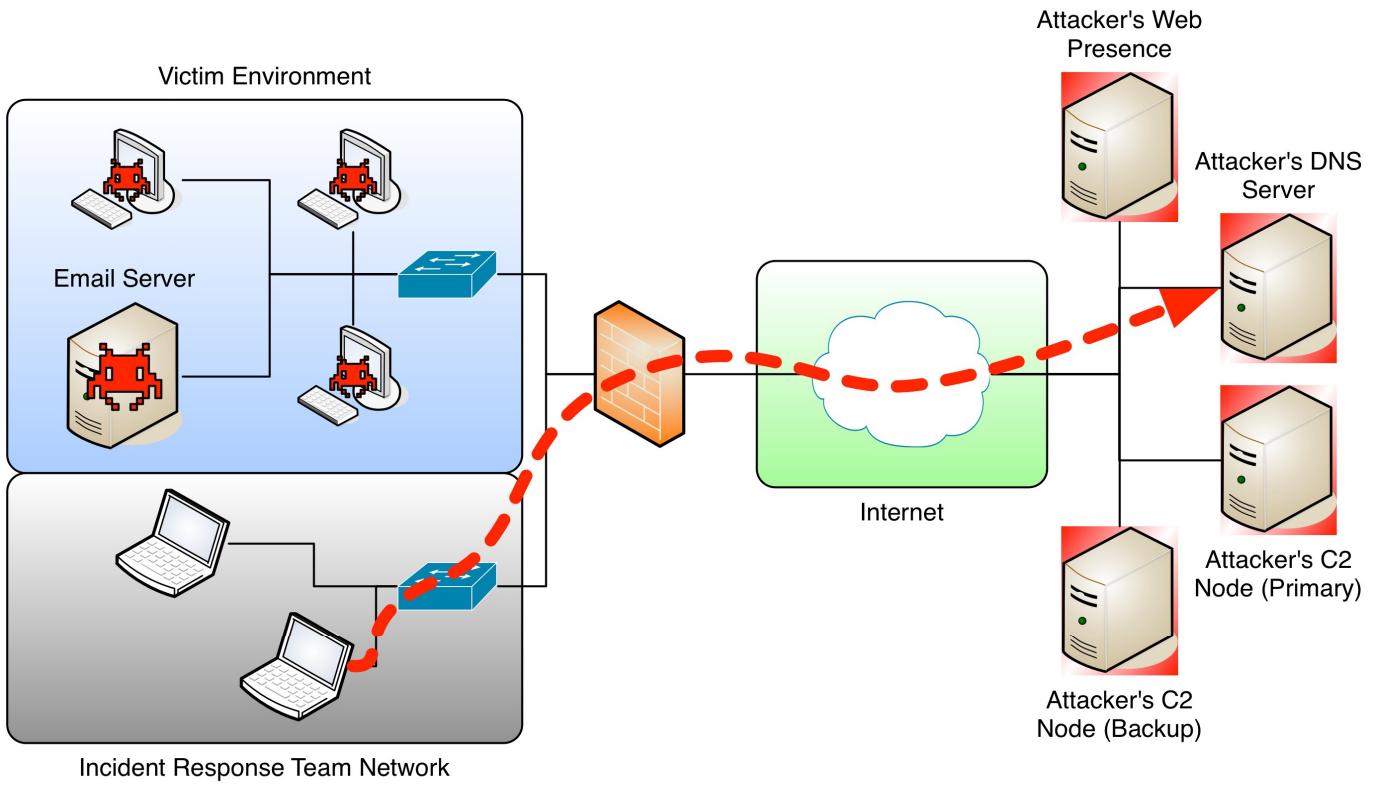
- Load web content from attacker's infrastructure
- Create public artifacts you cannot control
 - AV/malware scanning services
 - Social graphs
 - Location tracking



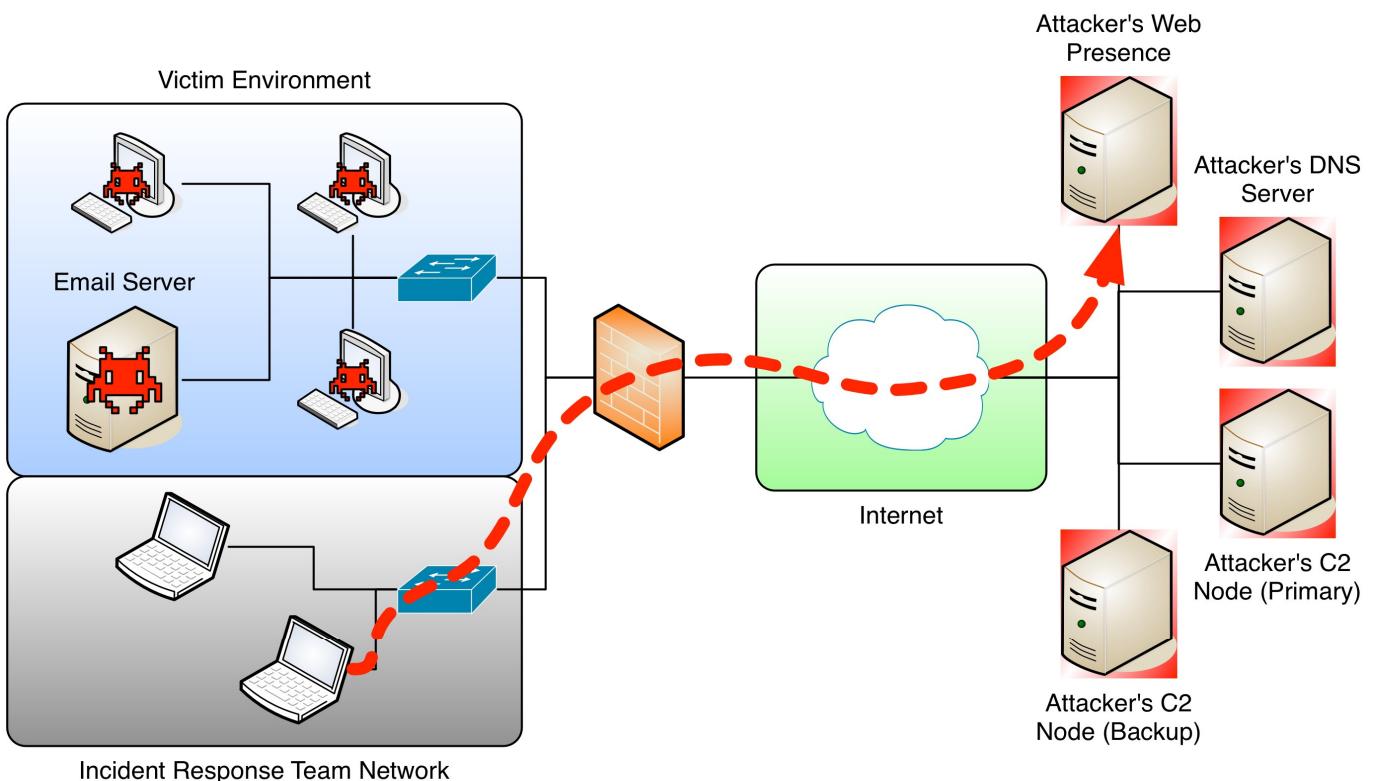
Similarly, basic Open-Source Intelligence (OSINT, not to be confused with Open-Source software) can provide the adversary with an early warning system of sorts. Especially if they control any sources of information (e.g., web servers) where you may find their name, investigators could quickly cause a similar situation to that of DNS lookups on the previous slide. An attacker that runs their own server can mine access logs, search terms, time frames, and other seemingly benign data sources to establish a pretty accurate picture of your research. Although many attackers won't have such a high level of counter-investigative capabilities, we must plan for well-funded and well-placed groups to have at least the same level of professional expertise as we do. Every technique we use to track an adversary can be used against us if we access infrastructure that they can also access.

Even in research using hosted platforms and social media sites, the collections and tracking procedures used by the site can compromise an investigation. A startling example of this is the LinkedIn.com “People who looked at this profile also looked at...” sidebar section. If an investigation team uses a single LinkedIn account—whether real or fake—profiles can quickly become associated by the LinkedIn.com software. It’s quite hair-raising to see all of your supposedly disconnected subjects/suspects show up in that sidebar at the same time. Location-based services are adding a layer of concern to this, as our mobile devices can give up a significant amount of sensitive information without our knowledge. Although more of an issue for physical investigations, the trend toward “unwitting transparency” is a real risk to our OPSEC.

Another way your own habits and behavior could compromise an investigation’s OPSEC is through “leaked” tracking cookies. Many advertisers and social media sites (who themselves are arguably just advertisers) use cookies to track your activity across many different sites. If an investigator accidentally logs onto their own social media profile, then conducts research related to a case, some level of their identity could be passed along to a third-party site—possibly even to the attackers themselves. The pervasive use of the Facebook “Like” widget, Google+ “+1” button, embedded Twitter posts, and other ubiquitous “Share this page” functions all involve countless cookies and other tracking features that could contribute to a loss of OPSEC.



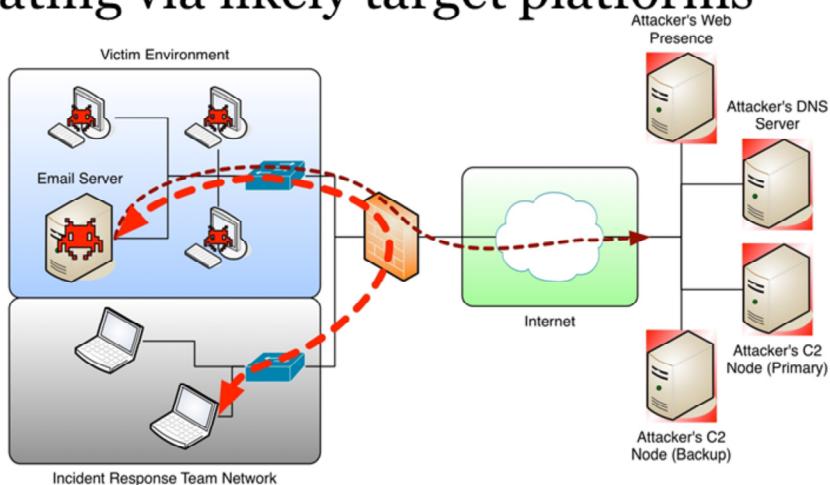
Hostname Lookups



OSINT Research

Live Research Complications: The Attacker Is Watching

- Everything you do can (and likely is) observed
- Avoid communicating via likely target platforms
 - Email servers
 - Collab portals
 - Business ops
 - Trouble tickets



Another major risk to consider, especially before an incident has been fully scoped, is that an attacker could easily be observing all of your actions within the compromised environment. While spreading within an environment, savvy attackers seize opportunities to monitor email, IM, and other communications. This enables them to identify the state of any incident response or investigations surrounding their activity. We must, therefore, assume that any such communications sent within a compromised environment could be fully read by the attacker.

Although this may sound like an exotic capability, remember that most email systems are now integrated into the domain environment (Active Directory, LDAP, or something similar). Therefore, a successful domain-level administrative compromise gives immediate and complete access to the email server as well—at no additional cost to the attacker. Many nation-state or organized crime adversaries have military or governmental backgrounds and seizing a victim's lines of communication has been a basic tenet of warfare and espionage for thousands of years. Little has changed today, aside from the methodology to achieve this advantage. Never underestimate the speed with which your communications will be targeted and/or compromised.

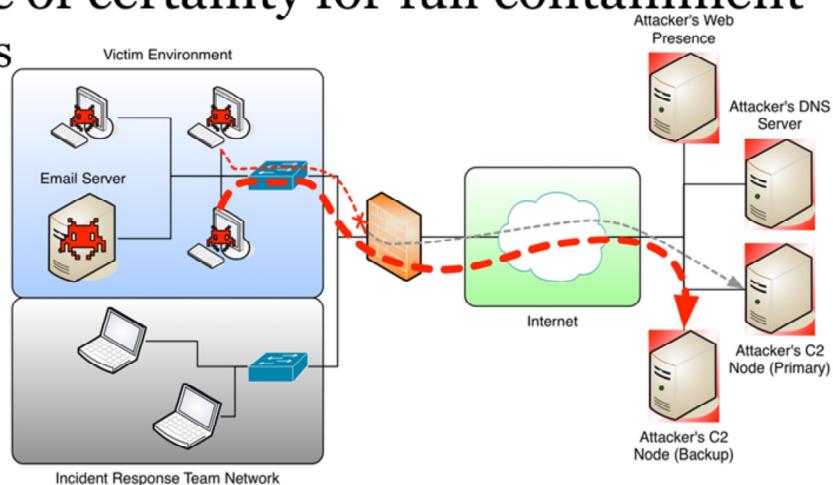
Put another way, expect that you and your environment are under surveillance. This is perhaps an evolution of the “Moscow Rules”^[1] adopted by Western intelligence services operating in Moscow before the fall of the Iron Curtain.

References:

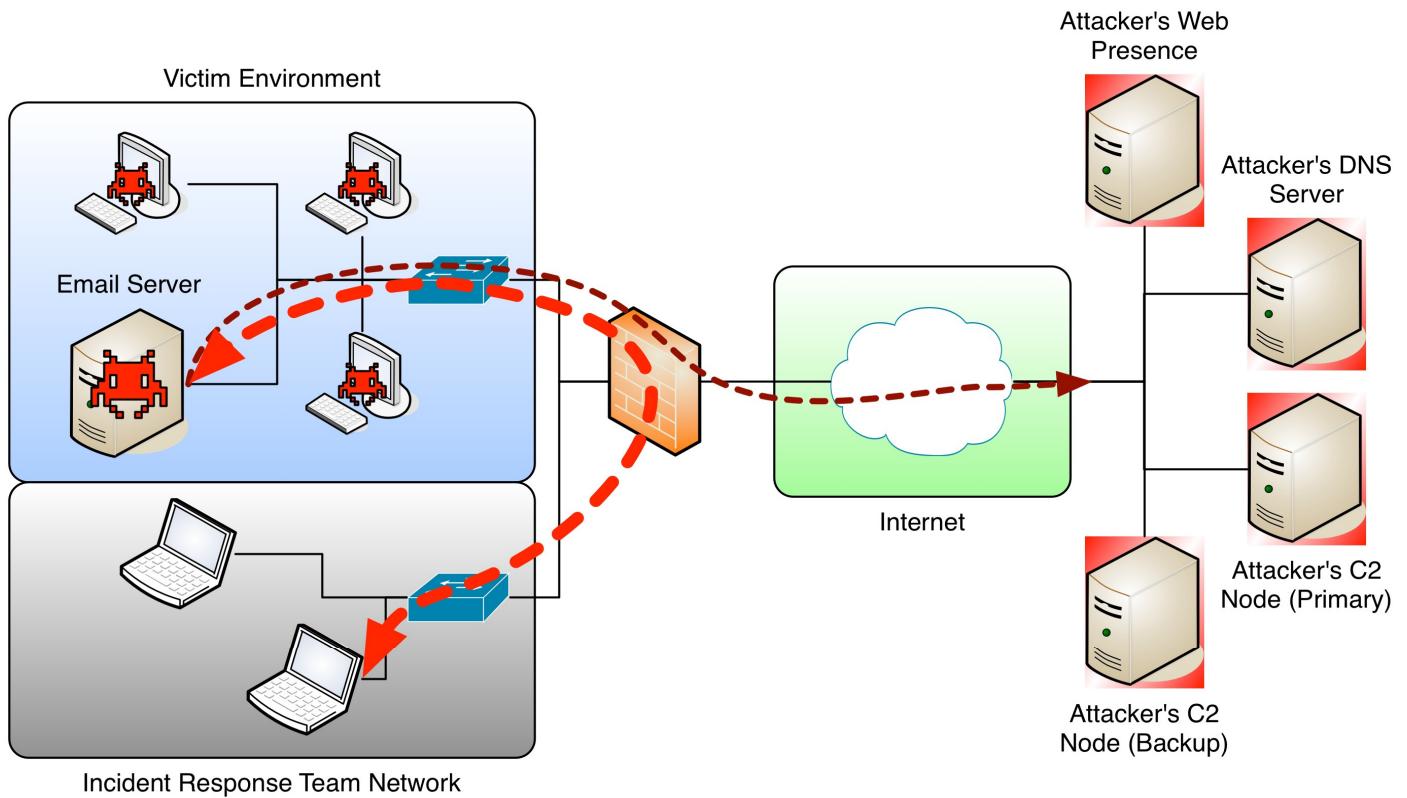
[1] <https://for572.com/wexz5>

Live Research Complications: Premature Traffic Block

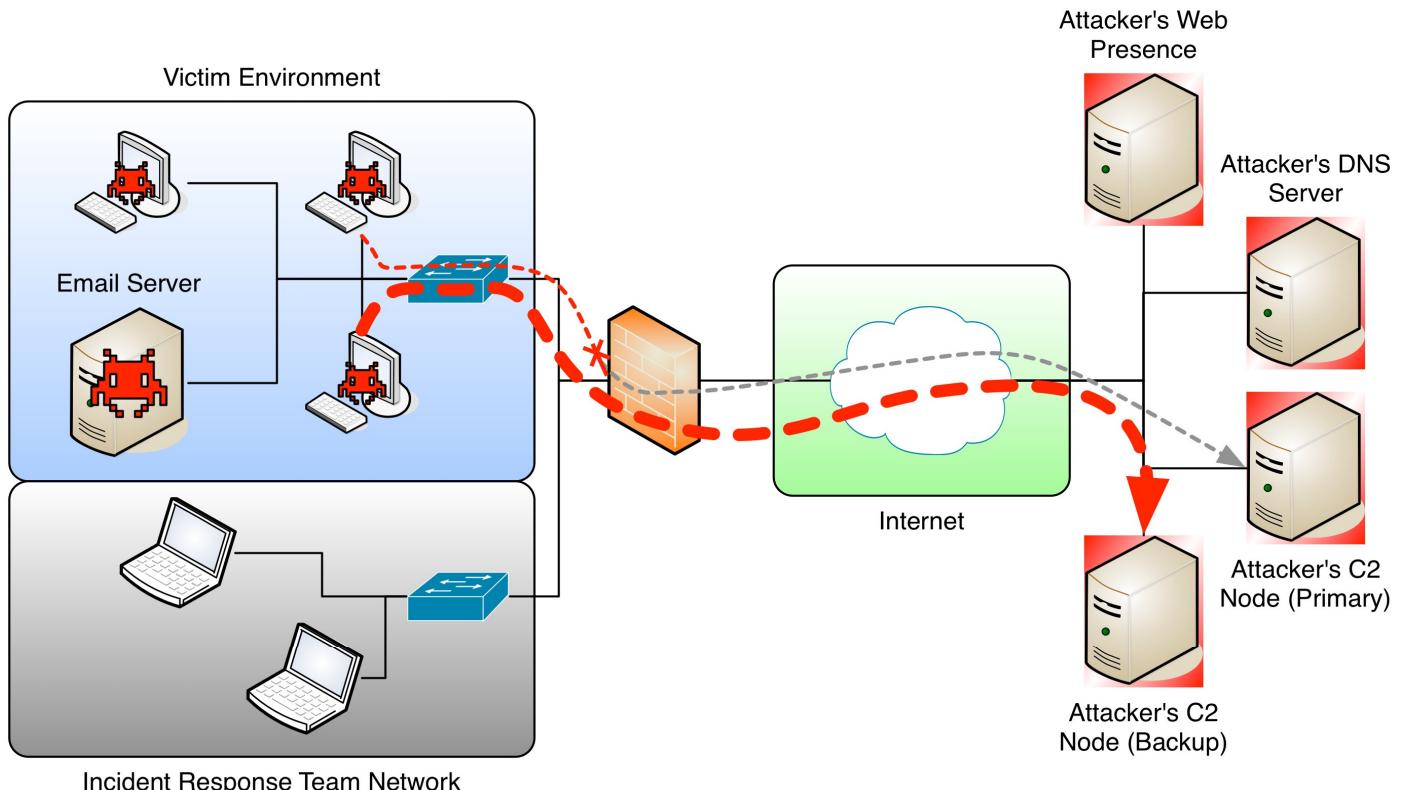
- Blocking C2 before full scoping and remediation
- Need high degree of certainty for full containment
 - Attacker often has 1+ backup C2
 - Avoid playing whack-a-mole



In what is possibly the most difficult “sell” to management, aggressive containment can pose a significant OPSEC risk. Often, acting too soon can prevent any hope of an effective remediation. For example, the legacy approach was to immediately block a newly identified command and control server, or to load a new malware traffic signature to the enterprise fleet of IPSes to prevent any further infestation. Attackers were severely trounced when this was done—but only for a short while. The attackers adapted and therefore so must the investigators.



The Attacker is Watching



Premature Traffic Block

Local Network Impact: DHCP/DNS Traffic

- DHCP/DNS are typically noninteractive
 - As soon as system links to switch, packets flow
- MDNS/UPnP broadcast services/software
 - Sends details via plaintext multicast
 - Often includes serial numbers, GUIDs, and other uniquely identifying data points

Although login sessions are typically a deliberate action taken by an investigator, there is also a great deal of automatic activity that occurs simply by connecting a system to the network. We discussed how DHCP is often the first evidence of a piece of hardware being introduced to the environment—this applies to the investigators as well! With most modern configurations, as soon as a device is connected to a switch port and a link is established, the device sends either a DHCPDISCOVER or DHCPREQUEST packet, seeking a lease from a DHCP server so it can proceed with network communications. Similarly, DNS traffic often follows, because most applications use hostnames rather than IP addresses while initiating communications.

Both of these protocols can provide critical evidence when examining suspect activities, so being able to rule out any investigator-initiated traffic will help to improve the overall signal-to-noise ratio. For example, seeing a known C2 hostname or IP address in a DNS query log could indicate further infection within the environment, or an examiner performing OSINT on the IP. (Of course, that would mean the examiner hasn't taken this class to learn about managing OPSEC posture yet!)

An interesting new development in DNS traffic has been the growing use of multicast DNS (also known as zero-configuration networking, Universal Plug and Play, or Bonjour). This protocol uses DNS traffic, though usually over UDP port 5353, and usually limited to local network segments. MDNS traffic advertises the hardware and software capabilities for a given device, such as whether it is a web server, file server, printer, etc. MDNS-aware clients use this traffic to find resources with which they want to interact. This protocol can be very noisy and provides a wealth of information about the devices using it. The simple act of an Apple system joining a network, for example, can create dozens or even hundreds of Bonjour packets to enter the network environment.

The next page contains several MDNS/bonjour messages, with interesting data points in bold text. Of course, these have been redacted to prevent possible misuse.

```

Apple Music v1.02.37: (Note various “*ID” values, music library name, system name and model)
{
    "id.orig_h": "fe80::18d3:938:fd7f:8bf4",
    "id.orig_p": 5353,
    "id.resp_h": "ff02::fb",
    "id.resp_p": 5353,
    "query": "philip hagen's library_music._hscp._tcp.local",
    "answers": [
        "TXT 20 model=MacBookPro16,1 TXT 10 osxvers=19 TXT 18
ecolor=157,157,160,
        "TXT 22 libid=A001FA33AE2A801A",
        "phil's macbook pro._atc._tcp.local",
        "phils-macbook-pro.local",
        "TXT 17 rpHN=e2e95c19fb22 TXT 12 rpFl=0x20000 TXT 11 rpVr=190.11 TXT 17
rpHA=43200f7e1b7 TXT 17 rpAD=d489f003a8e TXT 17 rpHI=3847168fe98e TXT 22
rpBA=C1:98:5C:21:40:8A",
        "_companion-link._tcp.local",
        "phil's macbook pro._companion-link._tcp.local",
        "phils-macbook-pro.local",
        "TXT 9 txtvers=1 TXT 10 dmv=131085 TXT 6 hQ=100 TXT 22
MID=0x7970a1A3391f2b7A TXT 14 Version=196621 TXT 19 iTSh Version=196624 TXT 9
DvSv=3081 TXT 10 DvTy=Music TXT 37 Machine Name=Philip Hagen's Library TXT 12
Ossi=0x1C941 TXT 10 iCSV=65540 TXT 39 hG=00000000-06e4-a753-6f22-60572b9bfda8
TXT 27 Database ID=233d2073D9441C6 TXT 23 Machine ID=4C8571c87A38 TXT 39
hC=b5dcae8d-4ab3-437e-9c82-ffffdc457256b,
        "_hscp._tcp.local",
        "philip hagen's library_music._hscp._tcp.local",
        "philip hagen's library_music._hscp._tcp.local",
        "phils-macbook-pro.local"
    ]
}

```

1Password v7.4.1 application

```

{
    "uid": "C11EwV2f4qUbDPNktf",
    "id.orig_h": "172.20.10.4",
    "id.orig_p": 5353,
    "id.resp_h": "224.0.0.251",
    "id.resp_p": 5353,
    "query": "bc0edb92-19a6-4210-937c-a14572da2c46@phil's macbook
pro._1password4._tcp.local",
    "answers": [
        "TXT 41 uuid=BC0EDB92-19A6-4210-937C-A14572DA2c46 TXT 24
updated=598505660.747534 TXT 23 name=Phil's MacBook Pro TXT 14
build=70401002,
        "_1password4._tcp.local",
        "bc0edb92-19a6-4210-937c-a14572da2c46@phil's macbook
pro._1password4._tcp.local",
        "phils-macbook-pro.local"
    ]
}

```

Apple iPhone Device Name

```
14:24:18.476019 IP 192.168.75.7.5353 > 224.0.0.251.5353: 0*- [0q] 1/0/5
(Cache flush) SRV Phil-Hagens-iPhone.local.:32498 0 0 (235)
```

Google Chrome browser automatically seeking a Google Chromecast device:

```
14:26:50.942621 IP 172.16.164.1.5353 > 224.0.0.251.5353: 0 PTR (QM)?
_googlecast._tcp.local.(40)
```

CentOS 7 system broadcasting system-level hostname (ivs-netflow.identityvector.com) and HTTP service name used for Ubiquiti UniFi network management (unifi.identityvector.net):

```
14:27:06.515911 IP 192.168.75.32.5353 > 224.0.0.251.5353: 0*- [0q] 3/0/0 PTR
UniFi Controller (unifi_identityvector_net)._http._tcp.local., (Cache flush)
SRV ivs-netflow-identityvector-com.local.:8080 0 0, (Cache flush) TXT
"path=/" "name=unifi.identityvector.net" (186)
```

QNAP TS-419P II Turbo NAS (Note device name, model number, firmware version information and serial number):

```
10:48:01.396033 IP 192.168.75.250.5353 > 224.0.0.251.5353: 0*- [0q] 4/0/0
(Cache flush) SRV HagenNAS1.local.:8080 0 0, (Cache flush) A 192.168.75.250,
PTR HagenNAS1._qddiscover._tcp.local., (Cache flush) TXT
"accessType=https,accessPort=443,model=TS-419P,displayModel=TS-419P
II,fwVer=4.0.5,fwBuildNum=20140117,serialNum=NAS_serial_number,webAdmPort=0,webAdmSslPort=0,webPort=0,webSslPort=0" (290)
```

Local Network Impact: Other Network Traffic

- “Automatically check for updates?”
“Send anonymous usage statistics?”
Licensing and anti-piracy verification
 - Most often use HTTP/HTTPS
- Creating traffic is not necessarily a problem
 - Need to minimize impact on the investigation
- Host-based activity notification may help
 - Little Snitch, GlassWire, OpenSnitch



Somewhat recent developments in software design have also created a situation in which network traffic is generated behind the scenes, without the user requesting or being aware of that activity. There are a number of situations that meet this description, but here are a few:

- Software often “calls out” to its developer to determine the latest version available and notifies the user if they need to update. With marketplaces such as Apple’s App Store and Google Play, these transactions can be accomplished with a single protocol for many different software titles. However, there are countless developers who have deployed their own version-check systems that you may see flying across the wire.
- Developers have also incorporated “usage statistic” collection in their works. These features, often briefly explained and enabled at installation, queue and send a variety of data points to the developer for logging and analysis. These are sent automatically, after the user agrees to their collection.
- Some software uses the network to validate licensing status and/or as a means of addressing piracy.

In all of these cases, the developer is free to use whatever protocol they want, but in practice, this most often occurs over HTTP or HTTPS because, in most environments, this traffic will almost certainly get through to its intended destination.

In all of these scenarios, it’s important to recognize and accept that introducing network traffic to the environment—potentially changing the environment being observed—is a reality that we need to handle. It can never be fully avoided, so our practices must adapt. It’s far more important to approach this problem by minimizing the impact on the environment and then fully understanding and explaining the residual activity that may be observed within the scope of the investigation. This also includes understanding how the environment may change this traffic, through the normal use of proxies, firewalls, routers, etc.

Some software solutions such as “Little Snitch” on macOS,^[1] GlassWire^[2] on Windows, OpenSnitch^[3] for Linux, may be helpful in preventing traffic “leakage”. However, as with any such tool, it is imperative that the forensically minded user knows the tool inside and out before using it operationally to avoid any accidental disclosure during an investigation.

References:

- [1] <https://for572.com/ko9wt>
- [2] <https://for572.com/h0fog>
- [3] <https://for572.com/mr-y9>

Research Risk Mitigations (I)

- Third-party research minimizes association
 - domaintools.com
 - archive.org, Google web cache (Watch out for images and other side-loaded resources!)
 - Proxy and VPN architectures
- Use air-gapped network when possible
 - No connectivity means no leak risk

Now that the “doom and gloom” is out of the way, we’ll take a look at some basic ways to mitigate the risks of compromising the OPSEC of an investigation. Of course, there are no perfect solutions in this business, but knowing how to make risks less severe is a key factor in our success.

An easy method is to use external, third-party sites to conduct network-centric research. By pushing your research off the victim’s network, the research more easily melts into background noise from the adversary’s point of view. Many of these services, such as domaintools.com, also add extra value to their offerings, such as cross-referencing across a larger data set, historical correlation, and more.

For some OSINT research requirements, sites like the Internet Archive Wayback Machine at archive.org, or even Google’s web cache can put reasonable distance between you and the source material. However, it’s important to fully vet a new service like this to understand its strengths and weaknesses. For example, the Google web cache does not cache or otherwise block image loads from the original source web server by default. This leads to a situation where the source server not only knows who is loading their content, but also that they’re using the Google web cache to do so. Depending on the adversary’s savvy, this could be a significant finding, resulting in a significant blow to your OPSEC posturing.

Another possible solution is to use any of the various privacy-focused proxy and VPN routing services or technologies. However, some of these services and providers have profiles of their own. For example, identifying an IP address as a TOR exit node is trivial on either a one-off or an automated basis.^[1]

However, it’s important to consider your adversaries’ skill level when determining which research methods are acceptable. Some attackers employ single-operation domains, so any research on that domain at all will uniquely identify the target that is conducting research.

Perhaps the most straightforward solution would be to prevent any network research whatsoever! It’s said that an unplugged computer encased in concrete and sunk to the bottom of the ocean is provably secure, although not very useful. In the same vein, though, a network-based investigation without network access will usually not be too

effective. So rather than a wholesale air-gap solution, it makes sense to consider what really **needs** online access, which actions are just more convenient to conduct with network access, and which actions are equally as effective on the quiet side of an air gap as they would be otherwise. Again, the key is to consider and mitigate.

References:

- [1] <https://for572.com/htp8z>

Research Risk Mitigations (2)

- Use separate network access
 - LTE hotspot in IR kit
 - Alternate access path
 - Throwaway VMs
- Never use forensic platforms for open research!



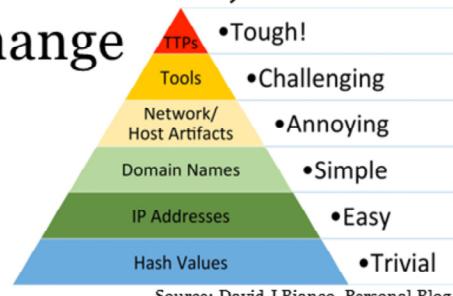
When conducting network-based research is absolutely required, a strategically appropriate solution must be used. As stated previously, operating within a victim environment can be very risky—especially before fully scoping the incident. For this reason, it is a good idea to use a fully separate connection, with no direct association to the victim (or the incident responders, if they are from an external team).

The most basic method is to keep an LTE hotspot in the “go bag” because it provides a fully separate line of communication. Another possible solution would be to install a separate land-based Internet connection for the IR team. This separate cable/DSL/fiber connection would be used solely for investigative purposes. New VPN-based technologies such as the IDVector^[1] provide a secure and anonymous means of communicating from a “hostile” environment such as public wireless. This USB dongle sets up a dedicated VPN path via any wireless environment, ensuring no network traffic leaks from the system to which it is connected. The IDVector suite includes an iOS client and will expand platform support in the future.

And, although this should go without saying, it is important enough to reiterate here:
NEVER use your forensic workstation or platform to perform Internet-based research! The risk of compromise of your entire analysis is simply too great to justify saving the time it would take to use a more appropriate research solution. Don’t become the “lesson learned” at the end of the investigation!

Guard Threat Intelligence Well

- Once profiled, attackers are easier to hunt; Once spotted, attackers are quick to change
 - Back out of environment entirely
 - Go to dormant state, wake up later
 - Change C2 servers and/or protocols
 - Take everything and leave quickly
 - Mandiant: PLA Unit 61398 retooled after APT-1 report
- Protect internally generated threat intelligence
 - Consider OPSEC implications of lost intelligence data
 - Share per local policy and industry best practices



Source: David J Bianco, Personal Blog

Engagement managers must walk a fine line between sharing necessary intelligence with system administrators, but doing so in a manner that does not give the attackers insight to the IR team's intel.

Today, any sufficiently advanced adversary considers the risk associated with being identified by the victim and includes several levels of contingencies to address that risk. So, a C2 block that is implemented before the full scope of the infestation is known, or the full capabilities of the malware have been determined, can be disastrous. Modern adversaries will quickly identify that the primary C2 was lost, then shift to a secondary method. Many will implement a “back off” period because they know the investigative team is closing in. They may defer or delay data theft transfer activity, pause C2 check-ins for days or weeks, or implement any number of other backup plans to maintain a foothold in the victim’s environment.

Sometimes, attackers may even move to “scuttle” an operation, stealing anything that’s already been staged for exfiltration, in expectation of soon being driven out of the environment. They may also opt to take destructive measures to conceal their presence or artifacts of their operations, including deploying ransomware as a smoke screen of sorts.

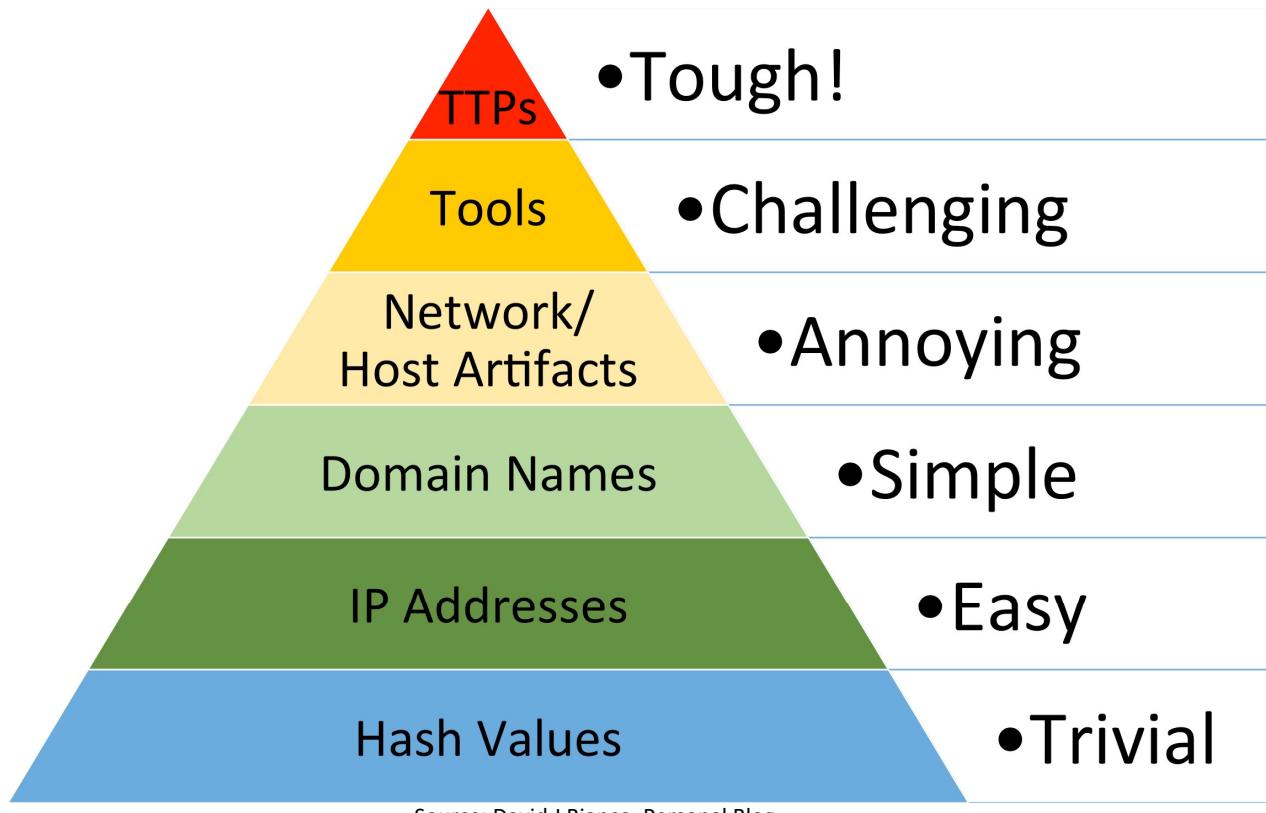
If an attacker ever does gain access to the team’s intelligence on them, they will certainly change their tactics to evade further detection and mitigation. This was documented by Mandiant when just three months after releasing their APT1 report^[1], the attack group they identified as PLA Unit 61398 was already retooling their operations and moving to new infrastructure^[2].

As covered in SANS FOR578, Cyber Threat Intelligence, generating your own threat intelligence is an important component of incident response and threat hunting. However, this intelligence would be incredibly useful to the attacker, allowing them to remain concealed and free to operate. Therefore, keeping that intelligence as closely held as possible is paramount to the operational security and eventual effectiveness of the incident responders’ actions.

The ease of changing tactics is reflected in David Bianco's "Pyramid of Pain"^[3] model. The PoP shows the increasing value of indicators in various categories due to the increasingly difficult nature for an attacker to change them. This model is a valuable perspective for incident responders to consult when considering the impact of losing threat intelligence and/or indicators during an investigation.

References:

- [1] <https://for572.com/c39t8>
- [2] <https://for572.com/j16kh>
- [3] <https://for572.com/pyramidofpain>



Plan to Share before the Incident: ISACs and ISAOs

- Information Sharing and Analysis Centers (ISACs)
 - Focus on critical infrastructure sectors
 - Defense Industrial Base (DIB), State/local governments, healthcare, financial services, aviation, etc.
- ISA Organizations (ISAOs)
 - Address sectors and groups not designated as critical infrastructure
 - Dozens of organizations accepted to the program



NationalCouncilofISACs



Note: ISACs and ISAOs are admittedly US-centric entities, but we acknowledge the international nature of our students. Therefore, please consider the information here as one possible avenue for organizations to share cyber threat intelligence. Of course those non-US-based students who support international organizations will still find relevance in the services such groups provide. However, there have been a great number of similar structures created across the world that may provide similar value to those in the information security industry regardless of where your geographical focus may be.

In 1998, US Presidential Decision Directive 63^[1] laid the groundwork to create groups that would share information about threats and vulnerabilities faced by designated critical infrastructure sectors. Soon thereafter, the notion of an Information Sharing and Analysis Center, or ISAC, was born. Today, nearly two dozen ISACs cover a wide variety of industry sectors. Some groups provide 24/7 warning functions, disseminate critical information about vulnerabilities and evolving threat actors, and more. ISACs typically aim to be more agile than similar (yet inherently bureaucratic) government organizations. This is a critical factor when dealing with aggressive, fast-paced threat groups.

These organizations are aligned by industry, which provides a narrower focus than the likes of the US-CERT or even The MITRE Corporation's CVE® program—which are both still very important, of course. Perhaps one of the most useful features the ISACs provide is a relationship structure in which member organizations can “get to know” their counterparts from other organizations within the same industry.

The National Council of ISACs^[2] provides an umbrella over all of the member ISACs and offers a good starting point for anyone wishing to learn more or become involved.

A more recent corollary to the ISAC program was the establishment of the Information Sharing and Analysis Organizations, or ISAOs^[3]. Despite the almost-too-similar name and resulting acronym, this is more a framework to form and maintain an information sharing group than a hierarchical set of named groups. They do not focus

specifically on a critical infrastructure sector, and anyone who follows the guidelines set forth by the ISAO Standards Organization itself can apply to be an ISAO. There are currently over 60 of these groups, with many more sure to come.

References:

- [1] <https://for572.com/u25c6>
- [2] <https://for572.com/87xdu>
- [3] <https://for572.com/q6c71>

Protect Threat Intelligence but Share Responsibly

- Community-defined “Traffic Light Protocol” (TLP)
 - Defines widely-recognized threat intel protection levels

Color	When should it be used?	Quick Reference
RED	Sources may use TLP:RED when information cannot be effectively acted upon by additional parties, and could lead to impacts on a party's privacy, reputation, or operations if misused.	Named recipients only
AMBER	Sources may use TLP:AMBER when information requires support to be effectively acted upon, but carries risks to privacy, reputation, or operations if shared outside of the organizations involved.	Organizational distribution
GREEN	Sources may use TLP:GREEN when information is useful for the awareness of all participating organizations as well as with peers within the broader community or sector.	Community-wide distribution
WHITE	Sources may use TLP:WHITE when information carries minimal or no foreseeable risk of misuse, in accordance with applicable rules and procedures for public release.	Unlimited

Regardless of whether you share information internally, within an ISAC or ISAO, or through other partnerships, knowing how to share—or not to share—data is important. To address this requirement, the security community developed the “Traffic Light Protocol”, or TLP,^[1] as a means of designating the sensitivity of intelligence that needs to be shared with partners but kept from adversaries. The TLP structure is not limited strictly to information security use, nor only to incident responders.

It is used in many nations and industry segments, so the definition for each level of protection is defined generically enough to apply widely. For example, someone who receives threat intelligence that is coated “TLP:RED” must not share that with any other party unless explicitly authorized by the originator. A “TLP:AMBER” communication can be shared within each recipient’s organization. “TLP:GREEN” designates information that can be freely shared within the overall community it is intended for, and “TLP:WHITE” indicates the information can be disseminated freely and publicly (within copyright constraints).

Those coming from a government or military background will be familiar with this style of dissemination restriction, and it is important for those in strictly commercial/civilian situations to understand so they can interact with the broader security community.

References:

[1] <https://for572.com/i68gk>

Community Resources to Consider

- GIAC Advisory Board Mailing List
 - Invitation-only – requires 90% on a GIAC exam
- Digital Forensics Discord Server
 - Info and join link: <https://for572.com/df-discord>



Building solid industry relationships doesn't always require a large organizational structure such as a ISAC/ISAO, or defined classification structures such as the TLPs.

Sometimes the most advantageous relationships can be built by being active in the broader digital forensics community. It doesn't take long to realize that our group of colleagues and peers is quite small.

GIAC provides an invitation-only mailing list for their Advisory Board members. The GIAC Advisory Board consists of certification holders who scored at least 90% on one or more GIAC exams. Those who earn an invitation to this group will receive instructions on joining the list directly from GIAC.

Another excellent community resource is the Digital Forensics Discord Server^[1], managed by Andrew Rathburn. This chat server is frequented by countless individuals and industry representatives and was the honored recipient of the 2020 Forensic 4:Cast award for “DFIR Resource of the Year”. Please note that this Discord server is not managed or endorsed by SANS.

References:

[1] <https://for572.com/df-discord>



Lab 5.3



Mini-Comprehensive Investigation

This page intentionally left blank.

Lab 5.3 Objectives: Mini-Comprehensive Investigation



- Review evidence from various sources and correlate the findings between them
- Understand the relative strengths and weaknesses of different sources of evidence and when each can be most useful during an investigation
- Use Wireshark's SSL/TLS decryption features to provide insight to otherwise impervious communications

This page intentionally left blank.

Lab 5.3 Takeaways: Mini-Comprehensive Investigation



- NetFlow provides overview, path to additional data
- HTTP server logs lack detail for POST actions
- pcap requires efficient data reduction
 - Encryption is a major hurdle; can't always mitigate
- Correlating evidence may require revisiting findings
 - Events rarely discovered in order they occurred
 - Reorder as needed to maintain chronological awareness
- Many activities leave artifacts of communication
 - Port scanning, brute-force attacks, data theft, etc.

This page intentionally left blank.



Encryption, Protocol Reversing, OPSEC, and Intel

©2022 Lewes Technology Consulting, LLC and Mat Oldham | All Rights Reserved | Version # FOR572_H01_01

Authors:

Phil Hagen, Lewes Technology Consulting, LLC
phil@lewestech.com | [@PhilHagen](https://twitter.com/PhilHagen)
Mat Oldham
mat.oldham@gmail.com | [@roujisecurity](https://twitter.com/roujisecurity)

Capstone Challenge Preparation

This page intentionally left blank.

Welcome to SRL

- Government-sponsored laboratory
- Specializes in metal alloys and bioengineering
- After two years of research, R&D may have identified the secret formula for Carbonadium
 - Carbonadium is the single most important Stark intellectual property
- Numerous other ongoing projects of varying degrees of sensitivity

This page intentionally left blank.

SRL Key Personnel

- mhill: CEO
 - Appointed by the President
 - IP: 172.16.7.11
- nfury: Sr. Manager
 - Twitter: @NickFury6
 - IP Address: 172.16.7.15
- tdungan: Sr. Researcher
 - Working on Carbonadium project
 - Twitter: @TimothyDungan
 - IP address: 172.16.6.11
- nromanoff: Proj. Manager
 - Twitter: @RomanoffNatasha
 - IP: 172.16.6.14
- rsydow: IT Admin
 - Twitter: *Not even once*
 - IP Address: 172.16.5.26
- cbarton: IT Security
 - Hired after 2012 breach
 - Twitter: *Not even once*
 - IP Address: 172.16.5.25

This page intentionally left blank.

Network Infrastructure

- Approximately 1k hosts
- Internal systems
 - Mix of Windows 7 and 10, Server 2012R2 and 2016, CentOS 7
- DMZ systems
 - Server 2012R2, CentOS 7
- Local and remote users
 - VPN, OWA, ActiveSync
- Mid-to-high security posturing
 - Clients firewalled from direct Internet access
 - TLS interception of web traffic at proxy
 - DNS reputation filtering
 - Inbound spam filtering
 - Aggressive network security monitoring and log collection/retention

This page intentionally left blank.

Threat Intelligence

- Believed Russian-linked group APT HAMMER targets metallurgy organizations
 - Use commodity as well as highly customized TTPs
 - Operate quickly and seemingly without worry of being detected
 - Well financed and organized
 - Known to be very effective



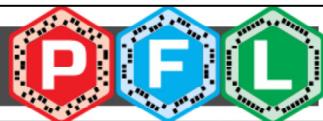
This page intentionally left blank.

Initial Suspicious Events: September 2018

- Web and email infrastructure intermittent
- IT troubleshooting and root cause analysis pointed to malicious activity
- On 5 September, IT Admin and Security Analyst initiated host-based forensics and incident response
- Network-based evidence was “frozen” so it would not be deleted with normal data retention pruning

This page intentionally left blank.

Available Evidence

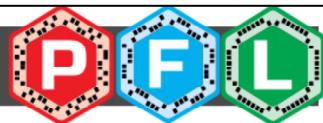


Evidence Type	Media File	Source	Start Date	End Date	Evidence Volume
NetFlow	A	base-fw: 105M flows; 3.8TB original traffic	2018-07-01	2018-09-19	1.6 GB
NetFlow	A	base-rtr: 32M flows; 931GB original traffic	2018-08-31	2018-09-19	492 MB
Log Data	A	Various, forwarded to base-elk: 206M log events from 13 sources	2018-05-24	2018-09-25	66 GB
Packet Capture	B	TAP-A (Sensor 1, eth2, Outside Firewall)	2018-08-14	2018-08-17	35 GB
Packet Capture	B	TAP-B (Sensor 1, eth3, DMZ)	2018-09-03	2018-09-06	2.6 GB
Packet Capture	B, C	TAP-D (Sensor 2, eth3, Services)	2018-09-03	2018-09-06	69 GB

All evidence volume figures refer to fully decompressed data

This page intentionally left blank.

Reminder: Previous Capstone-Consistent Labs



- Several course labs are consistent with the Capstone Scenario:
 - Lab 3.1: Visual NetFlow Analysis with SOF-ELK®
 - Lab 3.2: Tracking Lateral Movement with NetFlow
 - Lab 3.3: SMB Session Analysis & Reconstruction
 - Lab 4.1: Using NetworkMiner For Automated Extraction
 - Lab 5.1: SSL/TLS Profiling
 - Bonus Lab A: nfcapd Data Consolidation and Reduction
 - Bonus Lab B: Scaling With pcap_iterator.sh

This page intentionally left blank.

Your Mission, Should You Choose to Accept It...

- Did a breach occur?
- How did the breach initially occur?
- What systems were compromised?
- What material was taken?
- Where was the data exfiltrated?
- What malware was used?
- Any other indicators that can be used?

This page intentionally left blank.

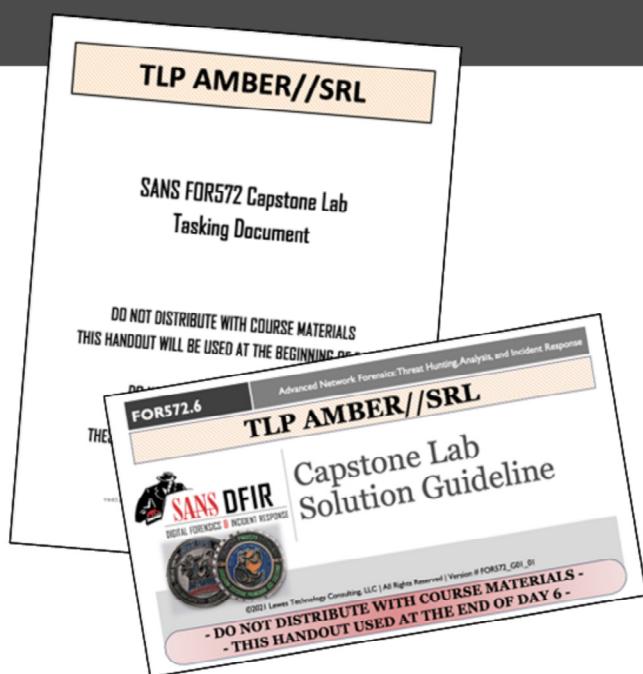
The Rules...

- Capstone is designed for teams of 3-4
- A “choose your own path” challenge
 - There is no singular correct solution!
- Use any tools you have—online or locally
 - Investigation OPSEC is still a significant concern
- Teams will present findings to the client
 - Maintain timeline and indicators of compromise
 - Identify knowledge gaps and how to address
 - Don’t wait until T-0:15:00 to start presentation!

This page intentionally left blank.

Capstone Documentation

- Instructor will provide additional documents and their passwords
 - More detailed kickoff information covered in Tasking Document
 - High-level overview of major findings covered in Solution Guideline



This page intentionally left blank.



This page intentionally left blank.

COURSE RESOURCES AND CONTACT INFORMATION

AUTHOR CONTACT



Phil Hagen/Lewes Technology Consulting, LLC
phil@lewesotech.com | @PhilHagen

Mat Oldham
mat.oldham@gmail.com | @roujisecurity

SANS INSTITUTE



11200 Rockville Pike, Suite 200
North Bethesda, MD 20852
301.654.SANS(7267)

DFIR RESOURCES



digital-forensics.sans.org
Twitter: @sansforensics

SANS EMAIL



GENERAL INQUIRIES: info@sans.org
REGISTRATION: registration@sans.org
TUITION: tuition@sans.org
PRESS/PR: press@sans.org

This page intentionally left blank.